

Applied Network Science with R

George G. Vega Yon, Ph.D.

2024-01-07

Contents

1 Preface	7
2 About the Author	9
I Applications	11
3 Introduction	13
4 R Basics	15
4.1 What is R	15
4.2 How to install packages	15
4.3 Prerequisites	15
4.4 A gentle Quick n' Dirty Introduction to R	16
5 Network Nomination Data	25
5.1 Data preprocessing	25
5.2 Creating a network	27
5.3 Network descriptive stats	32
5.4 Plotting the network in igraph	35
5.5 Statistical tests	41
6 Exponential Random Graph Models	43
6.1 A naïve example	44
6.2 Estimation of ERGMs	46
6.3 The ergm package	48
6.4 Running ERGMs	50
6.5 Model Goodness-of-Fit	53
6.6 More on MCMC convergence	65
6.7 Mathematical Interpretation	66
6.8 Markov independence	67
7 Using constraints in ERGMs	69
7.1 Example 1: Interlocking egos and disconnected alters	69
7.2 Example 2: Bi-partite networks	75

8 (Separable) Temporal Exponential Family Random Graph Models	81
9 Simulating and visualizing networks	83
9.1 Random Graph Models	84
9.2 Social Networks in Schools	84
9.3 Reading a network	87
9.4 Visualizing the network	90
10 Egocentric networks	101
10.1 Network files (graphml)	102
10.2 Person files	106
10.3 Ego files	109
10.4 Edgelist files	111
10.5 Putting all together	112
10.6 Saving the data	115
11 Hypothesis testing in networks	117
11.1 Comparing networks	117
11.2 Examples	120
12 Network diffusion with netdiffuseR	121
12.1 Network diffusion of innovation	121
12.2 The netdiffuseR R package	122
12.3 Simulation of diffusion processes	128
12.4 Statistical inference	137
13 Stochastic Actor Oriented Models	147
14 Power calculation in network studies	149
14.1 Example 1: Spillover effects in egocentric studies	149
14.2 Example 2: Spillover effects pre-post effect	153
14.3 Example 3: First difference	157
II Statistical Foundations	161
15 Bayes' Rule	163
16 Markov Chain	165
16.1 Metropolis Algorithm	165
16.2 Metropolis-Hastings	165
16.3 Likelihood-free MCMC	166
17 Power and sample size	167
17.1 Error types	167
17.2 Example 1: Sample size for a proportion	167
17.3 Example 2: Sample size for a proportion (vis)	169

<i>CONTENTS</i>	5
A Datasets	173
A.1 SNS data	173
References	175

Chapter 1

Preface

Statistical methods for networked systems are present in most disciplines. Nonetheless, despite the language differences between disciplines, many methods developed to study specific types of problems can be helpful outside of their original context.

This project began as a part of a workshop that took place at USC's **Center for Applied Network Analysis**. Now, it is a personal project that I use to gather and study statistical methods to analyze networks, emphasizing social and biological systems. Moreover, the book will use statistical computing methods as a core component when developing these topics.

In general, we will, besides R itself, we will be using RStudio. For data management, we will use dplyr and data.table. The network data management and visualization packages we will use are igraph, netdiffuseR, the statnet suite, and netplot. Some of you may be wondering "what about ggplot2 and friends? What about tidyverse", well, my short answer is I jumped into R before all of that was that popular.

You can access the book's source code at <https://github.com/gvegayon/appliedsnar>.

Chapter 2

About the Author

I am a Research Assistant Professor at the **University of Utah's Division of Epidemiology**, where I work on studying Complex Systems using Statistical Computing. I was born and raised in Chile. I have over ten years of experience developing scientific software focusing on high-performance computing, data visualization, and social network analysis. My training is in Public Policy (M.A. UAI, 2011), Economics (M.Sc. Caltech, 2015), and Biostatistics (Ph.D. USC, 2020).

I obtained my Ph.D. in Biostatistics under the supervision of **Prof. Paul Marjoram** and **Prof. Kayla de la Haye**, with my dissertation titled "*Essays on Bioinformatics and Social Network Analysis: Statistical and Computational Methods for Complex Systems.*"

If you'd like to learn more about me, please visit my website at <https://ggvy.cl>.

Part I

Applications

Chapter 3

Introduction

Social Network Analysis and Network Science, have a long scholarly tradition. From social diffusion models to protein-interaction networks, these complex systems disciplines cover a wide range of problems across scientific fields. Yet, although these could be seen as wildly different, the object under the microscope is the same, networks.

With a long history (and insufficient levels of inter-discipline collaboration, if you allow me to say) of scientific advances happening in a somewhat isolated fashion, the potential for cross-pollination between disciplines within network science is immense.

This book is an attempt to compile the many methods available in the realm of complexity sciences, provide an in-depth mathematical examination—when possible—, and provide a few examples illustrating their usage.

Chapter 4

R Basics

4.1 What is R

A good reference book for both novice and advanced users is “The Art of R programming” (Matloff 2011)¹

4.2 How to install packages

Nowadays there are two ways of installing R packages (that I’m aware of), either using `install.packages`, which is a function shipped with R, or using the `devtools` R package to install a package from some remote repository other than CRAN, here are a few examples:

```
# This will install the igraph package from CRAN
> install.packages("netdiffuseR")

# This will install the bleeding-edge version from the project's GitHub repo!
> devtools::install_github("USCCANA/netdiffuseR")
```

The first one, using `install.packages`, installs the CRAN version of `netdiffuseR`, whereas the line of code installs whatever version is published on <https://github.com/USCCANA/netdiffuseR>, which is usually called the development version.

In some cases, users may want/need to install packages from the command line as some packages need extra configuration to be installed. But we won’t need to look at it now.

4.3 Prerequisites

To install R just follow the instructions available at <http://cran.r-project.org>.

¹Here a free pdf version distributed by the author.

RStudio is the most popular Integrated Development Environment (IDE) for R that is developed by posit.co. While having RStudio is not a requirement for using netdiffuseR, it is highly recommended.

To get RStudio just visit <https://www.rstudio.com/products/rstudio/download/>.

4.4 A gentle Quick n' Dirty Introduction to R

Some common tasks in R

0. Getting help (and reading the manual) is *THE MOST IMPORTANT* thing you should know about. For example, if you want to read the manual (help file) of the `read.csv` function, you can type either of these: `r ?read.csv ?"read.csv" help(read.csv)`
`help("read.csv")` If you are not fully aware of what is the name of the function, you can always use the *fuzzy search* `r help.search("linear regression") ??"linear regression"`
1. In R you can create new objects by either using the assign operator (`<-`) or the equal sign `=`, for example, the following 2 are equivalent: `r a <- 1 a = 1` Historically the assign operator is the most commonly used.
2. R has several type of objects, the most basic structures in R are `vectors`, `matrix`, `list`, `data.frame`. Here is an example creating several of these (each line is enclosed with parenthesis so that R prints the resulting element):

```
(a_vector <- 1:9)

## [1] 1 2 3 4 5 6 7 8 9

(another_vect <- c(1, 2, 3, 4, 5, 6, 7, 8, 9))

## [1] 1 2 3 4 5 6 7 8 9

(a_string_vec <- c("I", "like", "netdiffuseR"))

## [1] "I"          "like"       "netdiffuseR"

(a_matrix <- matrix(a_vector, ncol = 3))

##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

(a_string_mat <- matrix(letters[1:9], ncol=3)) # Matrices can be of strings too

##      [,1] [,2] [,3]
## [1,] "a"   "d"   "g"
## [2,] "b"   "e"   "h"
## [3,] "c"   "f"   "i"
```

```
(another_mat <- cbind(1:4, 11:14)) # The `cbind` operator does "column bind"

##      [,1] [,2]
## [1,]    1   11
## [2,]    2   12
## [3,]    3   13
## [4,]    4   14

(another_mat2 <- rbind(1:4, 11:14)) # The `rbind` operator does "row bind"

##      [,1] [,2] [,3] [,4]
## [1,]    1    2    3    4
## [2,]   11   12   13   14

(a_string_mat <- matrix(letters[1:9], ncol = 3))

##      [,1] [,2] [,3]
## [1,] "a"  "d"  "g"
## [2,] "b"  "e"  "h"
## [3,] "c"  "f"  "i"

(a_list <- list(a_vector, a_matrix))

## [[1]]
## [1] 1 2 3 4 5 6 7 8 9
##
## [[2]]
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

(another_list <- list(my_vec = a_vector, my_mat = a_matrix)) # same but with names!

## $my_vec
## [1] 1 2 3 4 5 6 7 8 9
##
## $my_mat
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9

# Data frames can have multiple types of elements, it is a collection of lists
(a_data_frame <- data.frame(x = 1:10, y = letters[1:10]))

##      x y
## 1    1 a
```

```
## 2 2 b
## 3 3 c
## 4 4 d
## 5 5 e
## 6 6 f
## 7 7 g
## 8 8 h
## 9 9 i
## 10 10 j
```

3. Depending on the type of object, we can access to its components using indexing:

```
a_vector[1:3] # First 3 elements
## [1] 1 2 3

a_string_vec[3] # Third element
## [1] "netdiffuseR"

a_matrix[1:2, 1:2] # A sub matrix
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5

a_matrix[,3] # Third column
## [1] 7 8 9

a_matrix[3,] # Third row
## [1] 3 6 9

a_string_mat[1:6] # First 6 elements of the matrix. R stores matrices by column.
## [1] "a" "b" "c" "d" "e" "f"

# These three are equivalent
another_list[[1]]
## [1] 1 2 3 4 5 6 7 8 9

another_list$my_vec
## [1] 1 2 3 4 5 6 7 8 9

another_list[["my_vec"]]
## [1] 1 2 3 4 5 6 7 8 9

# Data frames are just like lists
a_data_frame[[1]]
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
a_data_frame[,1]

## [1] 1 2 3 4 5 6 7 8 9 10
a_data_frame[["x"]]

## [1] 1 2 3 4 5 6 7 8 9 10
a_data_frame$x

## [1] 1 2 3 4 5 6 7 8 9 10
```

4. Control-flow statements

```
# The oldfashion forloop
for (i in 1:10) {
  print(paste("I'm step", i, "/", 10))
}

## [1] "I'm step 1 / 10"
## [1] "I'm step 2 / 10"
## [1] "I'm step 3 / 10"
## [1] "I'm step 4 / 10"
## [1] "I'm step 5 / 10"
## [1] "I'm step 6 / 10"
## [1] "I'm step 7 / 10"
## [1] "I'm step 8 / 10"
## [1] "I'm step 9 / 10"
## [1] "I'm step 10 / 10"

# A nice ifelse

for (i in 1:10) {

  if (i %% 2) # Modulus operand
    print(paste("I'm step", i, "/", 10, "(and I'm odd)"))
  else
    print(paste("I'm step", i, "/", 10, "(and I'm even)"))

}

## [1] "I'm step 1 / 10 (and I'm odd)"
## [1] "I'm step 2 / 10 (and I'm even)"
## [1] "I'm step 3 / 10 (and I'm odd)"
## [1] "I'm step 4 / 10 (and I'm even)"
## [1] "I'm step 5 / 10 (and I'm odd)"
## [1] "I'm step 6 / 10 (and I'm even)"
```

```

## [1] "I'm step 7 / 10 (and I'm odd)"
## [1] "I'm step 8 / 10 (and I'm even)"
## [1] "I'm step 9 / 10 (and I'm odd)"
## [1] "I'm step 10 / 10 (and I'm even)"

# A while
i <- 10
while (i > 0) {
  print(paste("I'm step", i, "/", 10))
  i <- i - 1
}

## [1] "I'm step 10 / 10"
## [1] "I'm step 9 / 10"
## [1] "I'm step 8 / 10"
## [1] "I'm step 7 / 10"
## [1] "I'm step 6 / 10"
## [1] "I'm step 5 / 10"
## [1] "I'm step 4 / 10"
## [1] "I'm step 3 / 10"
## [1] "I'm step 2 / 10"
## [1] "I'm step 1 / 10"

```

5. R has a very nice set of pseudo-random number generation functions. In general, distribution functions have the following name structure:

- Random Number Generation: `r[name-of-the-distribution]`, e.g. `rnorm` for normal, `runif` for uniform.
- Density function: `d[name-of-the-distribution]`, e.g. `dnorm` for normal, `dunif` for uniform.
- Cumulative Distribution Function (CDF): `p[name-of-the-distribution]`, e.g. `pnorm` for normal, `punif` for uniform.
- Inverse (quantile) function: `q[name-of-the-distribution]`, e.g. `qnorm` for the normal, `qunif` for the uniform.

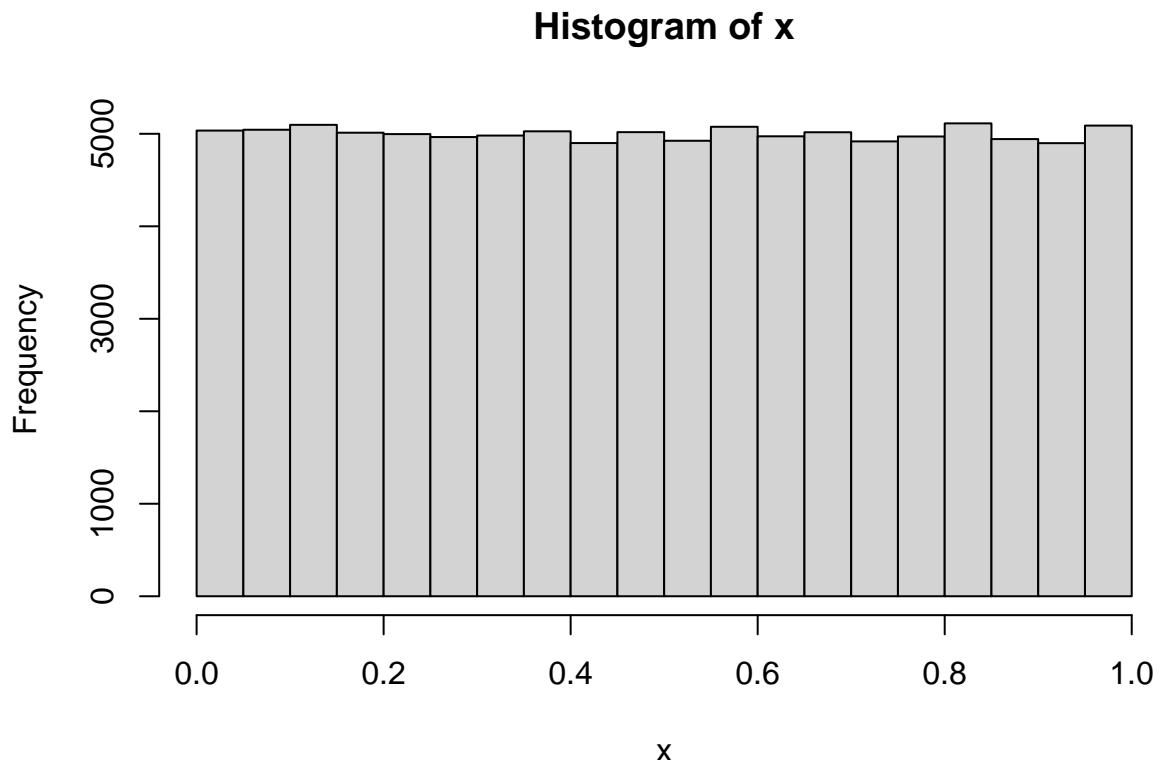
Here are some examples:

```

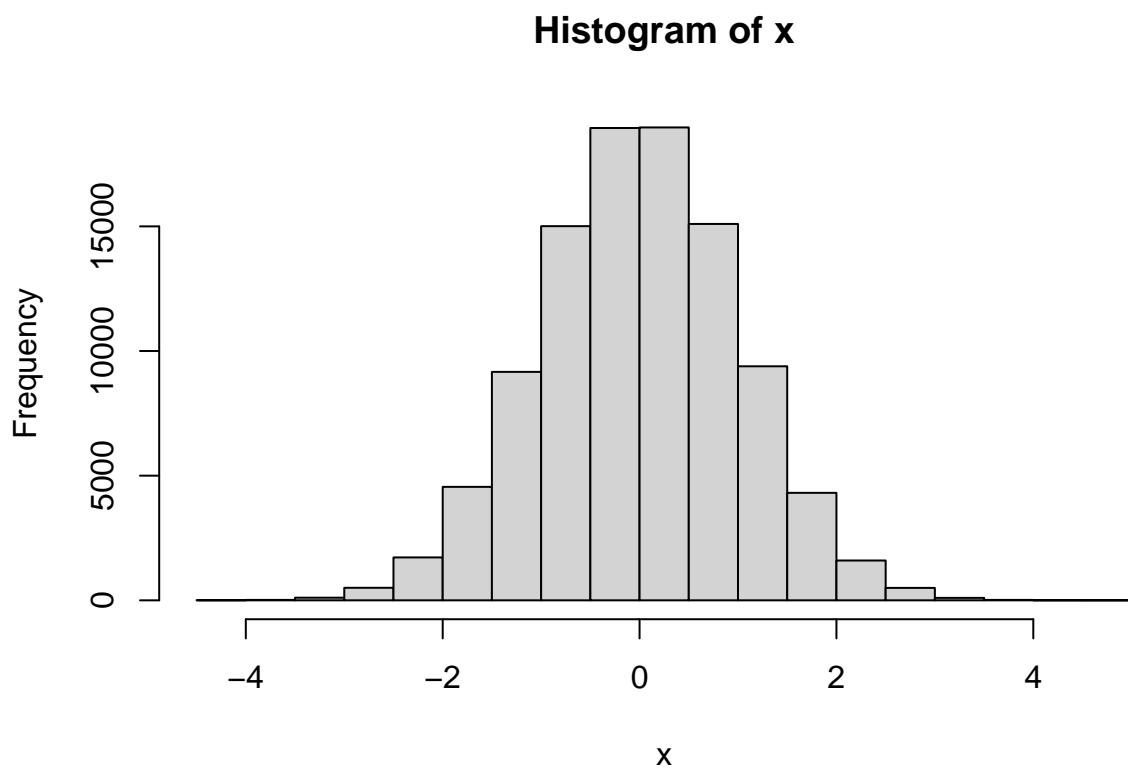
# To ensure reproducibility
set.seed(1231)

# 100,000 Unif(0,1) numbers
x <- runif(1e5)
hist(x)

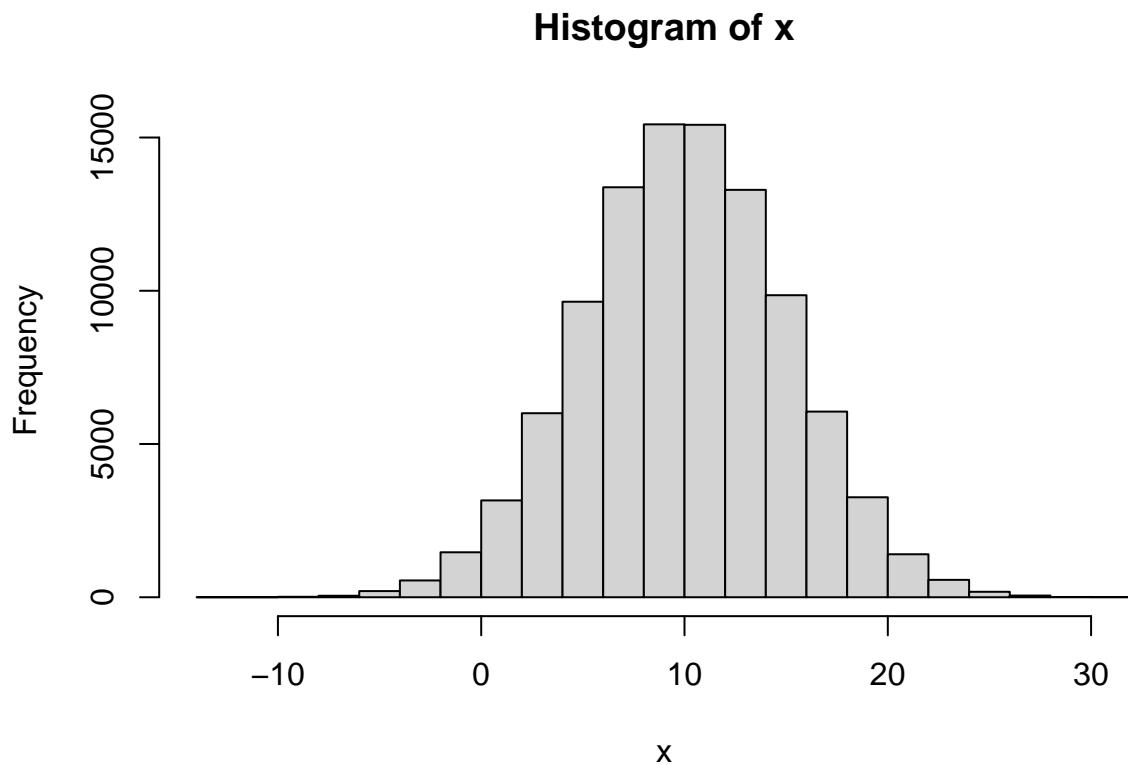
```



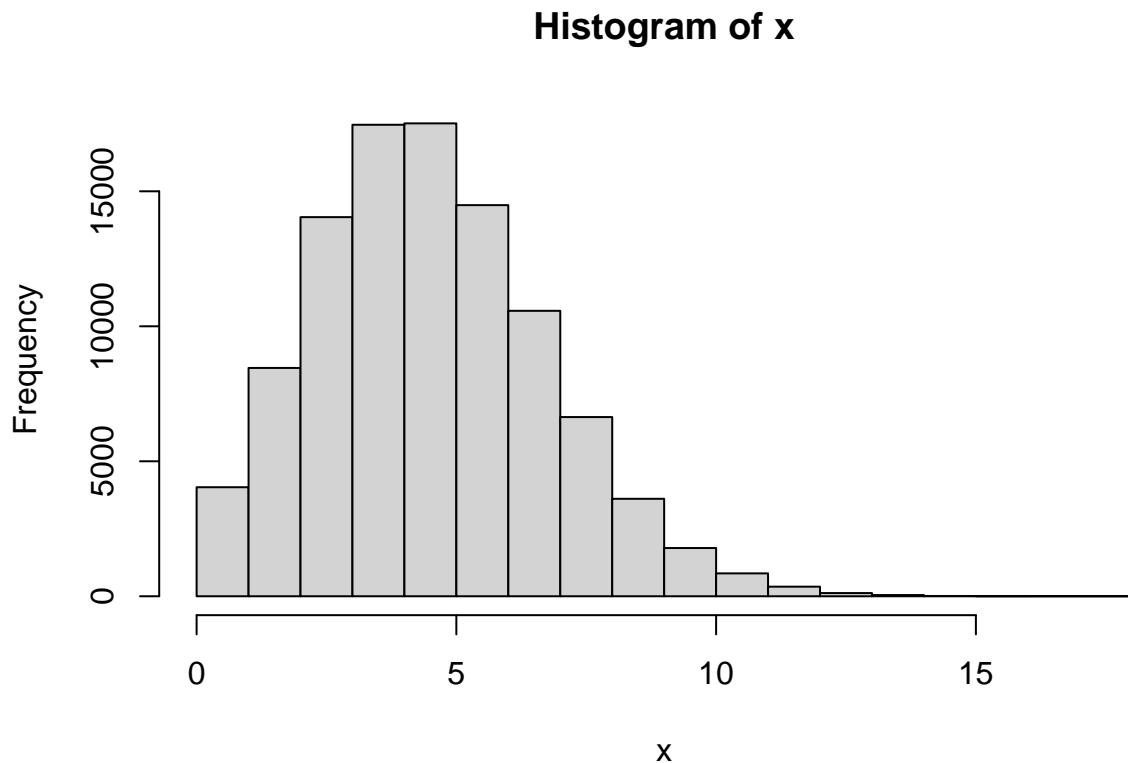
```
# 100,000 N(0,1) numbers
x <- rnorm(1e5)
hist(x)
```



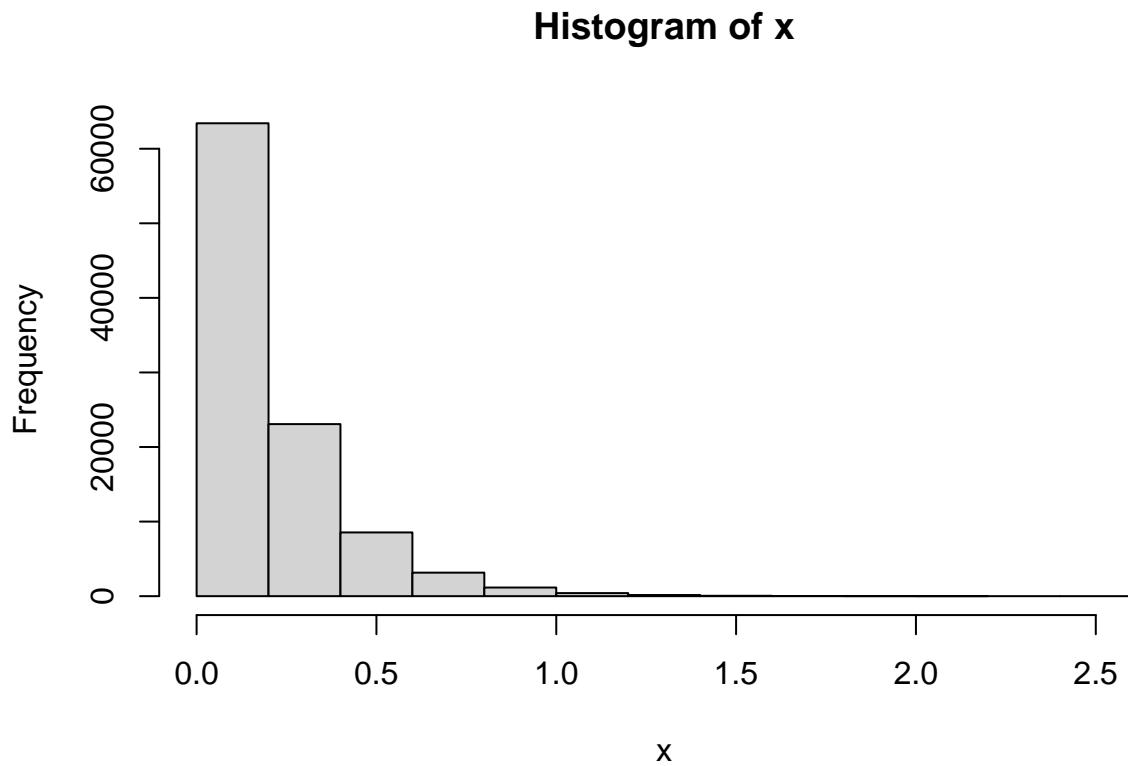
```
# 100,000 N(10,25) numbers
x <- rnorm(1e5, mean = 10, sd = 5)
hist(x)
```



```
# 100,000 Poisson(5) numbers
x <- rpois(1e5, lambda = 5)
hist(x)
```



```
# 100,000 rexp(5) numbers
x <- rexp(1e5, 5)
hist(x)
```



More distributions are available at [??Distributions](#).

For a nice intro to R, take a look at “The Art of R Programming” by Norman Matloff. For more advanced users, take a look at “Advanced R” by Hadley Wickham.

For this book, we need the following

R Core Team (2017b)

1. Install R from CRAN: <https://www.r-project.org/>
2. (optional) Install Rstudio: <https://rstudio.org>

While I find RStudio extremely useful, it is not necessary to use it with R.

Chapter 5

Network Nomination Data

You can download the data for this chapter here.

The codebook for the data provided here is in the appendix.

The goals for this chapter are:

1. Read the data into R,
2. Create a network with it,
3. Compute descriptive statistics
4. Visualize the network

5.1 Data preprocessing

5.1.1 Reading the data into R

R has several ways of reading data. Your data can be Raw plain files like CSV, tab-delimited, or specified by column width. To read plain-text data, you can use the `readr` package (Wickham, Hester, and Francois 2017). In the case of binary files, like Stata, Octave, or SPSS files, you can use the R package `foreign` (R Core Team 2017a). If your data is formatted as Microsoft spreadsheets, the `readxl` R package (Wickham and Bryan 2017) is the alternative to use. In our case, the data for this session is in Stata format:

```
library(foreign)

# Reading the data
dat <- foreign::read.dta("03-sns.dta")

# Taking a look at the data's first 5 columns and 5 rows
dat[1:5, 1:10]
```

```

##   photoid school hispanic female1 female2 female3 female4 grades1 grades2
## 1      1     111       1     NA     NA      0      0     NA     NA
## 2      2     111       1      0     NA     NA      0     3.0     NA
## 3      7     111       0      1      1      1      1     5.0    4.5
## 4     13     111       1      1      1      1      1     2.5    2.5
## 5     14     111       1      1      1      1      1     NA    3.0    3.5
##   grades3
## 1     3.5
## 2     NA
## 3     4.0
## 4     2.5
## 5     3.5

```

5.1.2 Creating a unique id for each participant

Now suppose that we want to create a unique id using the school and photo id. In this case, since both variables are numeric, a good way of doing it is to encode the id. For example, the last three numbers are the photoid and the first ones are the school id. To do this, we need to take into account the range of the variables:

```
(photo_id_ran <- range(dat$photoid))
```

```
## [1] 1 2074
```

As the variable spans up to 2074, we need to set the last 4 units of the variable to store the photoid. We will use `dplyr` (Wickham et al. 2017) and `magrittr` (Bache and Wickham 2014)] (the pipe operator, `%>%`) to create this variable, and we will call it... `id` (mind blowing, right?):

```
library(dplyr)
```

```

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##     filter, lag
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

```

```
library(magrittr)
```

```
(dat %>% mutate(id = school*10000 + photoid)) %>%
  head %>%
  select(school, photoid, id)
```

```
##   school photoid      id
```

```
## 1    111      1 1110001
## 2    111      2 1110002
## 3    111      7 1110007
## 4    111     13 1110013
## 5    111     14 1110014
## 6    111     15 1110015
```

Wow, what happened in the last three lines of code! What is that `%>%`? Well, that's the pipe operator, and it is an appealing way of writing nested function calls. In this case, instead of writing something like:

```
dat_filtered$id <- dat_filtered$school * 10000 + dat_filtered$photoid
subset(head(dat_filtered), select = c(school, photoid, id))
```

5.2 Creating a network

- We want to build a social network. For that, we either use an adjacency matrix or an edgelist.
- Each individual of the SNS data nominated 19 friends from school. We will use those nominations to create the social network.
- In this case, we will create the network by coercing the dataset into an edgelist.

5.2.1 From survey to edgelist

Let's start by loading a couple of handy R packages. We will load `tidyverse` (Wickham and Henry 2017) and `stringr` (Wickham 2017). We will use the first, `tidyverse`, to reshape the data. The second, `stringr`, will help us processing strings using *regular expressions*¹.

```
library(tidyverse)
library(stringr)
```

Optionally, we can use the `tibble` type of object, which is an alternative to the actual `data.frame`. This object is said to provide *more efficient methods for matrices and data frames*.

```
dat <- as_tibble(dat)
```

What I like from `tibbles` is that when you print them on the console, these actually look nice:

```
dat
```

```
## # A tibble: 2,164 x 100
##   photoid school hispanic female1 female2 female3 female4 grades1 grades2
```

¹Please refer to the help file ?'regular expression' in R. The R package `rex` (Ushey, Hester, and Krzyzanowski 2017) is a very nice companion for writing regular expressions. There's also a neat (but experimental) RStudio add-in that can be very helpful for understanding how regular expressions work, the `regexplain` add-in.

```

##      <int>  <int>  <dbl>  <int>  <int>  <int>  <int>  <dbl>  <dbl>
## 1     1     111     1     NA     NA     0     0     NA     NA
## 2     2     111     1     0     NA     NA     0     3     NA
## 3     7     111     0     1     1     1     1     5     4.5
## 4    13     111     1     1     1     1     1     2.5    2.5
## 5    14     111     1     1     1     1     NA     3     3.5
## 6    15     111     1     0     0     0     0     2.5    2.5
## 7    20     111     1     1     1     1     1     2.5    2.5
## 8    22     111     1     NA     NA     0     0     NA     NA
## 9    25     111     0     1     1     NA     1     4.5    3.5
## 10   27     111     1     0     NA     0     0     3.5    NA
## # i 2,154 more rows
## # i 91 more variables: grades3 <dbl>, grades4 <dbl>, eversmk1 <int>,
## #   eversmk2 <int>, eversmk3 <int>, eversmk4 <int>, everdrk1 <int>,
## #   everdrk2 <int>, everdrk3 <int>, everdrk4 <int>, home1 <int>, home2 <int>,
## #   home3 <int>, home4 <int>, sch_friend11 <int>, sch_friend12 <int>,
## #   sch_friend13 <int>, sch_friend14 <int>, sch_friend15 <int>,
## #   sch_friend16 <int>, sch_friend17 <int>, sch_friend18 <int>, ...
# Maybe too much piping... but its cool!
net <- dat %>%
  select(id, school, starts_with("sch_friend")) %>%
  gather(key = "varname", value = "content", -id, -school) %>%
  filter(!is.na(content)) %>%
  mutate(
    friendid = school*10000 + content,
    year      = as.integer(str_extract(varname, "(?<=[a-z])[0-9]")),
    nnom      = as.integer(str_extract(varname, "(?<=[a-z][0-9])[0-9]+"))
  )

```

Let's take a look at this step by step:

1. First, we subset the data: We want to keep `id`, `school`, `sch_friend*`. For the later, we use the function `starts_with` (from the `tidyselect` package). The latter allows us to select all variables that start with the word "sch_friend", which means that `sch_friend11`, `sch_friend12`, ... will be selected.

```
dat %>%
  select(id, school, starts_with("sch_friend"))
```

```

## # A tibble: 2,164 x 78
##       id school sch_friend11 sch_friend12 sch_friend13 sch_friend14
##      <dbl>  <int>        <int>        <int>        <int>        <int>
## 1 1110001     111         NA         NA         NA         NA
## 2 1110002     111        424        423        426        289
## 3 1110007     111        629        505         NA         NA
```

```

## 4 1110013    111      232      569      NA      NA
## 5 1110014    111      582      134      41      592
## 6 1110015    111      26       488      81      138
## 7 1110020    111      528      NA       492      395
## 8 1110022    111      NA       NA       NA      NA
## 9 1110025    111      135      185      553      84
## 10 1110027   111      346      168      559      5
## # i 2,154 more rows
## # i 72 more variables: sch_friend15 <int>, sch_friend16 <int>,
## #   sch_friend17 <int>, sch_friend18 <int>, sch_friend19 <int>,
## #   sch_friend110 <int>, sch_friend111 <int>, sch_friend112 <int>,
## #   sch_friend113 <int>, sch_friend114 <int>, sch_friend115 <int>,
## #   sch_friend116 <int>, sch_friend117 <int>, sch_friend118 <int>,
## #   sch_friend119 <int>, sch_friend21 <int>, sch_friend22 <int>, ...

```

2. Then, we reshape it to *long* format: By transposing all the sch_friend* to long format. We do this using the function gather (from the tidyverse package); an alternative to the reshape function, which I find easier to use. Let's see how it works:

```

dat %>%
  select(id, school, starts_with("sch_friend")) %>%
  gather(key = "varname", value = "content", -id, -school)

## # A tibble: 164,464 x 4
##       id school varname     content
##   <dbl> <int> <chr>     <int>
## 1 1110001    111 sch_friend11     NA
## 2 1110002    111 sch_friend11    424
## 3 1110007    111 sch_friend11    629
## 4 1110013    111 sch_friend11    232
## 5 1110014    111 sch_friend11    582
## 6 1110015    111 sch_friend11     26
## 7 1110020    111 sch_friend11    528
## 8 1110022    111 sch_friend11     NA
## 9 1110025    111 sch_friend11    135
## 10 1110027   111 sch_friend11    346
## # i 164,454 more rows

```

In this case, the key parameter sets the name of the variable that will contain the name of the variable that was reshaped, while value is the name of the variable that will hold the content of the data (that's why I named those like that). The -id, -school bit tells the function to “drop” those variables before reshaping. In other words, “reshape everything but id and school.”

Also, notice that we passed from 2164 rows to 19 (nominations) * 2164 (subjects) * 4 (waves) = 164464 rows, as expected.

3. As the nomination data can be empty for some cells, we need to take care of those cases, the NAs, so we filter the data:

```
dat %>%
  select(id, school, starts_with("sch_friend")) %>%
  gather(key = "varname", value = "content", -id, -school) %>%
  filter(!is.na(content))

## # A tibble: 39,561 x 4
##       id school varname     content
##   <dbl> <int> <chr>        <int>
## 1 1110002    111 sch_friend11     424
## 2 1110007    111 sch_friend11     629
## 3 1110013    111 sch_friend11     232
## 4 1110014    111 sch_friend11     582
## 5 1110015    111 sch_friend11      26
## 6 1110020    111 sch_friend11     528
## 7 1110025    111 sch_friend11     135
## 8 1110027    111 sch_friend11     346
## 9 1110029    111 sch_friend11     369
## 10 1110030   111 sch_friend11     462
## # i 39,551 more rows
```

4. And finally, we create three new variables from this dataset: friendid,, year, and nom_num (nomination number). All using regular expressions:

```
dat %>%
  select(id, school, starts_with("sch_friend")) %>%
  gather(key = "varname", value = "content", -id, -school) %>%
  filter(!is.na(content)) %>%
  mutate(
    friendid = school*10000 + content,
    year      = as.integer(str_extract(varname, "(?<=[a-z])[0-9]")),
    nnom      = as.integer(str_extract(varname, "(?<=[a-z][0-9])[0-9]+"))
  )

## # A tibble: 39,561 x 7
##       id school varname     content friendid    year  nnom
##   <dbl> <int> <chr>        <int>    <dbl> <int> <int>
## 1 1110002    111 sch_friend11     424    1110424     1     1
## 2 1110007    111 sch_friend11     629    1110629     1     1
## 3 1110013    111 sch_friend11     232    1110232     1     1
## 4 1110014    111 sch_friend11     582    1110582     1     1
## 5 1110015    111 sch_friend11      26    1110026     1     1
## 6 1110020    111 sch_friend11     528    1110528     1     1
## 7 1110025    111 sch_friend11     135    1110135     1     1
```

```
##  8 1110027    111 sch_friend11    346 1110346    1    1
##  9 1110029    111 sch_friend11    369 1110369    1    1
## 10 1110030    111 sch_friend11   462 1110462    1    1
## # i 39,551 more rows
```

The regular expression `(?<=[a-z])` matches a string preceded by any letter from *a* to *z*. In contrast, the expression `[0-9]` matches a single number. Hence, from the string "sch_friend12", the regular expression will only match the 1, as it is the only number followed by a letter. The expression `(?<=[a-z][0-9])` matches a string preceded by a lower case letter and a one-digit number. Finally, the expression `[0-9]+` matches a string of numbers—so it could be more than one. Hence, from the string "sch_friend12", we will get 2:

```
str_extract("sch_friend12", "(?<=[a-z])[0-9]")
## [1] "1"
str_extract("sch_friend12", "(?<=[a-z][0-9])[0-9]+")
## [1] "2"
```

And finally, the `as.integer` function coerces the returning value from the `str_extract` function from character to integer. Now that we have this edgelist, we can create an igraph object

5.2.2 igraph network

For coercing the edgelist into an igraph object, we will be using the `graph_from_data_frame` function in igraph (Csardi and Nepusz 2006). This function receives the following arguments: a data frame where the two first columns are “source” (ego) and “target” (alter), an indicator of whether the network is directed or not, and an optional data frame with vertices, in which’s first column should contain the vertex ids.

Using the optional `vertices` argument is a good practice since, by doing so, you are telling the function what ids that you are expecting to find. Using the original dataset, we will create a data frame name `vertices`:

```
vertexAttrs <- dat %>%
  select(id, school, hispanic, female1, starts_with("eversmk"))
```

Now, let's now use the function `graph_from_data_frame` to create an igraph object:

```
library(igraph)

ig_year1 <- net %>%
  filter(year == "1") %>%
  select(id, friendid, nnom) %>%
  graph_from_data_frame()
```

```

    vertices = vertexAttrs
  )

## Error in graph_from_data_frame(., vertices = vertexAttrs): Some vertex names in edge l

```

Ups! It seems that individuals are making nominations to other students not included in the survey. How to solve that? Well, it all depends on what you need to do! In this case, we will go for the *quietly-remove-em'-and-don't-tell* strategy:

```

ig_year1 <- net %>%
  filter(year == "1") %>%
  # Extra line, all nominations must be in ego too.
  filter(friendid %in% id) %>%
  select(id, friendid, nnom) %>%
  graph_from_data_frame(
    vertices = vertexAttrs
  )

ig_year1

## IGRAPH aa7fef5 DN-- 2164 9514 --
## + attr: name (v/c), school (v/n), hispanic (v/n), female1 (v/n),
## | eversmk1 (v/n), eversmk2 (v/n), eversmk3 (v/n), eversmk4 (v/n), nnom
## | (e/n)
## + edges from aa7fef5 (vertex names):
## [1] 1110007->1110629 1110013->1110232 1110014->1110582 1110015->1110026
## [5] 1110025->1110135 1110027->1110346 1110029->1110369 1110035->1110034
## [9] 1110040->1110390 1110041->1110557 1110044->1110027 1110046->1110030
## [13] 1110050->1110086 1110057->1110263 1110069->1110544 1110071->1110167
## [17] 1110072->1110289 1110073->1110014 1110075->1110352 1110084->1110305
## [21] 1110086->1110206 1110093->1110040 1110094->1110483 1110095->1110043
## + ... omitted several edges

```

So there we have our network with 2164 nodes and 9514 edges. The following steps: get some descriptive stats and visualize our network.

5.3 Network descriptive stats

While we could do all networks at once, in this part, we will focus on computing some network statistics for one of the schools only. We start by school 111. The first question that you should be asking yourself now is, “how can I get that information from the igraph object?.” Vertex and edges attributes can be accessed via the V and E functions, respectively; moreover, we can list what vertex/edge attributes are available:

```
list.vertex.attributes(ig_year1)

## [1] "name"      "school"     "hispanic"   "female1"    "eversmk1"   "eversmk2"   "eversmk3"
## [8] "eversmk4"

list.edge.attributes(ig_year1)

## [1] "nnom"
```

Just like we would do with data frames, accessing vertex attributes is done via the dollar sign operator \$. Together with the V function; for example, accessing the first ten elements of the variable `hispanic` can be done as follows:

```
V(ig_year1)$hispanic[1:10]
```

```
## [1] 1 1 0 1 1 1 1 1 0 1
```

Now that you know how to access vertex attributes, we can get the network corresponding to school 111 by identifying which vertices are part of it and pass that information to the `induced_subgraph` function:

```
# Which ids are from school 111?
school111ids <- which(V(ig_year1)$school == 111)

# Creating a subgraph
ig_year1_111 <- induced_subgraph(
  graph = ig_year1,
  vids  = school111ids
)
```

The `which` function in R returns a vector of indices indicating which elements pass the test, returning true and false, otherwise. In our case, it will result in a vector of indices of the vertices which have the attribute `school` equal to 111. With the subgraph, we can compute different centrality measures² for each vertex and store them in the `igraph` object itself:

```
# Computing centrality measures for each vertex
V(ig_year1_111)$indegree  <- degree(ig_year1_111, mode = "in")
V(ig_year1_111)$outdegree <- degree(ig_year1_111, mode = "out")
V(ig_year1_111)$closeness <- closeness(ig_year1_111, mode = "total")
V(ig_year1_111)$betweeness <- betweenness(ig_year1_111, normalized = TRUE)
```

From here, we can go back to our old habits and get the set of vertex attributes as a data frame so we can compute some summary statistics on the centrality measurements that we just got

```
# Extracting each vectex features as a data.frame
stats <- as_data_frame(ig_year1_111, what = "vertices")
```

²For more information about the different centrality measurements, please take a look at the “Centrality” article on Wikipedia.

```
# Computing quantiles for each variable
stats_degree <- with(stats, {
  cbind(
    indegree = quantile(indegree, c(.025, .5, .975), na.rm = TRUE),
    outdegree = quantile(outdegree, c(.025, .5, .975), na.rm = TRUE),
    closeness = quantile(closeness, c(.025, .5, .975), na.rm = TRUE),
    betweenness = quantile(betweenness, c(.025, .5, .975), na.rm = TRUE)
  )
})

stats_degree

##          indegree  outdegree   closeness  betweenness
## 2.5%        0       0.0005915148 0.000000000
## 50%         4       0.0007487833 0.001879006
## 97.5%       16      0.0008838413 0.016591048
```

The with function is somewhat similar to what dplyr allows us to do when we want to work with the dataset but without mentioning its name everytime that we ask for a variable. Without using the with function, the previous could have been done as follows:

```
stats_degree <-
  cbind(
    indegree = quantile(stats$indegree, c(.025, .5, .975), na.rm = TRUE),
    outdegree = quantile(stats$outdegree, c(.025, .5, .975), na.rm = TRUE),
    closeness = quantile(stats$closeness, c(.025, .5, .975), na.rm = TRUE),
    betweenness = quantile(stats$betweenness, c(.025, .5, .975), na.rm = TRUE)
)
```

Now we will compute some statistics at the graph level:

```
cbind(
  size = vcount(ig_year1_111),
  nedges = ecount(ig_year1_111),
  density = edge_density(ig_year1_111),
  recip = reciprocity(ig_year1_111),
  centr = centr_betw(ig_year1_111)$centralization,
  pathLen = mean_distance(ig_year1_111)
)

##          size  nedges   density     recip      centr pathLen
## [1,]  533 2638 0.009303277 0.3731513 0.02179154 4.23678
```

Triadic census

```
triadic <- triad_census(ig_year1_111)
triadic
```

```
## [1] 24059676    724389   290849     3619     3383     4401     3219     2997
## [9]      407       33       836      235      163      137      277       85
```

To get a nicer view of this, we can use a table that I retrieved from `?triad_census`. Moreover, we can normalize the triadic object by its sum instead of looking at raw counts. That way, we get proportions instead³

```
knitr:::kable(cbind(
  Pcent = triadic/sum(triadic)*100,
  read.csv("triadic_census.csv")
), digits = 2)
```

Pcent	code	description
95.88	003	A,B,C, the empty graph.
2.89	012	A->B, C, the graph with a single directed edge.
1.16	102	A<->B, C, the graph with a mutual connection between two vertices.
0.01	021D	A<-B->C, the out-star.
0.01	021U	A->B<-C, the in-star.
0.02	021C	A->B->C, directed line.
0.01	111D	A<->B<-C.
0.01	111U	A<->B->C.
0.00	030T	A->B<-C, A->C.
0.00	030C	A<-B<-C, A->C.
0.00	201	A<->B<->C.
0.00	120D	A<-B->C, A<->C.
0.00	120U	A->B<-C, A<->C.
0.00	120C	A->B->C, A<->C.
0.00	210	A->B<->C, A<->C.
0.00	300	A<->B<->C, A<->C, the complete graph.

5.4 Plotting the network in igraph

5.4.1 Single plot

Let's take a look at how does our network looks like when we use the default parameters in the `plot` method of the `igraph` object:

```
plot(ig_year1)
```

Not very nice, right? A couple of things with this plot:

1. We are looking at all schools simultaneously, which does not make sense. So, instead of plotting `ig_year1`, we will focus on `ig_year1_111`.

³During our workshop, Prof. De la Haye suggested using $\binom{n}{3}$ as a normalizing constant. It turns out that `sum(triadic) = choose(n, 3)!` So either approach is correct.

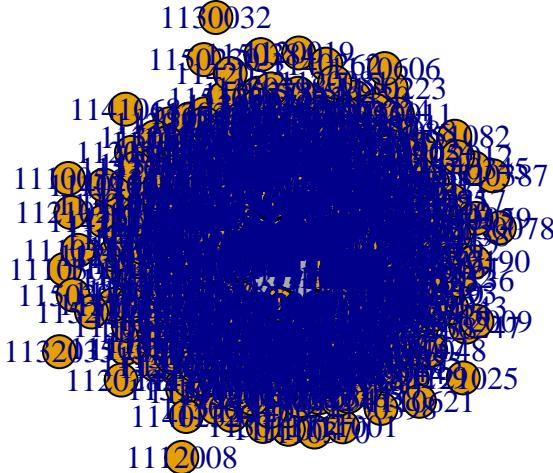


Figure 5.1: A not very nice network plot. This is what we get with the default parameters in `igraph`.

2. All the vertices have the same size and are overlapping. Instead of using the default size, we will size the vertices by indegree using the `degree` function and passing the vector of degrees to `vertex.size`.⁴
3. Given the number of vertices in these networks, the labels are not useful here. So we will remove them by setting `vertex.label = NA`. Moreover, we will reduce the size of the arrows' tip by setting `edge.arrow.size = 0.25`.
4. And finally, we will set the color of each vertex to be a function of whether the individual is Hispanic or not. For this last bit we need to go a bit more of programming:

```
col_hispanic <- V(ig_year1_111)$hispanic + 1
col_hispanic <- coalesce(col_hispanic, 3)
col_hispanic <- c("steelblue", "tomato", "white")[col_hispanic]
```

Line by line, we did the following:

1. The first line added one to all no NA values so that the 0s (non-Hispanic) turned to 1s and the 1s (Hispanic) turned to 2s.
2. The second line replaced all NAs with the number three so that our vector `col_hispanic` now ranges from one to three with no NAs in it.
3. In the last line, we created a vector of colors. Essentially, what we are doing here is telling R to create a vector of length `length(col_hispanic)` by selecting elements by index from the vector `c("steelblue", "tomato", "white")`. This way, if, for example, the first element of the vector `col_hispanic` was a 3, our new vector of colors would have a "white" in it.

⁴Figuring out what is the optimal vertex size is a bit tricky. Without getting too technical, there's no other way of getting *nice* vertex size other than just playing with different values of it. A nice solution to this is using `netdiffuseR::igraph_vertex_rescale` which rescales the vertices so that these keep their aspect ratio to a predefined proportion of the screen.

To make sure we know we are right, let's print the first 10 elements of our new vector of colors together with the original hispanic column:

```
cbind(
  original = V(ig_year1_111)$hispanic[1:10],
  colors   = col_hispanic[1:10]
)

##      original colors
## [1,] "1"      "tomato"
## [2,] "1"      "tomato"
## [3,] "0"      "steelblue"
## [4,] "1"      "tomato"
## [5,] "1"      "tomato"
## [6,] "1"      "tomato"
## [7,] "1"      "tomato"
## [8,] "1"      "tomato"
## [9,] "0"      "steelblue"
## [10,] "1"     "tomato"
```

With our nice vector of colors, now we can pass it to `plot.igraph` (which we call implicitly by just calling `plot`), via the `vertex.color` argument:

```
# Fancy graph
set.seed(1)
plot(
  ig_year1_111,
  vertex.size     = degree(ig_year1_111)/10 +1,
  vertex.label    = NA,
  edge.arrow.size = .25,
  vertex.color    = col_hispanic
)
```

Nice! So it does look better. The only problem is that we have a lot of isolates. Let's try again by drawing the same plot without isolates. To do so, we need to filter the graph, for which we will use the function `induced_subgraph`

```
# Which vertices are not isolates?
which_ids <- which(degree(ig_year1_111, mode = "total") > 0)

# Getting the subgraph
ig_year1_111_sub <- induced_subgraph(ig_year1_111, which_ids)

# We need to get the same subset in col_hispanic
col_hispanic <- col_hispanic[which_ids]
```




Figure 5.2: Friends network in time 1 for school 111.

```
# Fancy graph
set.seed(1)
plot(
  ig_year1_111_sub,
  vertex.size      = degree(ig_year1_111_sub)/5 +1,
  vertex.label     = NA,
  edge.arrow.size = .25,
  vertex.color     = col_hispanic
)
```

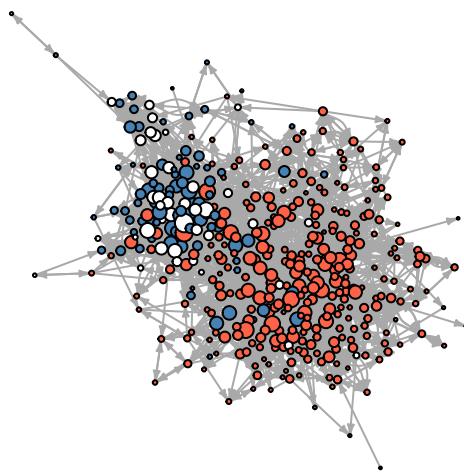


Figure 5.3: Friends network in time 1 for school 111. The graph excludes isolates.

Now that's better! An interesting pattern that shows up is that individuals seem to cluster by whether they are Hispanic or not.

We can write this as a function to avoid copying and pasting the code n times (supposing that we want to create a plot similar to this n times). We do the latter in the following subsection.

5.4.2 Multiple plots

When you are repeating yourself repeatedly, it is a good idea to write down a sequence of commands as a function. In this case, since we will be running the same type of plot for all schools/waves, we write a function in which the only things that change are: (a) the school id, and (b) the color of the nodes.

```
myplot <- function(
  net,
  schoolid,
  mindgr = 1,
  vcol   = "tomato",
  ...) {

  # Creating a subgraph
  subnet <- induced_subgraph(
    net,
    which(degree(net, mode = "all") >= mindgr & V(net)$school == schoolid)
  )

  # Fancy graph
  set.seed(1)
  plot(
    subnet,
    vertex.size      = degree(subnet)/5,
    vertex.label     = NA,
    edge.arrow.size = .25,
    vertex.color     = vcol,
    ...
  )
}
```

The function definition:

1. The `myplot <- function([arguments]) {[body of the function]}` tells R that we are going to create a function called `myplot`.
2. We declare four specific arguments: `net`, `schoolid`, `mindgr`, and `vcol`. These are an igraph object, the school id, the minimum degree that vertices must have to be included in the figure, and the color of the vertices. Observe that, compared to other programming languages, R does not require declaring the data types.
3. The ellipsis object, `....`, is an especial object in R that allows us to pass other arguments without specifying which. If you take a look at the `plot` bit in the function body, you will see that we also added `....`. We use the ellipsis to pass extra arguments (different from the ones that we explicitly defined) directly to `plot`. In practice, this implies that we can,

for example, set the argument `edge.arrow.size` when calling `myplot`, even though we did not include it in the function definition! (See `?dotsMethods` in R for more details).

In the following lines of code, using our new function, we will plot each schools' network in the same plotting device (window) with the help of the `par` function, and add legend with the legend:

```
# Plotting all together
oldpar <- par(no.readonly = TRUE)
par(mfrow = c(2, 3), mai = rep(0, 4), oma= c(1, 0, 0, 0))
myplot(ig_year1, 111, vcol = "tomato")
myplot(ig_year1, 112, vcol = "steelblue")
myplot(ig_year1, 113, vcol = "black")
myplot(ig_year1, 114, vcol = "gold")
myplot(ig_year1, 115, vcol = "white")
par(oldpar)

# A fancy legend
legend(
  "bottomright",
  legend = c(111, 112, 113, 114, 115),
  pt.bg  = c("tomato", "steelblue", "black", "gold", "white"),
  pch    = 21,
  cex    = 1,
  bty    = "n",
  title  = "School"
)
```

So what happened here?

- `oldpar <- par(no.readonly = TRUE)` This line stores the current parameters for plotting. Since we are going to be changing them, we better make sure we are able to go back!.
- `par(mfrow = c(2, 3), mai = rep(0, 4), oma=rep(0, 4))` Here we are setting various things at the same time. `mfrow` specifies how many *figures* will be drawn, and in what order. In particular, we are asking the plotting device to make room for $2 \times 3 = 6$ figures organized in two rows and three columns drawn by row.

`mai` specifies the size of the margins in inches, setting all margins equal to zero (which is what we are doing now) gives more space to the graph. The same is true for `oma`. See `?par` for more info.

- `myplot(ig_year1, ...)` This is simply calling our plotting function. The neat part of this is that, since we set `mfrow = c(2, 3)`, R takes care of *distributing* the plots in the device.
- `par(oldpar)` This line allows us to restore the plotting parameters.

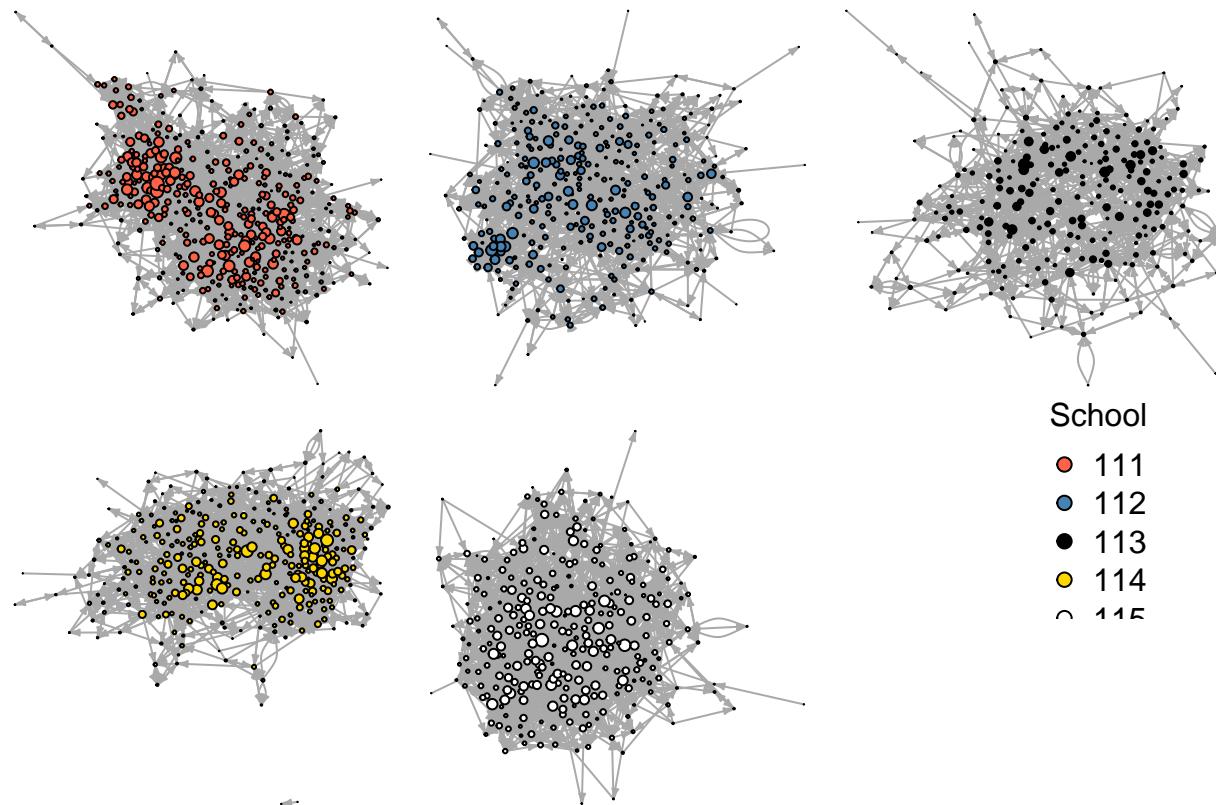


Figure 5.4: All 5 schools in time 1. Again, the graphs exclude isolates.

5.5 Statistical tests

5.5.1 Is nomination number correlated with indegree?

Hypothesis: Individuals that, on average, are among the first nominations of their peers are more popular

```
# Getting all the data in long format
edgelist <- as_long_data_frame(ig_year1) %>%
  as_tibble

# Computing indegree (again) and average nomination number
# Include "On a scale from one to five how close do you feel"
# Also for egocentric friends (A. Friends)
indeg_nom_cor <- group_by(edgelist, to, to_name, to_school) %>%
  summarise(
    indeg = length(nnom),
    nom_avg = 1/mean(nnom)
  ) %>%
  rename(
    school = to_school
  )
```

```
)  
  
## `summarise()` has grouped output by 'to', 'to_name'. You can override using the  
## `.`groups` argument.  
indeg_nom_cor  
  
## # A tibble: 1,561 x 5  
## # Groups:   to, to_name [1,561]  
##       to to_name school indeg nom_avg  
##   <dbl> <chr>    <int>  <int>   <dbl>  
## 1     2 1110002    111    22   0.222  
## 2     3 1110007    111     7   0.175  
## 3     4 1110013    111     6   0.171  
## 4     5 1110014    111    19   0.134  
## 5     6 1110015    111     3   0.15  
## 6     7 1110020    111     6   0.154  
## 7     9 1110025    111     6   0.214  
## 8    10 1110027    111    13   0.220  
## 9    11 1110029    111    14   0.131  
## 10   12 1110030    111     6   0.222  
## # ... with 1,551 more rows  
  
# Using pearson's correlation  
with(indeg_nom_cor, cor.test(indeg, nom_avg))  
  
##  
## Pearson's product-moment correlation  
##  
## data: indeg and nom_avg  
## t = -12.254, df = 1559, p-value < 2.2e-16  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## -0.3409964 -0.2504653  
## sample estimates:  
##            cor  
## -0.2963965  
save.image("03.rda")
```

Chapter 6

Exponential Random Graph Models

I strongly suggest reading the vignette included in the `ergm` R package

```
vignette("ergm", package="ergm")
```

The purpose of ERGMs, in a nutshell, is to describe parsimoniously the local selection forces that shape the global structure of a network. To this end, a network dataset, like those depicted in Figure 1, may be considered as the response in a regression model, where the predictors are things like “propensity for individuals of the same sex to form partnerships” or “propensity for individuals to form triangles of partnerships”. In Figure 1(b), for example, it is evident that the individual nodes appear to cluster in groups of the same numerical labels (which turn out to be students’ grades, 7 through 12); thus, an ERGM can help us quantify the strength of this intra-group effect.

— (David R. Hunter et al. 2008)

In a nutshell, we use ERGMs as a parametric interpretation of the distribution of \mathbf{Y} , which takes the canonical form:

$$\Pr(\mathbf{Y} = \mathbf{y} | \theta, \mathcal{V}) = \frac{\exp\{\theta^T \mathbf{g}(\mathbf{y})\}}{\kappa(\theta, \mathcal{V})}, \quad \mathbf{y} \in \mathcal{Y}$$

Where $\theta \in \Omega \subset \mathbb{R}^q$ is the vector of model coefficients and $\mathbf{g}(\mathbf{y})$ is a q -vector of statistics based on the adjacency matrix \mathbf{y} .

Model (6) may be expanded by replacing $\mathbf{g}(\mathbf{y})$ with $\mathbf{g}(\mathbf{y}, \mathbf{X})$ to allow for additional covariate information \mathbf{X} about the network. The denominator,

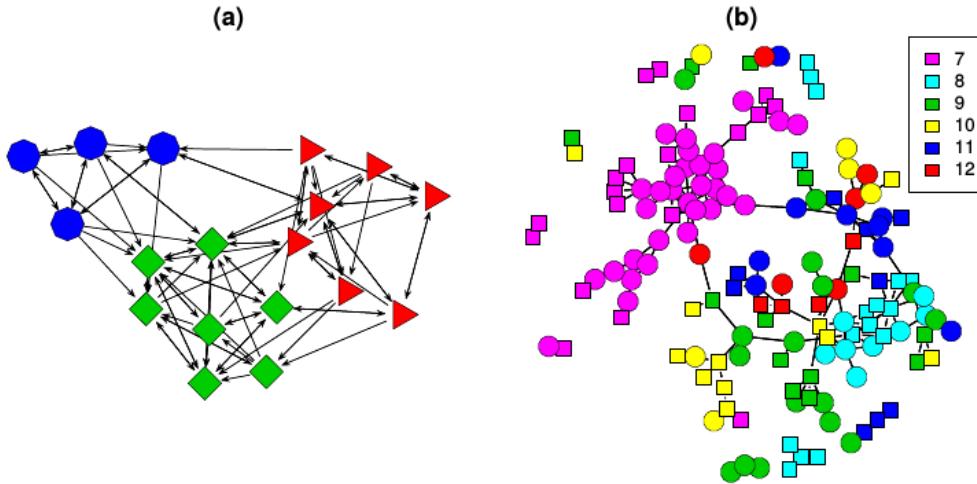


Figure 1: The (a) `samplike` and (b) `faux.mesa.high` networks described in Section 2. The values of nodal covariates may be indicated using various colors, shapes, and labels of nodes.

Figure 6.1: Source: Hunter et al. (2008)

$$\kappa(\theta, \mathcal{Y}) = \sum_{\mathbf{y} \in \mathcal{Y}} \exp \{ \theta^T \mathbf{g}(\mathbf{y}) \}$$

, is the normalizing factor that ensures that equation (6) is a legitimate probability distribution. Even after fixing \mathcal{Y} to be all the networks that have size n , the size of \mathcal{Y} makes this type of statistical model hard to estimate as there are $N = 2^{n(n-1)}$ possible networks! (David R. Hunter et al. 2008)

Recent developments include new forms of dependency structures to take into account more general neighborhood effects. These models relax the one-step Markovian dependence assumptions, allowing investigation of longer-range configurations, such as longer paths in the network or larger cycles (Pattison and Robins 2002). Models for bipartite (Faust and Skvoretz 1999) and tripartite (Mische and Robins 2000) network structures have been developed. (David R. Hunter et al. 2008, 9)

6.1 A naïve example

In the simplest case, ERGMs equate a logistic regression. By simple, I mean cases in which there are no Markovian terms—motifs involving more than one edge—for example, the Bernoulli graph. In the Bernoulli graph, ties are independent of each other, so the presence/absence of a tie between nodes i and j won't affect the presence/absence of a tie between nodes k and l .

Let's fit an ERGM using the `sampson` dataset included in the `ergm` package.

```
library(ergm)
data("sampson")
samplike
## Network attributes:
##   vertices = 18
##   directed = TRUE
##   hyper = FALSE
##   loops = FALSE
##   multiple = FALSE
##   total edges= 88
##   missing edges= 0
##   non-missing edges= 88
##
## Vertex attribute names:
##   cloisterville group vertex.names
##
## Edge attribute names:
##   nominations
```

Using `ergm` to fit a Bernoulli graph requires using the `edges` term, which counts how many ties are in the graph:

```
ergm_fit <- ergm(samplike ~ edges)
## Starting maximum pseudolikelihood estimation (MPLE):
## Obtaining the responsible dyads.
## Evaluating the predictor and response matrix.
## Maximizing the pseudolikelihood.
## Finished MPLE.
## Evaluating log-likelihood at the estimate.
```

Since this is equivalent to a logistic regression, we can use the `glm` function to fit the same model. First, we need to prepare the data so we can pass it to `glm`:

We can now fit the GLM model:

```
glm_fit <- glm(y~1, family=binomial("logit"))
```

The coefficients of both ERGM and GLM should match:

```
glm_fit
##
## Call: glm(formula = y ~ 1, family = binomial("logit"))
##
## Coefficients:
## (Intercept)
##      -0.9072
##
## Degrees of Freedom: 305 Total (i.e. Null); 305 Residual
## Null Deviance:      367.2
## Residual Deviance: 367.2    AIC: 369.2
ergm_fit
##
## Call:
## ergm(formula = samplike ~ edges)
##
## Maximum Likelihood Coefficients:
##   edges
## -0.9072
```

Furthermore, in the case of the Bernoulli graph, we can get the estimate using the Logit function:

```
pr <- mean(y)
# Logit function:
# Alternatively we could have used log(pr) - log(1-pr)
qlogis(pr)
## [1] -0.9071582
```

Again, the same result. The Bernoulli graph is not the only ERGM model that can be fitted using a Logistic regression. Moreover, if all the terms of the model are non-Markov terms, ergm automatically defaults to a Logistic regression.

6.2 Estimation of ERGMs

The ultimate goal is to perform statistical inference on the proposed model. In a *standard* setting, we would be able to use Maximum-Likelihood-Estimation (MLE), which consists of finding the model parameters θ that, given the observed data, maximize the likelihood of the model. For the latter, we generally use Newton's method. Newton's method requires been able to compute the log-likelihood of the model, which in ERGMs can be challenging.

For ERGMs, since part of the likelihood involves a normalizing constant that is a function of all possible networks, this is not as straightforward as in the regular setting. Because of it, most estimation methods rely on simulations.

In `statnet`, the default estimation method is based on a method proposed by (Geyer and Thompson 1992), Markov-Chain MLE, which uses Markov-Chain Monte Carlo for simulating networks and a modified version of the Newton-Raphson algorithm to estimate the parameters.

The idea of MC-MLE for this family of statistical models is that we can approximate the expectation of normalizing constant ratios using the law of large numbers. In particular, the following:

$$\begin{aligned}
 \frac{\kappa(\theta, \mathcal{Y})}{\kappa(\theta_0, \mathcal{Y})} &= \frac{\sum_{\mathbf{y} \in \mathcal{Y}} \exp \{ \theta^T \mathbf{g}(\mathbf{y}) \}}{\sum_{\mathbf{y} \in \mathcal{Y}} \exp \{ \theta_0^T \mathbf{g}(\mathbf{y}) \}} \\
 &= \sum_{\mathbf{y} \in \mathcal{Y}} \left(\frac{1}{\sum_{\mathbf{y} \in \mathcal{Y}} \exp \{ \theta_0^T \mathbf{g}(\mathbf{y}) \}} \times \exp \{ \theta^T \mathbf{g}(\mathbf{y}) \} \right) \\
 &= \sum_{\mathbf{y} \in \mathcal{Y}} \left(\frac{\exp \{ \theta_0^T \mathbf{g}(\mathbf{y}) \}}{\sum_{\mathbf{y} \in \mathcal{Y}} \exp \{ \theta_0^T \mathbf{g}(\mathbf{y}) \}} \times \exp \{ (\theta - \theta_0)^T \mathbf{g}(\mathbf{y}) \} \right) \\
 &= \sum_{\mathbf{y} \in \mathcal{Y}} (\Pr(Y = \mathbf{y} | \mathcal{Y}, \theta_0) \times \exp \{ (\theta - \theta_0)^T \mathbf{g}(\mathbf{y}) \}) \\
 &= E_{\theta_0} (\exp \{ (\theta - \theta_0)^T \mathbf{g}(\mathbf{y}) \})
 \end{aligned}$$

In particular, the MC-MLE algorithm uses this fact to maximize the ratio of log-likelihoods. The objective function itself can be approximated by simulating m networks from the distribution with parameter θ_0 :

$$l(\theta) - l(\theta_0) \approx (\theta - \theta_0)^T \mathbf{g}(\mathbf{y}_{obs}) - \log \left[\frac{1}{m} \sum_{i=1}^m \exp \{ (\theta - \theta_0)^T \} \mathbf{g}(\mathbf{Y}_i) \right]$$

For more details, see (David R. Hunter et al. 2008). A sketch of the algorithm follows:

1. Initialize the algorithm with an initial guess of θ , call it $\theta^{(t)}$ (must be a rather OK guess)
2. While (no convergence) do:
 - a. Using $\theta^{(t)}$, simulate M networks by means of small changes in the \mathbf{Y}_{obs} (the observed network). This part is done by using an importance-sampling method which weights each proposed network by its likelihood conditional on $\theta^{(t)}$
 - b. With the networks simulated, we can do the Newton step to update the parameter $\theta^{(t)}$ (this is the iteration part in the `ergm` package): $\theta^{(t)} \rightarrow \theta^{(t+1)}$.

- c. If convergence has been reached (which usually means that $\theta^{(t)}$ and $\theta^{(t+1)}$ are not very different), then stop; otherwise, go to step a.

For more details see (Lusher, Koskinen, and Robins 2012; Admiraal and Handcock 2006; T. A. Snijders 2002; Wang et al. 2009) provides details on the algorithm used by PNet (which is the same as the one used in RSiena). (Lusher, Koskinen, and Robins 2012) provides a short discussion on the differences between ergm and PNet.

6.3 The ergm package

The ergm R package (Handcock et al. 2017)

From the previous section:¹

```
library(igraph)
library(magrittr)
library(dplyr)

load("03.rda")
```

In this section, we will use the ergm package (from the statnet suit of packages (Handcock et al. 2016),) and the intergraph (Bojanowski 2015) package. The latter provides functions to go back and forth between igraph and network objects from the igraph and network packages respectively²

```
library(ergm)
library(intergraph)
```

As a rather important side note, the order in which R packages are loaded matters. Why is this important to mention now? Well, it turns out that at least a couple of functions in the network package have the same name as some functions in the igraph package. When the ergm package is loaded, since it depends on network, it will load the network package first, which will *mask* some functions in igraph. This becomes evident once you load ergm after loading igraph:

The following objects are masked from ‘package:igraph’:

```
add.edges, add.vertices, %c%, delete.edges, delete.vertices, get.edge.attribute, get.edg
get.vertex.attribute, is.bipartite, is.directed, list.edge.attributes, list.vertex.attri
set.edge.attribute, set.vertex.attribute
```

What are the implications of this? If you call the function list.edge.attributes for an object of class igraph R will return an error as the first function that matches that name comes from the network package! To avoid this you can use the double colon notation:

¹You can download the 03.rda file from this link.

²Yes, the classes have the same name as the packages.

```
igraph::list.edge.attributes(my_igraph_object)
network::list.edge.attributes(my_network_object)
```

Anyway... Using the `asNetwork` function, we can coerce the `igraph` object into a network object so we can use it with the `ergm` function:

```
# Creating the new network
network_111 <- intergraph::asNetwork(ig_year1_111)

# Running a simple ergm (only fitting edge count)
ergm(network_111 ~ edges)
## [1] "Warning: This network contains loops"
## Starting maximum pseudolikelihood estimation (MPLE):
## Evaluating the predictor and response matrix.
## Maximizing the pseudolikelihood.
## Finished MPLE.
## Stopping at the initial estimate.
## Evaluating log-likelihood at the estimate.
##
## Call:
## ergm(formula = network_111 ~ edges)
##
## Maximum Likelihood Coefficients:
##   edges
## -4.734
```

So what happened here? We got a warning. It turns out that our network has loops (didn't think about it before!). Let's take a look at that with the `which_loop` function

```
E(ig_year1_111)[which_loop(ig_year1_111)]
## + 1/2638 edge from 837e5f4 (vertex names):
## [1] 1110111->1110111
```

We can get rid of these using the `igraph:: - .igraph`. Let's remove the isolates using the same operator

```
# Creating the new network
network_111 <- ig_year1_111

# Removing loops
network_111 <- network_111 - E(network_111)[which(which_loop(network_111))]

# Removing isolates
network_111 <- network_111 - which(degree(network_111, mode = "all") == 0)

# Converting the network
```

```
network_111 <- intergraph::asNetwork(network_111)

asNetwork(simplify(ig_year1_111)) ig_year1_111 %>% simplify %>% asNetwork
```

A problem that we have with this data is the fact that some vertices have missing values in the variables `hispanic`, `female1`, and `eversmk1`. For now, we will proceed by imputing values based on the averages:

```
for (v in c("hispanic", "female1", "eversmk1")) {
  tmpv <- network_111 %v% v
  tmpv[is.na(tmpv)] <- mean(tmpv, na.rm = TRUE) > .5
  network_111 %v% v <- tmpv
}
```

6.4 Running ERGMs

Proposed workflow:

1. Estimate the simplest model, adding one variable at a time.
2. After each estimation, run the `mcmc.diagnostics` function to see how good (or bad) behaved the chains are.
3. Run the `gof` function and verify how good the model matches the network's structural statistics.

What to use:

1. `control.ergm`: Maximum number of iterations, seed for Pseudo-RNG, how many cores
2. `ergm.constraints`: Where to sample the network from. Gives stability and (in some cases) faster convergence as by constraining the model you are reducing the sample size.

Here is an example of a couple of models that we could compare³

```
ans0 <- ergm(
  network_111 ~
    edges +
    nodematch("hispanic") +
    nodematch("female1") +
    nodematch("eversmk1") +
    mutual
  ,
  constraints = ~bd(maxout = 19),
  control = control.ergm(
    seed      = 1,
```

³Notice that this document may not include the usual messages that the `ergm` command generates during the estimation procedure. This is just to make it more printable-friendly.

```

MCMLE.maxit = 10,
parallel     = 4,
CD.maxit     = 10
)
)

```

So what are we doing here:

1. The model is controlling for:
 - a. edges Number of edges in the network (as opposed to its density)
 - b. nodematch("some-variable-name-here") Includes a term that controls for homophily/heterophily
 - c. mutual Number of mutual connections between $(i, j), (j, i)$. This can be related to, for example, triadic closure.

For more on control parameters, see (Morris, Handcock, and Hunter 2008).

```

ans1 <- ergm(
  network_111 ~
    edges +
    nodematch("hispanic") +
    nodematch("female1") +
    nodematch("eversmk1")

  ,
  constraints = ~bd(maxout = 19),
  control = control.ergm(
    seed       = 1,
    MCMLE.maxit = 10,
    parallel   = 4,
    CD.maxit   = 10
  )
)

```

This example takes longer to compute

```

ans2 <- ergm(
  network_111 ~
    edges +
    nodematch("hispanic") +
    nodematch("female1") +
    nodematch("eversmk1") +
    mutual +
    balance

  ,
  constraints = ~bd(maxout = 19),
)

```

```

control = control.ergm(
  seed      = 1,
  MCMLE.maxit = 10,
  parallel   = 4,
  CD.maxit   = 10
)
)

```

Now, a nice trick to see all regressions in the same table, we can use the `texreg` package (Leifeld 2013) which supports `ergm` outputs!

```

library(texreg)
## Version: 1.39.3
## Date: 2023-11-09
## Author: Philip Leifeld (University of Essex)
##
## Consider submitting praise using the praise or praise_interactive functions.
## Please cite the JSS article in your publications -- see citation("texreg").
##
## Attaching package: 'texreg'
## The following object is masked from 'package:magrittr':
##
##     extract
screenreg(list(ans0, ans1, ans2))
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.
##
## =====
##          Model 1        Model 2        Model 3
## -----
## edges      -5.63 ***     -5.49 ***     -5.60 ***

```

```

##          (0.05)      (0.06)      (0.06)
## nodematch.hispanic 0.22 ***   0.30 ***   0.22 ***
##                      (0.04)      (0.05)      (0.04)
## nodematch.female1  0.87 ***   1.17 ***   0.87 ***
##                      (0.04)      (0.05)      (0.04)
## nodematch.eversmk1 0.33 ***   0.45 ***   0.34 ***
##                      (0.04)      (0.04)      (0.04)
## mutual             4.10 ***   1.75 ***
##                      (0.07)      (0.14)
## balance            0.01 ***
##                      (0.00)
## -----
## AIC                -40017.80  -37511.87  -39989.59
## BIC                -39967.46  -37471.60  -39929.18
## Log Likelihood    20013.90  18759.94  20000.79
## =====
## *** p < 0.001; ** p < 0.01; * p < 0.05

```

Or, if you are using rmarkdown, you can export the results using LaTeX or html, let's try the latter to see how it looks like here:

```

library(texreg)
texreg(list(ans0, ans1, ans2))
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.
## Warning: This object was fit with 'ergm' version 4.1.2 or earlier. Summarizing
## it with version 4.2 or later may return incorrect results or fail.
## Warning in nobs.ergm(object): The number of observed dyads in this network is
## ill-defined due to complex constraints on the sample space. Disable this
## warning with 'options(ergm.loglik.warn_dyads=FALSE)'.

```

6.5 Model Goodness-of-Fit

In raw terms, once each chain has reach stationary distribution, we can say that there are no problems with autocorrelation and that each sample point is iid. This implies that, since we

	Model 1	Model 2	Model 3
edges	-5.63*** (0.05)	-5.49*** (0.06)	-5.60*** (0.06)
nodematch.hispanic	0.22*** (0.04)	0.30*** (0.05)	0.22*** (0.04)
nodematch.female1	0.87*** (0.04)	1.17*** (0.05)	0.87*** (0.04)
nodematch.eversmk1	0.33*** (0.04)	0.45*** (0.04)	0.34*** (0.04)
mutual	4.10*** (0.07)		1.75*** (0.14)
balance			0.01*** (0.00)
AIC	-40017.80	-37511.87	-39989.59
BIC	-39967.46	-37471.60	-39929.18
Log Likelihood	20013.90	18759.94	20000.79

*** $p < 0.001$; ** $p < 0.01$; * $p < 0.05$

Table 6.1: Statistical models

are running the model with more than 1 chain, we can use all the samples (chains) as a single dataset.

Recent changes in the ergm estimation algorithm mean that these plots can no longer be used to ensure that the mean statistics from the model match the observed network statistics. For that functionality, please use the GOF command: gof(object, GOF=~model).

—?ergm::mcmc.diagnostics

Since ans0 is the one model which did best, let's take a look at it's GOF statistics. First, lets see how the MCMC did. For this we can use the mcmc.diagnostics function including in the package. This function is actually a wrapper of a couple of functions from the coda package (Plummer et al. 2006) which is called upon the \$sample object which holds the *centered* statistics from the sampled networks. This last point is important to consider since at first look it can be confusing to look at the \$sample object since it neither matches the observed statistics, nor the coefficients.

When calling the function mcmc.diagnostics(ans0, centered = FALSE), you will see a lot of output including a couple of plots showing the trace and posterior distribution of the *uncentered* statistics (centered = FALSE). In the next code chunks we will reproduce the output from the mcmc.diagnostics function step by step using the coda package. First we need to *uncenter* the sample object:

```
# Getting the centered sample
sample_centered <- ans0$sample

# Getting the observed statistics and turning it into a matrix so we can add it
# to the samples
observed <- summary(ans0$formula)
```

```

observed <- matrix(
  observed,
  nrow = nrow(sample_centered[[1]]),
  ncol = length(observed),
  byrow = TRUE
)

# Now we uncenter the sample
sample_uncentered <- lapply(sample_centered, function(x) {
  x + observed
})

# We have to make it an mcmc.list object
sample_uncentered <- coda::mcmc.list(sample_uncentered)

```

Under the hood:

1. *Empirical means and sd, and quantiles:*

```

summary(sample_uncentered)
##
## Iterations = 1769472:10944512
## Thinning interval = 65536
## Number of chains = 4
## Sample size per chain = 141
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## edges       2485 60.26   2.5372      3.753
## nodematch.hispanic 1838 51.25   2.1578      3.662
## nodematch.female1 1888 52.78   2.2224      3.779
## nodematch.eversmk1 1759 50.82   2.1400      3.072
## mutual        493 23.40   0.9855      1.967
##
## 2. Quantiles for each variable:
##
##           2.5%  25%  50%  75% 97.5%
## edges       2373 2444 2482 2530 2612
## nodematch.hispanic 1736 1803 1839 1872 1947
## nodematch.female1 1791 1851 1885 1923 1993
## nodematch.eversmk1 1662 1725 1758 1794 1858
## mutual        449  476  493  509   537

```

2. Cross correlation:

```
coda::crosscorr(sample_uncentered)
##                                     edges nodematch.hispanic nodematch.female1
## edges                               1.0000000   0.8657369   0.8851587
## nodematch.hispanic  0.8657369      1.0000000   0.7713632
## nodematch.female1  0.8851587      0.7713632      1.0000000
## nodematch.eversmk1 0.8445651      0.7122693   0.7572735
## mutual                            0.7726517   0.6801783   0.7482026
##                                     nodematch.eversmk1    mutual
## edges                               0.8445651  0.7726517
## nodematch.hispanic      0.7122693  0.6801783
## nodematch.female1       0.7572735  0.7482026
## nodematch.eversmk1      1.0000000  0.6873242
## mutual                            0.6873242  1.0000000
```

3. Autocorrelation: For now, we will only look at autocorrelation for chain one. Autocorrelation should be small (in a general MCMC setting). If autocorrelation is high, then it means that your sample is not iid (no Markov property). A way out to solve this is *thinning* the sample.

```
coda::autocorr(sample_uncentered)[[1]]
## , , edges
##
##                                     edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
## Lag 0      1.000000000   0.861920590   0.90235072   0.86215333
## Lag 65536   0.415060923   0.326775063   0.43751588   0.38274418
## Lag 327680  0.063993999   0.002238453   0.09094189   0.05143792
## Lag 655360  0.002497326   -0.105210070   -0.02414091   0.00143358
## Lag 3276800 0.026845190   0.068616366   0.03686125   0.03652383
##                                     mutual
## Lag 0      0.785264416
## Lag 65536   0.428519050
## Lag 327680  0.074020671
## Lag 655360  0.009422505
## Lag 3276800 0.018126669
##
## , , nodematch.hispanic
##
##                                     edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
## Lag 0      0.86192059   1.000000000   0.76137201   0.74623272
## Lag 65536   0.32680263   0.336764054   0.30353156   0.32690588
## Lag 327680  0.05778076   0.004465856   0.07267341   0.03757479
## Lag 655360  0.07704457   0.024226503   0.03252125   0.08420548
## Lag 3276800 -0.02970399   0.021278122   -0.02753467   -0.03018601
```

```

##               mutual
## Lag 0      0.70578514
## Lag 65536  0.35558587
## Lag 327680 0.05282736
## Lag 655360 0.08176601
## Lag 3276800 -0.07743174
##
## , , nodematch.female1
##
##               edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
## Lag 0      0.902350724    0.76137201    1.000000000   0.77769826
## Lag 65536  0.453418914    0.37756721    0.51290498   0.41954866
## Lag 327680 0.055464012   -0.01058737    0.09841770   0.04272154
## Lag 655360 0.009910833   -0.06123858   -0.03186870   0.04679847
## Lag 3276800 0.004163166    0.04057544    0.01548719   -0.01288236
##               mutual
## Lag 0      0.76981085
## Lag 65536  0.46327442
## Lag 327680 0.03629824
## Lag 655360 0.01987496
## Lag 3276800 -0.00949882
##
## , , nodematch.eversmk1
##
##               edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
## Lag 0      0.86215333    0.746232721   0.77769826   1.000000000
## Lag 65536  0.37539678    0.297591397   0.41717478   0.448697559
## Lag 327680 0.02105523   -0.040132752   0.03760486   0.019124328
## Lag 655360 0.04566425    0.003387581   0.04761067   -0.006388743
## Lag 3276800 0.05048735    0.084790008   0.07108989   0.045582057
##               mutual
## Lag 0      0.7053009595
## Lag 65536  0.4020746950
## Lag 327680 0.0183308894
## Lag 655360 0.0840948296
## Lag 3276800 0.0009713556
##
## , , mutual
##
##               edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
## Lag 0      0.78526442    0.70578514    0.769810849   0.70530096
## Lag 65536  0.50645801    0.44741607    0.532817503   0.47751208
## Lag 327680 0.12979152    0.06061696    0.147380566   0.10930214

```

```

## Lag 655360 -0.06393205      -0.13217821      -0.008121728      -0.03814393
## Lag 3276800 -0.01707605      0.03244214      -0.023750630      0.02781638
##               mutual
## Lag 0        1.0000000000
## Lag 65536    0.580271013
## Lag 327680   0.091309576
## Lag 655360   -0.003521212
## Lag 3276800  -0.025558756

```

4. *Geweke Diagnostic*: From the function's help file:

"If the samples are drawn from the stationary distribution of the chain, the two means are equal and Geweke's statistic has an asymptotically standard normal distribution. [...] The Z-score is calculated under the assumption that the two parts of the chain are asymptotically independent, which requires that the sum of frac1 and frac2 be strictly less than 1."

—?coda::geweke.diag

Let's take a look at a single chain:

```

coda::geweke.diag(sample_uncentered)[[1]]
##
## Fraction in 1st window = 0.1
## Fraction in 2nd window = 0.5
##
##           edges nodematch.hispanic nodematch.female1 nodematch.eversmk1
##           -0.7115          -1.7204          -0.1841          0.6952
##           mutual
##           -1.2891

```

5. *(not included) Gelman Diagnostic*: From the function's help file:

Gelman and Rubin (1992) propose a general approach to monitoring convergence of MCMC output in which $m > 1$ parallel chains are run with starting values that are overdispersed relative to the posterior distribution. Convergence is diagnosed when the chains have 'forgotten' their initial values, and the output from all chains is indistinguishable. The gelman.diag diagnostic is applied to a single variable from the chain. It is based a comparison of within-chain and between-chain variances, and is similar to a classical analysis of variance. —?coda::gelman.diag

As a difference from the previous diagnostic statistic, this uses all chains simultaneously:

```

coda::gelman.diag(sample_uncentered)
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## edges            1.03      1.10

```

```

## nodematch.hispanic    1.03    1.10
## nodematch.female1     1.05    1.14
## nodematch.eversmk1    1.04    1.12
## mutual                 1.05    1.14
##
## Multivariate psrf
##
## 1.05

```

As a rule of thumb, values that are in the [.9, 1.1] are good.

One nice feature of the `mcmc.diagnostics` function is the nice trace and posterior distribution plots that it generates. If you have the R package `latticeExtra` (Sarkar and Andrews 2016), the function will override the default plots used by `coda::plot.mcmc` and use `lattice` instead, creating a nicer looking plots. The next code chunk calls the `mcmc.diagnostic` function, but we suppress the rest of the output (see figure ??).

```

# [2022-03-13] This line is failing for what it could be an ergm bug
# mcmc.diagnostics(ans0, center = FALSE) # Suppressing all the output

```

If we called the function `mcmc.diagnostics`, this message appears at the end:

MCMC diagnostics shown here are from the last round of simulation, prior to computation of final parameter estimates. Because the final estimates are refinements of those used for this simulation run, these diagnostics may underestimate model performance. To directly assess the performance of the final model on in-model statistics, please use the GOF command: `gof(ergmFitObject, GOF=~model)`.

—`mcmc.diagnostics(ans0)`

Not that bad (although the `mutual` term could do better)!⁴ First, observe that in the figure we see four different lines; why is that? Since we were running in parallel using four cores, the algorithm ran four chains of the MCMC algorithm. An eyeball test is to see if all the chains moved at about the same place; in such a case, we can start thinking about model convergence from the MCMC perspective.

Once we are sure to have reach convergence on the MCMC algorithm, we can start thinking about how well does our model predicts the observed network's properties. Besides the statistics that define our ERGM, the `gof` function's default behavior show GOF for:

- a. In degree distribution,
- b. Out degree distribution,
- c. Edge-wise shared partners, and
- d. Geodesics

Let's take a look at it

⁴The statnet wiki website as a very nice example of (very) bad and good MCMC diagnostics plots here.

```

# Computing and printing GOF estatistics
ans_gof <- gof(ans0)
ans_gof
##
## Goodness-of-fit for in-degree
##
##          obs min  mean max MC p-value
## idegree0 13   0  1.38  5   0.00
## idegree1 34   2  7.84 15   0.00
## idegree2 37   13 22.13 36   0.00
## idegree3 48   26 40.12 60   0.12
## idegree4 37   44 59.60 74   0.00
## idegree5 47   52 67.86 89   0.00
## idegree6 42   48 65.06 82   0.00
## idegree7 39   38 53.40 69   0.04
## idegree8 35   26 40.30 57   0.46
## idegree9 21   17 26.48 43   0.36
## idegree10 12   9  16.36 24   0.34
## idegree11 19   2  8.73 15   0.00
## idegree12 4    0  4.80 13   1.00
## idegree13 7    0  2.23  6   0.00
## idegree14 6    0  0.99  4   0.00
## idegree15 3    0  0.32  2   0.00
## idegree16 4    0  0.20  2   0.00
## idegree17 3    0  0.09  1   0.00
## idegree18 3    0  0.09  1   0.00
## idegree19 2    0  0.00  0   0.00
## idegree20 1    0  0.02  1   0.04
## idegree22 1    0  0.00  0   0.00
##
## Goodness-of-fit for out-degree
##
##          obs min  mean max MC p-value
## odegree0 4    0  1.33  6   0.14
## odegree1 28   2  7.67 15   0.00
## odegree2 45   13 22.07 33   0.00
## odegree3 50   27 40.96 55   0.16
## odegree4 54   44 59.37 75   0.50
## odegree5 62   40 67.12 92   0.62
## odegree6 40   46 64.49 81   0.00
## odegree7 28   39 54.33 78   0.00
## odegree8 13   28 39.73 53   0.00
## odegree9 16   17 27.48 45   0.00

```

```

## odegree10 20 7 15.74 26      0.22
## odegree11 8 2 9.25 16       0.86
## odegree12 11 1 4.98 12      0.04
## odegree13 13 0 2.12 7       0.00
## odegree14 6 0 0.72 3       0.00
## odegree15 6 0 0.41 3       0.00
## odegree16 7 0 0.17 2       0.00
## odegree17 4 0 0.03 1       0.00
## odegree18 3 0 0.01 1       0.00
## odegree19 0 0 0.02 1       1.00
##
## Goodness-of-fit for edgewise shared partner
##
##          obs  min   mean   max MC p-value
## esp0 1032 1991 2195.84 2289      0
## esp1 755 170  235.05 441       0
## esp2 352 1    14.86  86       0
## esp3 202 0    0.84   14      0
## esp4 79 0    0.08   2       0
## esp5 36 0    0.00   0       0
## esp6 14 0    0.00   0       0
## esp7 4 0    0.00   0       0
## esp8 1 0    0.00   0       0
##
## Goodness-of-fit for minimum geodesic distance
##
##          obs  min   mean   max MC p-value
## 1 2475 2319 2446.67 2571      0.6
## 2 10672 12048 13548.63 14726     0.0
## 3 31134 48016 54997.05 60096     0.0
## 4 50673 78122 80230.34 82613     0.0
## 5 42563 15553 20682.28 27611     0.0
## 6 18719 383  1227.29 2611      0.0
## 7 4808 0    40.46  214      0.0
## 8 822 0    0.87   16      0.0
## 9 100 0    0.00   0       0.0
## 10 7 0    0.00   0       0.0
## Inf 12333 0  1132.41 3324      0.0
##
## Goodness-of-fit for model statistics
##
##          obs  min   mean   max MC p-value
## edges 2475 2319 2446.67 2571      0.60



```

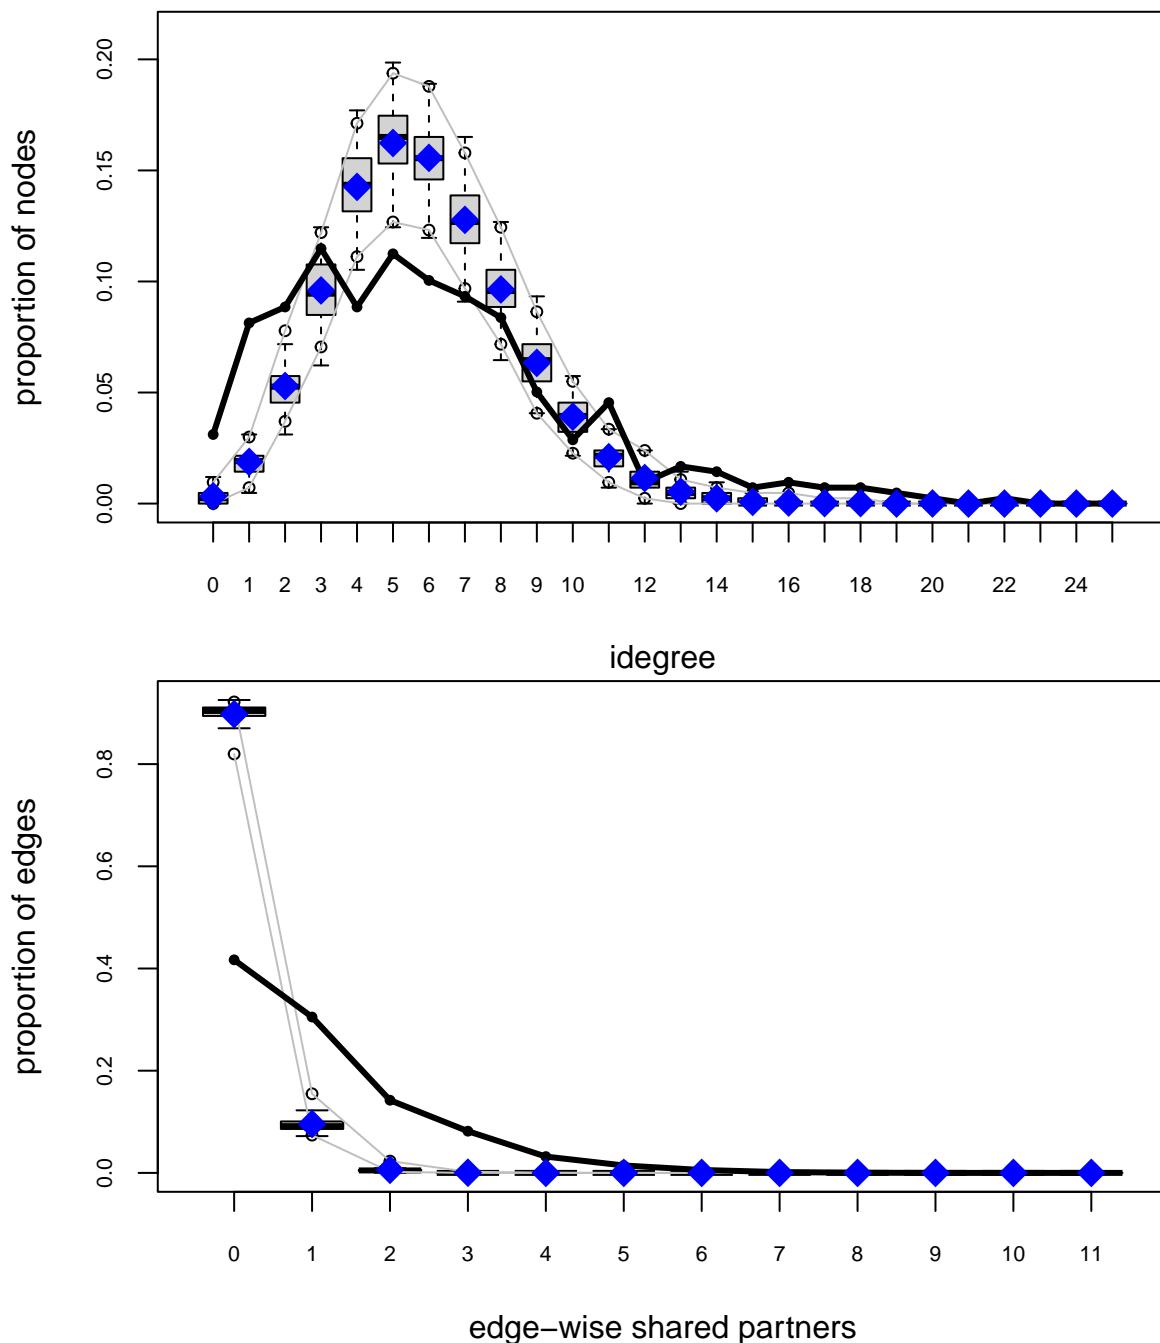
```

## nodematch.hispanic 1832 1710 1807.58 1903      0.56
## nodematch.female1 1879 1760 1852.63 1970      0.60
## nodematch.eversmk1 1755 1614 1724.70 1824      0.56
## mutual           486  441  475.38  514       0.66

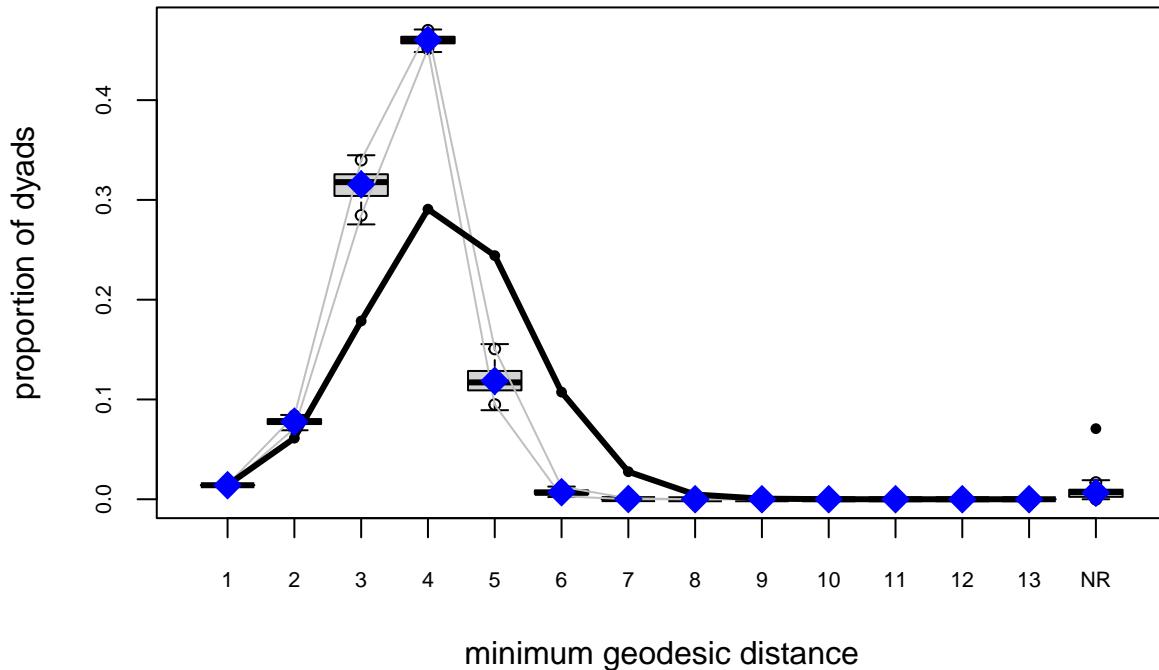
# Plotting GOF statistics
plot(ans_gof)

```





Goodness-of-fit diagnostics



Try the following configuration instead

```
ans0_bis <- ergm(
  network_111 ~
    edges +
    nodematch("hispanic") +
    nodematch("female1") +
    mutual +
    esp(0:3) +
    indegree(0:10)
  ,
  constraints = ~bd(maxout = 19),
  control = control.ergm(
    seed      = 1,
    MCMLE.maxit = 15,
    parallel   = 4,
    CD.maxit   = 15,
    MCMC.samplesize = 2048*4,
    MCMC.burnin = 30000,
    MCMC.interval = 2048*4
  )
)
```

Increase the sample size, so the curves are smoother, longer intervals (thinning), which reduces autocorrelation, and a larger burin. All this together to improve the Gelman test statistic. We also added idegree from 0 to 10, and esp from 0 to 3 to explicitly match those statistics in our model.

```
knitr:::include_graphics("awful-chains.png")
```

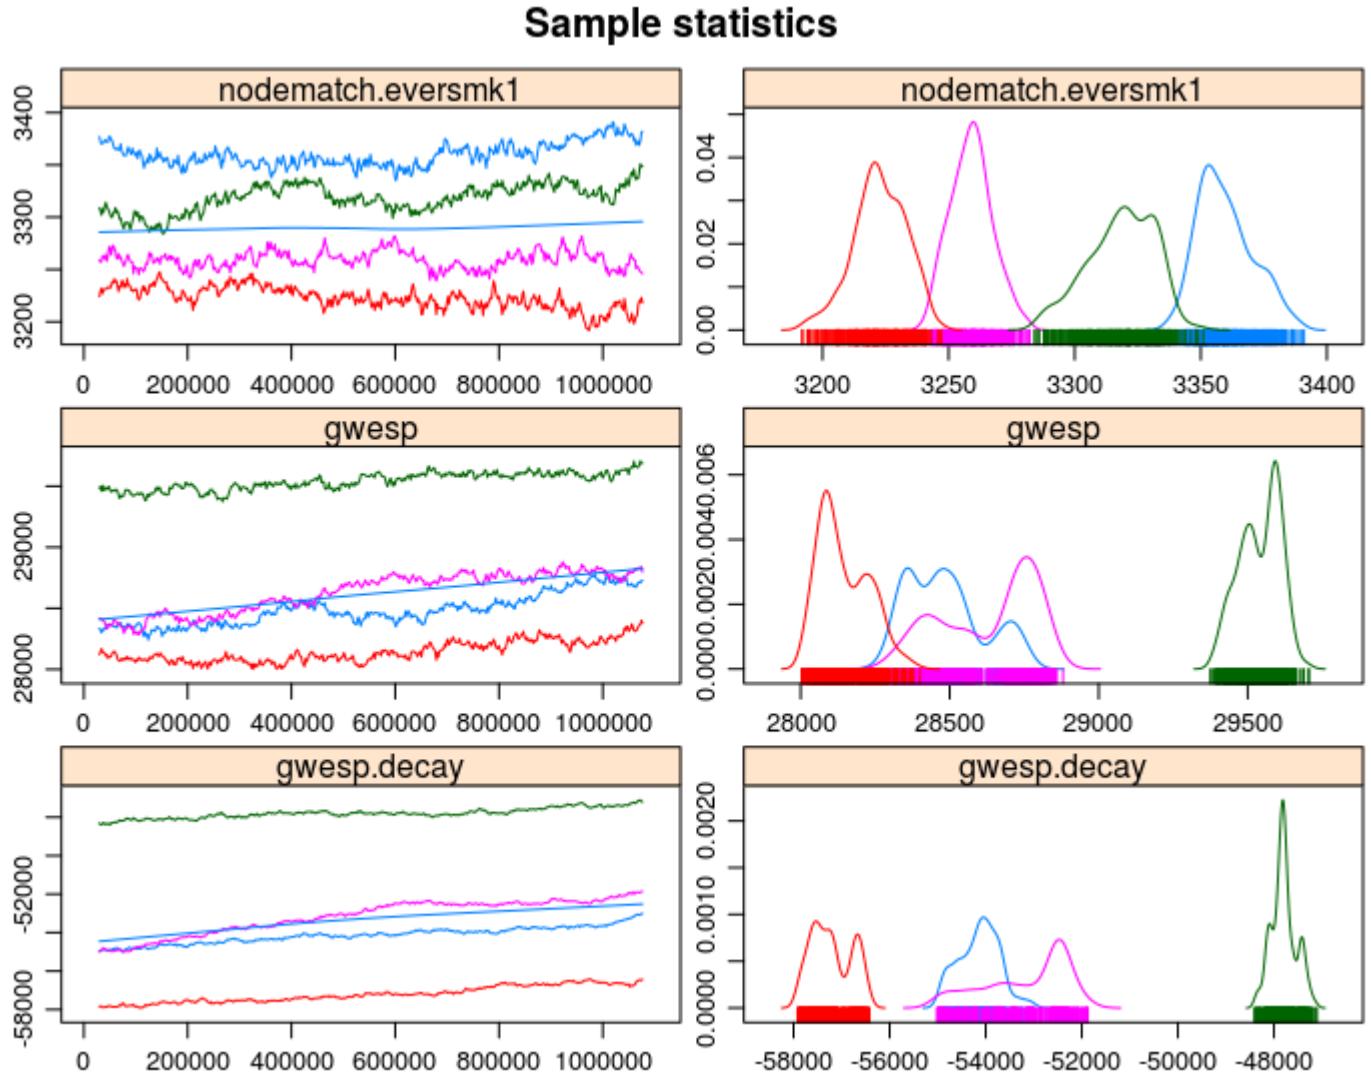


Figure 6.2: An example of a terrible ERGM (no convergence at all). Also, a good example of why running multiple chains can be useful

6.6 More on MCMC convergence

For more on this issue, I recommend reviewing chapter 1 and chapter 6 from the Handbook of MCMC (Brooks et al. 2011). Both chapters are free to download from the book's website.

For GOF take a look at section 6 of ERGM 2016 Sunbelt tutorial, and for a more technical review, you can take a look at (David R. Hunter, Goodreau, and Handcock 2008).

6.7 Mathematical Interpretation

One of the most critical parts of statistical modeling is interpreting the results, if not the most important. In the case of ERGMs, a key aspect is based on change statistics. Suppose that we would like to know how likely the tie y_{ij} is to happen, given the rest of the network. We can compute such probabilities using what literature sometimes describes as the Gibbs-sampler.

In particular, the log-odds of the ij tie occurring conditional on the rest of the network can be written as:

$$\text{logit}(\mathbb{P}(y_{ij} = 1 | y_{-ij})) = \theta^t \Delta \delta(y_{ij} : 0 \rightarrow 1), \quad (6.1)$$

with $\delta(y_{ij} : 0 \rightarrow 1) \equiv s(\mathbf{y})_{ij}^+ - s(\mathbf{y})_{ij}^-$ as the vector of change statistics, in other words, the difference between the sufficient statistics when $y_{ij} = 1$ and its value when $y_{ij} = 0$. To show this, we write the following:

$$\begin{aligned} \mathbb{P}(y_{ij} = 1 | y_{-ij}) &= \frac{\mathbb{P}(y_{ij} = 1, x_{-ij})}{\mathbb{P}(y_{ij} = 1, y_{-ij}) + \mathbb{P}(y_{ij} = 0, y_{-ij})} \\ &= \frac{\exp\{\theta^t s(\mathbf{y})_{ij}^+\}}{\exp\{\theta^t s(\mathbf{y})_{ij}^+\} + \exp\{\theta^t s(\mathbf{y})_{ij}^-\}} \end{aligned}$$

Applying the logit function to the previous equation, we obtain:

$$\begin{aligned} &= \log \left\{ \frac{\exp\{\theta^t s(\mathbf{y})_{ij}^+\}}{\exp\{\theta^t s(\mathbf{y})_{ij}^+\} + \exp\{\theta^t s(\mathbf{y})_{ij}^-\}} \right\} - \log \left\{ \frac{\exp\{\theta^t s(\mathbf{y})_{ij}^-\}}{\exp\{\theta^t s(\mathbf{y})_{ij}^+\} + \exp\{\theta^t s(\mathbf{y})_{ij}^-\}} \right\} \\ &= \log \{\exp\{\theta^t s(\mathbf{y})_{ij}^+\}\} - \log \{\exp\{\theta^t s(\mathbf{y})_{ij}^-\}\} \\ &= \theta^t (s(\mathbf{y})_{ij}^+ - s(\mathbf{y})_{ij}^-) \\ &= \theta^t \Delta \delta(y_{ij} : 0 \rightarrow 1) \end{aligned}$$

Henceforth, the conditional probability of node n gaining function k can be written as:

$$\mathbb{P}(y_{ij} = 1 | y_{-ij}) = \frac{1}{1 + \exp\{-\theta^t \Delta \delta(y_{ij} : 0 \rightarrow 1)\}} \quad (6.2)$$

i.e., a logistic probability.

6.8 Markov independence

The challenge of analyzing networks is their interdependent nature. Nonetheless, in the absence of such interdependence, ERGMs are equivalent to logistic regression. Conceptually, if all the statistics included in the model do not involve two or more dyads, then the model is non-Markovian in the sense of Markov graphs.

Mathematically, to see this, it suffices to show that the ERGM probability can be written as the product of each dyads' probabilities.

$$\mathbb{P}(\mathbf{y}|\theta) = \frac{\exp\{\theta^t s(\mathbf{y})\}}{\sum_y \exp\{\theta^t s(\mathbf{y})\}} = \frac{\prod_{ij} \exp\{\theta^t s(\mathbf{y})_{ij}\}}{\sum_y \exp\{\theta^t s(\mathbf{y})\}}$$

Where $s()_{ij}$ is a function such that $s(\mathbf{y}) = \sum_{ij} s(\mathbf{y})_{ij}$. We now need to deal with the normalizing constant. To see how that can be separated, let's start from the result:

$$\begin{aligned} &= \prod_{ij} (1 + \exp\{\theta^t s(\mathbf{y})_{ij}\}) \\ &= (1 + \exp\{\theta^t s(\mathbf{y})_{11}\})(1 + \exp\{\theta^t s(\mathbf{y})_{12}\}) \dots (1 + \exp\{\theta^t s(\mathbf{y})_{nn}\}) \\ &= 1 + \exp\{\theta^t s(\mathbf{y})_{11}\} + \exp\{\theta^t s(\mathbf{y})_{11}\} \exp\{\theta^t s(\mathbf{y})_{12}\} + \dots + \prod_{ij} \exp\{\theta^t s(\mathbf{y})_{ij}\} \\ &= 1 + \exp\{\theta^t s(\mathbf{y})_{11}\} + \exp\{\theta^t(s(\mathbf{y})_{11} + s(\mathbf{y})_{12})\} + \dots + \prod_{ij} \exp\{\theta^t s(\mathbf{y})_{ij}\} \\ &= \sum_{\mathbf{y} \in \mathcal{Y}} \exp\{\theta^t s(\mathbf{y})\} \end{aligned}$$

Where the last equality follows from the fact that the sum *is* the sum over all possible combinations of networks, starting from $\exp(0) = 1$, up to $\exp(all)$. This way, we can now write:

$$\frac{\prod_{ij} \exp\{\theta^t s(\mathbf{y})_{ij}\}}{\sum_y \exp\{\theta^t s(\mathbf{y})\}} = \prod_{ij} \frac{\exp\{\theta^t s(\mathbf{y})_{ij}\}}{1 + \exp\{\theta^t s(\mathbf{y})_{ij}\}} \quad (6.3)$$

Related to this, block-diagonal ERGMs can be estimated as independent models, one per block. To see more about this, read (SNIJDERS 2010). Likewise, since independence depends–pun intended–on partitioning the objective function, as pointed by Snijders, non-linear functions make the model dependent, e.g., $s(\mathbf{y}) = \sqrt{\sum_{ij} y_{ij}}$, the square root of the edgecount is no longer a bernoulli graph.

Chapter 7

Using constraints in ERGMs

Exponential Random Graph Models [ERGMs] can represent a variety of network classes. We often look at “regular” social networks like students in schools, colleagues in the workplace, or families. Nonetheless, some social networks we study have features that restrict how connections can occur. Typical examples are bi-partite graphs and multilevel networks. There are two classes of vertices in bi-partite networks, and ties can only occur between classes. On the other hand, Multilevel networks may feature multiple classes with inter-class ties somewhat restricted. In both cases, structural constraints exist, meaning that some configurations may not be plausible.

Mathematically, what we are trying to do is, instead of assuming that all network configurations are possible:

$$\{\mathbf{y} \in \mathcal{Y} : y_{ij} = 0, \forall i = j\}$$

we want to go a bit further avoiding loops, namely:

$$\{\mathbf{y} \in \mathcal{Y} : y_{ii} = 0, \forall i; \mathbf{y} \in C\}$$

where C is a constraint, for example, only networks with no triangles. The `ergm` R package has built-in capabilities to deal with some of these cases. Nonetheless, we can specify models with arbitrary structural constraints built into the model. The key is in using offset terms.

7.1 Example 1: Interlocking egos and disconnected alters

Imagine that we have two sets of vertices. The first, group E, are egos part of an egocentric study. The second group, called A, is composed of people mentioned by egos in E but were not surveyed. Assume that individuals in A can only connect to individuals in E; moreover,

individuals in E have no restrictions on connecting. In other words, only two types of ties exist: E-E and A-E. The question is now, how can we enforce such a constraint in an ERGM?

Using offsets, and in particular, setting coefficients to `-Inf` provides an easy way to restrict the support set of ERGMs. For example, if we wanted to constrain the support to include networks with no triangles, we would add the term `offset(triangle)` and use the option `offset.coef = -Inf` to indicate that realizations including triangles are not possible. Using R:

```
ergm(net ~ edges + offset(triangle), offset.coef = -Inf)
```

In this model, a Bernoulli graph, we reduce the sample space to networks with no triangles. In our example, such a statistic should only take non-zero values whenever ties within the A class happen. We can use the `nodematch()` term to do that. Formally

$$\text{NodeMatch}(x) = \sum_{i,j} y_{ij} \mathbf{1}(x_i = x_j)$$

This statistic will sum over all ties in which source (i) and target (j)'s X attribute is equal. One way to make this happen is by creating an auxiliary variable that equals, e.g., 0 for all vertices in A, and a unique value different from zero otherwise. For example, if we had 2 As and three Es, the data would look something like this: $\{0, 0, 1, 2, 3\}$. The following code block creates an empty graph with 50 nodes, 10 of which are in group E (ego).


```
library(ergm, quietly = TRUE)
library(sna, quietly = TRUE)

n <- 50
n_egos <- 10
net <- as.network(matrix(0, ncol = n, nrow = n), directed = TRUE)

# Let's assign the groups
net %v% "is.ego" <- c(rep(TRUE, n_egos), rep(FALSE, n - n_egos))
net %v% "is.ego"

## [1] TRUE FALSE FALSE
## [13] FALSE FALSE
## [25] FALSE FALSE
## [37] FALSE FALSE
## [49] FALSE FALSE

gplot(net, vertex.col = net %v% "is.ego")
```



To create the auxiliary variable, we will use the following function:

```
# Function that creates an aux variable for the ergm model
make_aux_var <- function(my_net, is_ego_dummy) {

  n_vertex <- length(my_net %v% is_ego_dummy)
  n_ego_   <- sum(my_net %v% is_ego_dummy)

  # Creating an auxiliary variable to identify the non-informant non-informant ties
  my_net %v% "aux_var" <- ifelse(
    !my_net %v% is_ego_dummy, 0, 1:(n_vertex - n_ego_)

  )

  my_net
}
```

Calling the function in our data results in the following:

```
net <- make_aux_var(net, "is.ego")

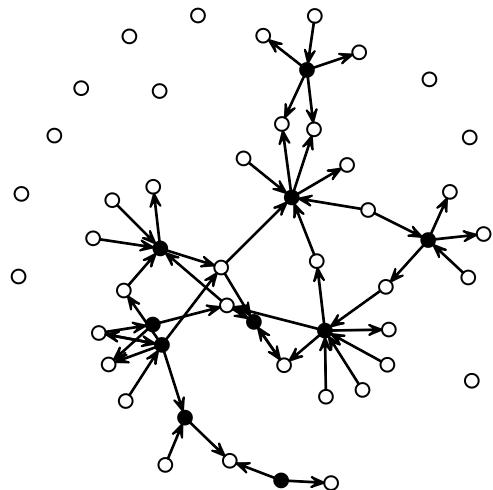
# Taking a look over the first 15 rows of data
cbind(
  Is_Ego = net %v% "is.ego",
  Aux     = net %v% "aux_var"
) |> head(n = 15)

##      Is_Ego Aux
## [1,]      1   1
## [2,]      1   2
## [3,]      1   3
## [4,]      1   4
## [5,]      1   5
```

```
## [6,]    1   6
## [7,]    1   7
## [8,]    1   8
## [9,]    1   9
## [10,]   1  10
## [11,]   0   0
## [12,]   0   0
## [13,]   0   0
## [14,]   0   0
## [15,]   0   0
```

We can now use this data to simulate a network in which ties between A-class vertices are not possible:

```
set.seed(2828)
net_sim <- simulate(net ~ edges + nodematch("aux_var"), coef = c(-3.0, -Inf))
gplot(net_sim, vertex.col = net_sim %v% "is.ego")
```



As you can see, this network has only ties of the type E-E and A-E. We can double-check by (i) looking at the counts and (ii) visualizing each induced-subgraph separately:

```
summary(net_sim ~ edges + nodematch("aux_var"))

##               edges nodematch.aux_var
##             49                  0

net_of alters <- get.inducedSubgraph(
  net_sim, which((net_sim %v% "aux_var") == 0)
)

net_of egos <- get.inducedSubgraph(
  net_sim, which((net_sim %v% "aux_var") != 0)
)
```

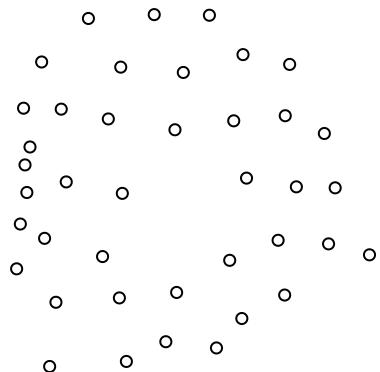
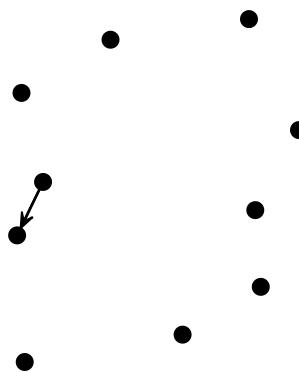
```
# Counts
summary(net_of alters ~ edges + nodematch("aux_var"))

##          edges nodematch.aux_var
##                 0                  0

summary(net_of_egos ~ edges + nodematch("aux_var"))

##          edges nodematch.aux_var
##                 1                  0

# Figures
op <- par(mfcol = c(1, 2))
gplot(net_of alters, vertex.col = net_of alters %v% "is.ego", main = "A")
gplot(net_of_egos, vertex.col = net_of_egos %v% "is.ego", main = "E")
```

A**E**

```
par(op)
```

Now, to fit an ERGM with this constraint, we simply need to make use of the offset terms. Here is an example:

```
ans <- ergm(
  net_sim ~ edges + offset(nodematch("aux_var")), # The model (notice the offset)
  offset.coef = -Inf                            # The offset coefficient
)
## Starting maximum pseudolikelihood estimation (MPLE):
## Obtaining the responsible dyads.
## Evaluating the predictor and response matrix.
## Maximizing the pseudolikelihood.
## Finished MPLE.
## Evaluating log-likelihood at the estimate.
summary(ans)
```

```

## Call:
## ergm(formula = net_sim ~ edges + offset(nodematch("aux_var")),
##       offset.coef = -Inf)
##
## Maximum Likelihood Results:
##
##                               Estimate Std. Error MCMC % z value Pr(>|z|)
## edges                  -2.843     0.147      0   -19.34  <1e-04 ***
## offset(nodematch.aux_var) -Inf      0.000      0    -Inf  <1e-04 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Null Deviance: 1233.8 on 890 degrees of freedom
## Residual Deviance: 379.4 on 888 degrees of freedom
##
## AIC: 381.4 BIC: 386.2 (Smaller is better. MC Std. Err. = 0)
##
## The following terms are fixed by offset and are not estimated:
## offset(nodematch.aux_var)


```

This ERGM model—which by the way only featured dyadic-independent terms, and thus can be reduced to a logistic regression—restricts the support by excluding all networks in which ties within the class A exists. To finalize, let's look at a few simulations based on this model:

```

set.seed(1323)
op <- par(mfcol = c(2,2), mar = rep(1, 4))
for (i in 1:4) {
  gplot(simulate(ans), vertex.col = net %v% "is.ego", vertex.cex = 2)
  box()
}

```



par(op)

All networks with no ties between A nodes.

7.2 Example 2: Bi-partite networks

In the case of bipartite networks (sometimes called affiliation networks,) we can use `ergm`'s terms for bipartite graphs to corroborate what we discussed here. For example, the two-star term. Let's start simulating a bipartite network using the edges and two-star parameters. Since the k-star term is usually complex to fit (tends to generate degenerate models,) we will take advantage of the `Log()` transformation function in the `ergm` package to smooth the term.¹

The bipartite network that we will be simulating will have 100 actors and 50 entities. Actors, which we will map to the first level of the `ergm` terms, this is, `b1star` `b1nodematch`, etc. will send ties to the entities, the second level of the bipartite ERGM. To create a bipartite network, we will create an empty matrix of size `nactors` x `nentities`; thus, actors are represented by rows and entities by columns.

¹After writing this example, it became apparent the use of the `Log()` transformation function may not be ideal. Since many terms used in ERGMs can be zero, e.g., triangles, the term `Log(~ ostar(2))` is undefined when `ostar(2) = 0`. In practice, the `ERGM` package sets a lower limit for the log of 0, so, instead of having `Log(0) ~ -Inf`, they set it to be a really large negative number. This causes all sorts of issues to the estimates; in our example, an overestimation of the population parameter and a positive log-likelihood. Therefore, I wouldn't recommend using this transformation too often.

```
# Parameters for the simulation
nactors <- 100
nentities <- floor(nactors/2)
n <- nactors + nentities

# Creating an empty bipartite network (baseline)
net_b <- network(
  matrix(0, nrow = nactors, ncol = nentities), bipartite = TRUE
)


# Simulating the bipartite ERGM,
net_b <- simulate(net_b ~ edges + Log(~b1star(2)), coef = c(-3, 1.5), seed = 55)
```

Let's see what we got here:

```
summary(net_b ~ edges + Log(~b1star(2)))
```

```
##      edges Log~b1star2
## 245.000000    5.746203
```

```
netplot::nplot(net_b, vertex.col = (1:n <= nactors) + 1)
```



Notice that the first `nactors` vertices in the network are the actors, and the remaining are the entities. Now, although the `ergm` package features bipartite network terms, we can still fit a

bipartite ERGM without explicitly declaring the graph as such. In such case, the `b1star(2)` term of a bipartite network is equivalent to an `ostar(2)` in a directed graph. Likewise, `b2star(2)` in a bipartite graph matches the `istar(2)` term in a directed graph. This information will be relevant when fitting the ERGM. Let's transform the bipartite network into a directed graph. The following code block does so:

```
# Identifying the edges
net_not_b <- which(as.matrix(net_b) != 0, arr.ind = TRUE)

# We need to offset the endpoint of the ties by nactors
# so that the ids go from 1 through (nactors + nentities)
net_not_b[,2] <- net_not_b[,2] + nactors

# The resulting graph is a directed network
net_not_b <- network(net_not_b, directed = TRUE)
```

Now we are almost done. As before, we need to use node-level covariates to put the constraints in our model. For this ERGM to reflect an ERGM on a bipartite network, we need two constraints:

1. Only ties from actors to entities are allowed, and
2. entities can only receive ties.

The corresponding offset terms for this model are: `nodematch("is.actor") ~ -Inf`, and `nodeocov("isnot.actor") ~ -Inf`. Mathematically:

$$\text{NodeMatch}(x = \text{"is.actor"}) = \sum_{i < j} y_{ij} \mathbb{1}(x_i = x_j)$$

$$\text{NodeOCov}(x = \text{"isnot.actor"}) = \sum_i x_i \times \sum_{j < i} y_{ij}$$

In other words, we are setting that ties between nodes of the same class are forbidden, and outgoing ties are forbidden for entities. Let's create the vertex attributes needed to use the aforementioned terms:

```
net_not_b %v% "is.actor" <- as.integer(1:n <= nactors)
net_not_b %v% "isnot.actor" <- as.integer(1:n > nactors)
```

Finally, to make sure we have done all well, let's look how both networks—bipartite and unimodal—look side by side:

```
# First, let's get the layout
fig <- netplot::nplot(net_b, vertex.col = (1:n <= nactors) + 1)
gridExtra::grid.arrange(
  fig,
  netplot::nplot(
    net_not_b, vertex.col = (1:n <= nactors) + 1,
```

```

    layout = fig$.layout
      ),
  ncol = 2, nrow = 1
)

# Looking at the counts
summary(net_b ~ edges + b1star(2) + b2star(2))

##   edges b1star2 b2star2
##     245     313     645

summary(net_not_b ~ edges + ostar(2) + istar(2))

##   edges ostar2 istar2
##     245     313     645

```

With the two networks matching, we can now fit the ERGMs with and without offset terms and compare the results between the two models:

```

# ERGM with a bipartite graph
res_b <- ergm(
  # Main formula
  net_b ~ edges + Log(~b1star(2)),

  # Control parameters
  control = control.ergm(seed = 1)
)

# ERGM with a digraph with constraints
res_not_b <- ergm(
  # Main formula
  net_not_b ~ edges + Log(~ostar(2)) +

  # Offset terms
  offset(nodematch("is.actor")) + offset(nodecov("isnot.actor")),
)

```

```

offset.coef = c(-Inf, -Inf),

# Control parameters
control = control.ergm(seed = 1)
)

```

Here are the estimates (using the `texreg` R package for a prettier output):

```

texreg::screenreg(list(Bipartite = res_b, Directed = res_not_b))

##
## =====
##               Bipartite     Directed
## -----
## edges           -3.14 ***      -3.11 ***
##                   (0.15)        (0.14)
## Log~b1star2    21.89
##                   (17.13)
## Log~ostar2      19.66
##                   (16.75)
## offset(nodematch.is.actor)      -Inf
## 
## offset(nodeocov.isnot.actor)      -Inf
## 
## -----
## AIC            1958.00      -2134192392498170112.00
## BIC            1971.03      -2134192392498170112.00
## Log Likelihood -977.00      1067096196249085056.00
## =====
## *** p < 0.001; ** p < 0.01; * p < 0.05

```

As expected, both models yield the “same” estimate. The minor differences observed between the models are how the `ergm` package performs the sampling. In particular, in the bipartite case, `ergm` has special routines for making the sampling more efficient, having a higher acceptance rate than that of the model in which the bipartite graph was not explicitly declared. We can tell this by inspecting rejection rates:

```

data.frame(
  Bipartite = coda::rejectionRate(res_b$sample[[1]]) * 100,
  Directed = coda::rejectionRate(res_not_b$sample[[1]][, -c(3,4)]) * 100
) |> knitr::kable(digits = 2, caption = "Rejection rate (percent)")

```

The ERGM fitted with the offset terms has a much higher rejection rate than that of the ERGM fitted with the bipartite ERGM.

Finally, the fact that we can fit ERGMs using offset does not mean that we need to use it

Table 7.1: Rejection rate (percent)

	Bipartite	Directed
edges	2.48	3.67
Log~b1star2	1.24	2.04

ALL the time. Unless there is a very good reason to go around ergm's capabilities, I wouldn't recommend fitting bipartite ERGMs as we just did, as the authors of the package have included (MANY) features to make our job easier.

Chapter 8

(Separable) Temporal Exponential Family Random Graph Models

This tutorial is great! https://statnet.org/trac/raw-attachment/wiki/Sunbelt2016/tergm_tutorial.pdf

Chapter 9

Simulating and visualizing networks

In this chapter, we will build and visualize artificial networks using Exponential Random Graph Models [ERGMs.] Together with chapter 3, this will be an extended example of how to read network data and visualize it using some of the available R packages out there.

For this chapter, we will be using the following R packages:

- `ergm`: To simulate and estimate ERGMs.
- `sna`: To visualize networks.
- `igraph`: Also to visualize networks.
- `intergraph`: To convert between `igraph` and `network` objects.
- `netplot`: Again, for visualization.
- `netdiffuseR`: For a single function we use for adjusting vertex size in `igraph`.
- `rgexf`: For building interactive (html) figures.

You can use the following codeblock to install any missing package:

```
# Creating the list to install
pkgs <- c(
  "ergm", "sna", "igraph", "intergraph", "netplot", "netdiffuseR", "rgexf"
)

# Checking if we can load them and install them if not available
for (pkg in pkgs) {
  if (!require(pkg, character.only = TRUE)) {

    # If not present, will install it
    install.packages(pkg, character.only = TRUE)

  # And load it!
}
```

```

library(pkg, character.only = TRUE)

}
}

```

A recorded version is available [here](#).

9.1 Random Graph Models

While there are tons of social network data, we will use an artificial one for this chapter. We do this as it is always helpful to have more examples simulating Random networks. For this chapter, we will classify random graph models for sampling and generating networks into three categories:

1. **Exogenous**: Graphs where the structure is determined by a macro rule, e.g., expected density, degree distribution, or degree-sequence. In these cases, ties are assigned to comply with a macro-property.
2. **Endogenous**: This category includes all Random Graphs generated based on endogenous information, e.g., small-world, scale-free, etc. Here, a tie creation rule gives origin to a macro property, for example, preferential attachment in scale-free networks.
3. **Exponential Random Graph Models**: Overall, since ERGMs compose a family of statistical models, we can always (or almost always) find a model specification that matches the previous categories. Whereas we are thinking about degree sequence, preferential attachment, or a mix of both, ERGMs can be the baseline for any of those models.

The latter, ERGMs, are a generalization that covers all classes. Because of that, we will use ERGMs to generate our artificial network.

9.2 Social Networks in Schools

A common type of network we analyze is friendship networks. In this case, we will use ERGMs to simulate friendship networks within a school. In our simulated world, these networks will be dominated by the following phenomena

- Low density,
- Race homophily,
- Structural balance,
- And age homophily.

If you have been paying attention to the previous chapters, you will notice that, out of these five properties, only one constitutes Markov graphs. Within a tie, homophily and density only depend on ego and alter. In race homophily, only ego and alter's race matter for the tie formation, but, in the case of Structural balance, ego is more likely to befriend alter if a friend of ego is friends with alter, i.e., "the friend of my friend is my friend."

The simulation steps are as follows:

1. Draw a population of n students and randomly distribute race and age across them.
2. Create a network object.
3. Simulate the ties in the empty network.

Here is the code

```
set.seed(712)
n <- 200

# Step 1: Students
race <- sample(c("white", "non-white"), n, replace = TRUE)
age <- sample(c(10, 14, 17), n, replace = TRUE)

# Step 2: Create an empty network
library(ergm)
library(network)
net <- network.initialize(n)

net %v% "race" <- race
net %v% "age" <- age

# Step 3: Simulate a graph
net_sim <- simulate(
  net ~ edges +
  nodematch("race") +
  ttriad +
  absdiff("age"),
  coef = c(-4, .5, .25, -.5)
)
```

What just happened? Here is a line-by-line breakout:

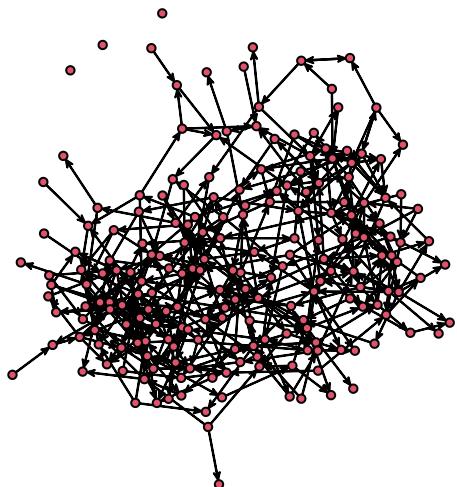
1. `set.seed(712)` Since this is a random simulation, we need to fix a seed so it is reproducible. Otherwise, results would change with every iteration.
2. `n <- 200` We are assigning the value 200 to the object `n`. This will make things easier as, if needed, changing the size of the networks can be done at the top of the code.
3. `race <- sample(c("white", "non-white"), n, replace = TRUE)` We are sampling 200, or actually, `n` values from the vector `c("white", "non-white")` with replacement.
4. `age <- sample(c(10, 14, 17), n, replace = TRUE)` Same as before, but with ages!
5. `library(ergm)` Loading the `ergm` R package, which we need to simulate the networks!
6. `library(network)` Loading the `network` R package, which we need to create the empty

graph.

7. `net <- network.initialize(n)` Creating an empty graph of size n.
8. `net %v% "race" <- race` Using the `%v%` operator, we can access vertices features in the network object. Since race does not exist in the network yet, the operator just creates it. Notice that the number of vertices matches the length of the race vector.
9. `net %v% "age" <- age` Same as with race!
10. `net_sim <- simulate(` Simulating an ERGM! A couple of observations here:
 - a. The LHS (left-hand-side) of the equation has the network, net
 - b. The RHS (you guessed it) has the terms that govern the process.
 - c. For low density, we used the edges term with a corresponding -4.0 for the parameter.
 - d. For race homophily, we used the nodematch("race") with a corresponding 0.5 parameter value.
 - e. For structural balance, we use the ttriad term with parameter 0.25.
 - f. For age homophily, we use the absdiff("age") term with parameter -0.5. This is, in rigor, a term capturing heterophily. Nonetheless, heterophily is the opposite of homophily.

Let's take a quick look at the resulting graph

```
library(sna)
gplot(net_sim)
```



We can now start to see whether we got what we wanted! Before that, let's save the network as a plain-text file so we can practice reading networks back in R!

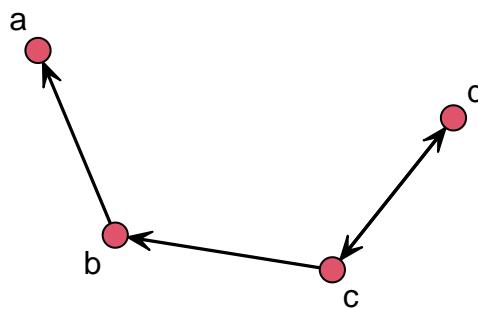
```
write.csv(
  x       = as.edgelist(net_sim),
  file    = "06-edgelist.csv",
  row.names = FALSE
```

```
)
write.csv(
  x           = as.data.frame(net_sim, unit = "vertices"),
  file        = "06-nodes.csv",
  row.names  = FALSE
)
```

9.3 Reading a network

The first step to analyzing network data is to read it in. Many times you'll find data in the form of an adjacency matrix. Other times, data will come in the form of an edgelist. Another common format is the adjacency list, which is a compressed version of an edgelist. Let's see how the formats look like for the following network:

```
example_graph <- matrix(0L, 4, 4, dimnames = list(letters[1:4], letters[1:4]))
example_graph[c(2, 7)] <- 1L
example_graph["c", "d"] <- 1L
example_graph["d", "c"] <- 1L
example_graph <- as.network(example_graph)
set.seed(1231)
gplot(example_graph, label = letters[1:4])
```



- **Adjacency matrix** a matrix of size n by n where the ij -th entry represents the tie between i and j . In a directed network, we say i connects to j , so the i -th row shows the ties i sends to the rest of the network. Likewise, in a directed graph, the j -th column shows the ties sent to j . For undirected graphs, the adjacency matrix is usually upper or lower diagonal. The adjacency matrix of an undirected graph is symmetric, so we don't need to report the same information twice. For example:

```
as.matrix(example_graph)
```

```
##   a b c d
## a 0 0 0 0
## b 1 0 0 0
## c 0 1 0 1
```

```
## d 0 0 1 0
```

- **Edge list** a matrix of size $|E|$ by 2, where $|E|$ is the number of edges. Each entry represents a tie in the graph.

```
as.edgelist(example_graph)
```

```
##      [,1] [,2]
## [1,]    2    1
## [2,]    3    2
## [3,]    3    4
## [4,]    4    3
## attr(,"n")
## [1] 4
## attr(,"vnames")
## [1] "a" "b" "c" "d"
## attr(,"directed")
## [1] TRUE
## attr(,"bipartite")
## [1] FALSE
## attr(,"loops")
## [1] FALSE
## attr(,"class")
## [1] "matrix_edgelist" "edgelist"      "matrix"          "array"
```

The command turns the network object into a matrix with a set of attributes (which are also printed.)

- **Adjacency list** This data format uses less space than edgelists as ties are grouped by ego (source.)

```
igraph::as_adj_list(intergraph::asIgraph(example_graph))
```

```
## [[1]]
## + 1/4 vertex, from 38b50ee:
## [1] 2
##
## [[2]]
## + 2/4 vertices, from 38b50ee:
## [1] 1 3
##
## [[3]]
## + 3/4 vertices, from 38b50ee:
## [1] 2 4 4
##
## [[4]]
## + 2/4 vertices, from 38b50ee:
```

```
## [1] 3 3
```

The function `igraph::as_adj_list` turns the `igraph` object into a list of type adjacency list. In plain text it would look something like this:

```
2
1 3
2 4 4
3 3
```

Here we will deal with an edgelist that includes node information. In my opinion, this is one of the best ways to share network data. Let's read the data into R using the function `read.csv`:

```
edges <- read.csv("06-edgelist.csv")
nodes <- read.csv("06-nodes.csv")
```

We now have two objects of class `data.frame`, `edges` and `nodes`. Let's inspect them using the `head` function:

```
head(edges)
```

```
##   V1   V2
## 1 2   7
## 2 2   41
## 3 3   5
## 4 3   16
## 5 4  138
## 6 5   9
```

```
head(nodes)
```

```
##   vertex.names      race age
## 1                 1 non-white 10
## 2                 2     white 10
## 3                 3     white 17
## 4                 4 non-white 14
## 5                 5 non-white 17
## 6                 6 non-white 14
```

It is always important to look at the data before creating the network. Most common errors happen before reading the data in and could go undetected in many cases. A few examples:

- Headers in the file could be treated as data, or the files may not have headers.
- Ego/alter columns may show in the wrong order. Both the `igraph` and `network` packages take the first and second columns of edgelists as ego and alter.
- Isolates, which wouldn't show in the edgelist, may be missing from the node information set. This is one of the most common errors.
- Nodes showing in the edgelist may be missing from the nodelist.

Both `igraph` and `network` have functions to read edgelist with a corresponding nodelist; the functions `graph_from_data_frame` and `as.network`, respectively. Although, for both cases, you can avoid using a nodelist, it is highly recommended as then you will (a) make sure that isolates are included and (b) become aware of possible problems in the data. A frequent error in `graph_from_data_frame` is nodes present in the edgelist but not in the set of nodes.

```
net_ig <- igraph::graph_from_data_frame(
  d      = edges,
  directed = TRUE,
  vertices = nodes
)
```

Using `as.network` from the `network` package:

```
net_net <- network::as.network(
  x      = edges,
  directed = TRUE,
  vertices = nodes
)
```

As you can see, both syntaxes are very similar. The main point here is that the more explicit we are, the better. Nevertheless, R can be brilliant; being *shy*, i.e., not throwing warnings or errors, is not uncommon. In the next section, we will finally start visualizing the data.

9.4 Visualizing the network

We will focus on three different attributes that we can use for this visualization: Node size, node shape, and node color. While there are no particular rules, some ideas you can follow are:

- **Node size** Use it to describe a continuous measurement. This feature is often used to highlight important nodes, e.g., using one of the many available degree measurements.
- **Node shape** Shapes can be used to represent categorical values. A good figure will not feature too many of them; less than four would make sense.
- **Node color** Like shapes, colors can be used to represent categorical values, so the same idea applies. Furthermore, it is not crazy to use both shape and color to represent the same feature.

Notice that we have not talked about layout algorithms. The R packages to build graphs usually have internal rules to decide what algorithm to use. We will discuss that later on. Let's start by size.

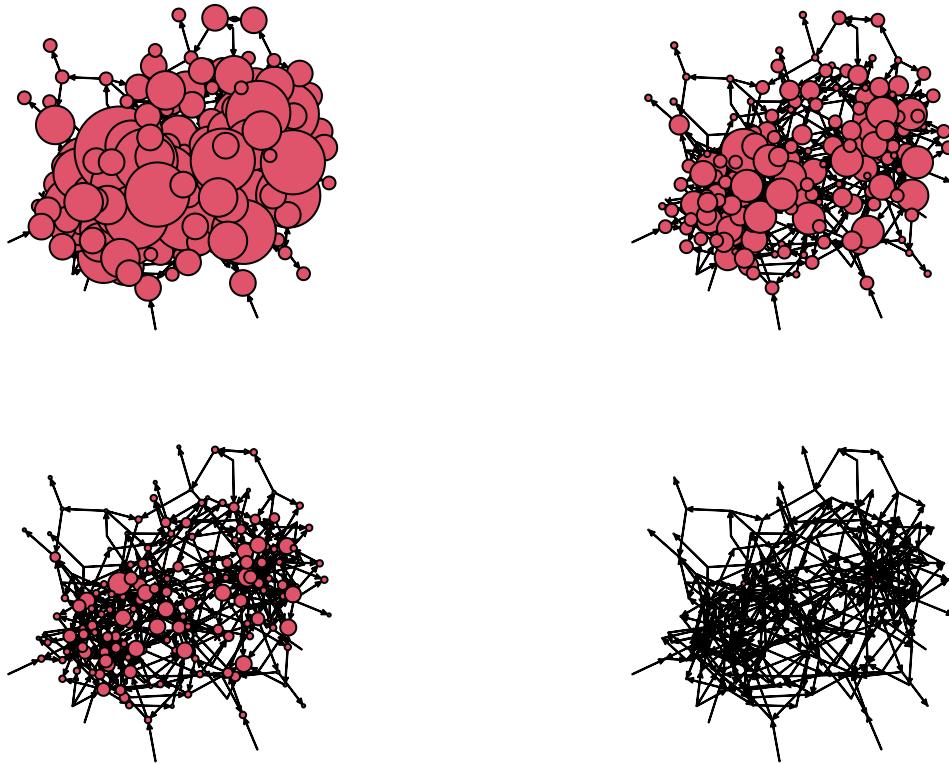
9.4.1 Vertex size

Finding the right scale can be somewhat difficult. We will draw the graph four times to see what size would be the best:

```
# Sized by indegree
net_sim %v% "indeg" <- sna::degree(net_sim, cmode = "indegree")

# Changing device config
op <- par(mfrow = c(2, 2), mai = c(.1, .1, .1, .1))

# Plotting
glayout <- gplot(net_sim, vertex.cex = (net_sim %v% "indeg") * 2)
gplot(net_sim, vertex.cex = net_sim %v% "indeg", coord = glayout)
gplot(net_sim, vertex.cex = (net_sim %v% "indeg")/2, coord = glayout)
gplot(net_sim, vertex.cex = (net_sim %v% "indeg")/10, coord = glayout)
```



```
# Restoring device config
par(op)
```

Line-by-line we did the following:

1. `net_sim %v% "indeg" <- degree(net_sim, cmode = "indegree")` Created a new vertex attribute called `indegree` and assigned it to the network object. The `indegree` is calculated using the `degree` function from the `sna` package. Since `igraph` also has a `degree` function, we are making sure that R uses `sna`'s and not `igraph`'s. The package:: function notation is useful for these cases.
2. `op <- par(mfrow = c(2, 2), mai = c(.1, .1, .1, .1))` This changes the graphical

device information to (a) `mfrow = c(2,2)` have a 2x2 grid by row, meaning that new figures will be added left to right and then top to bottom, and (b) set the margins in the figure to be 0.1 inches in all four sizes.

3. `glayout <- gplot(net_sim, vertex.cex = (net_sim %v% "indeg") * 2)` generating the plot **and** recording the layout. The `gplot` function returns a matrix of size # vertices by 2 with the positions of the vertices. We are also passing the `vertex.cex` argument, which we use to specify the size of each vertex. In our case, we decided to size the vertices proportional to their indegree *times two*.


4. `gplot(net_sim, vertex.cex = net_sim %v% "indeg", coord = glayout)`, again, we are drawing the graph using the coordinates of the previous draw, but now the vertices are half the size of the original figure.

The other two calls are similar to four. If we used `igraph`, setting the size can be more accessible thanks to the `netdiffuseR` R package. Let's start by converting our network to an `igraph` object with the R package `intergraph`.

```
library(intergraph)
library(igraph)


# Converting the network object to an igraph object
net_sim_i <- asIgraph(net_sim)

# Plotting with igraph
plot(
  net_sim_i,
  vertex.size = netdiffuseR::rescale_vertex_igraph(
    vertex.size = V(net_sim_i)$indeg,
    minmax.relative.size = c(.01, .1)
  ),
  layout      = glayout,
  vertex.label = NA
)
```



We could also have tried netplot, which should make things easier and make a better use of the space:

```
library(netplot)
nplot(
  net_sim, layout = glayout,
  vertex.color = "tomato",
  vertex.frame.color = "darkred"
)
```



With a good idea for size, we can now start looking into vertex color.

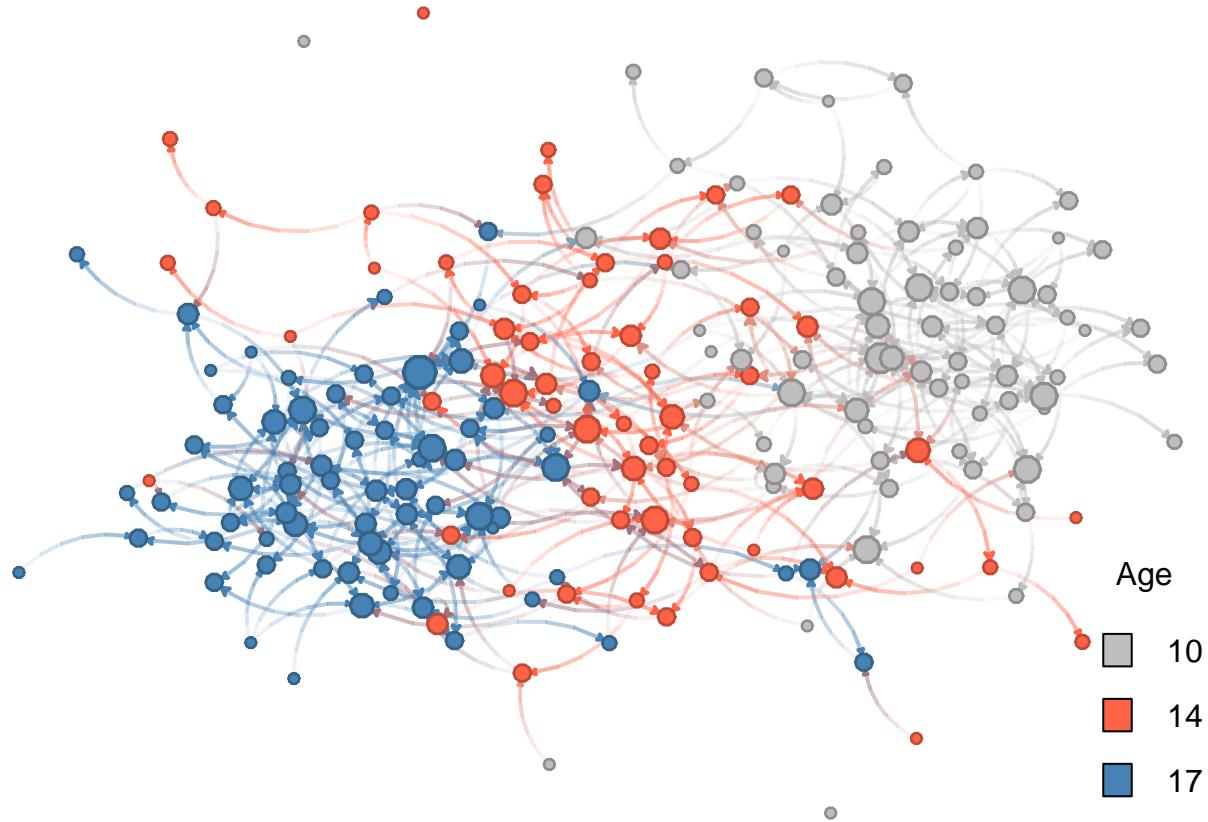
9.4.2 Vertex color

For the color, we will use vertex age. Although age is, by definition, continuous, we only have three values for age. Because of this, we can treat age as categorical. Instead of using `nplot` we will go ahead with `nplot_base`. As of this version of the book, the `netplot` package does not have an easy way to add legends with the core function, `nplot`; therefore, we use `nplot_base` which is compatible with the R function `legend`, as we will now see:

```
# Specifying colors for each vertex
vcolors_palette <- c("10" = "gray", "14" = "tomato", "17" = "steelblue")
vcolors <- vcolors_palette[as.character(net_sim %v% "age")]
net_sim %v% "color" <- vcolors

# Plotting
nplot_base(
  net_ig,
  layout = glayout,
  vertex.color = net_sim %v% "color",
)

# Color legend
legend(
  "bottomright",
  legend = names(vcolors_palette),
  fill   = vcolors_palette,
  bty    = "n",
  title  = "Age"
)
```



Line by line, this is what we just did:

1. `vcolors <- c("10" = "gray", "14" = "tomato", "17" = "steelblue")` we created a character vector with three elements, "gray", "tomato", and "blue". Furthermore, the vector has names assigned to it, "10", "14", and "17"—the ages we have in the network—so that we can access its elements by indexing by name, e.g., if we type `vcolors["10"]` R returns the value "gray".
2. `vcolors <- vcolors[as.character(net_sim %v% "age")]` there are several things going on in this line. First, we extract the attribute "age" from the network using the `%v%` operator. We then transform the resulting vector from integer type to a character type with the function `as.character`. Finally, using the resulting **character vector** with values "10", "14", "17", ..., we retrieve values from `vcolors` name-indexing. The resulting vector is of length equal to the vertex count in the network.
3. `net_sim %v% "color" <- vcolors` creates a new vertex attribute, `color`. The assigned value is the result from subsetting `vcolors` by the ages of each vertex.
4. `nplot_base(...)` finally draws the network. We pass the previously computed vertex coordinates and vertex colors with the new attribute `color`.
5. `legend(...)` Let's see one parameter at a time:
 - a. "bottomright" tells the overall position of the legend
 - b. `legend = names(vcolors)` passes the actual legend (text); in our case the ages of

individuals.

- c. `fill = vcolors` passes the colors associated with the text.
- d. `bty = "n"` suppresses wrapping the legend within a box.
- e. `title = "Age"` sets the title to be "Age".

9.4.3 Vertex shape

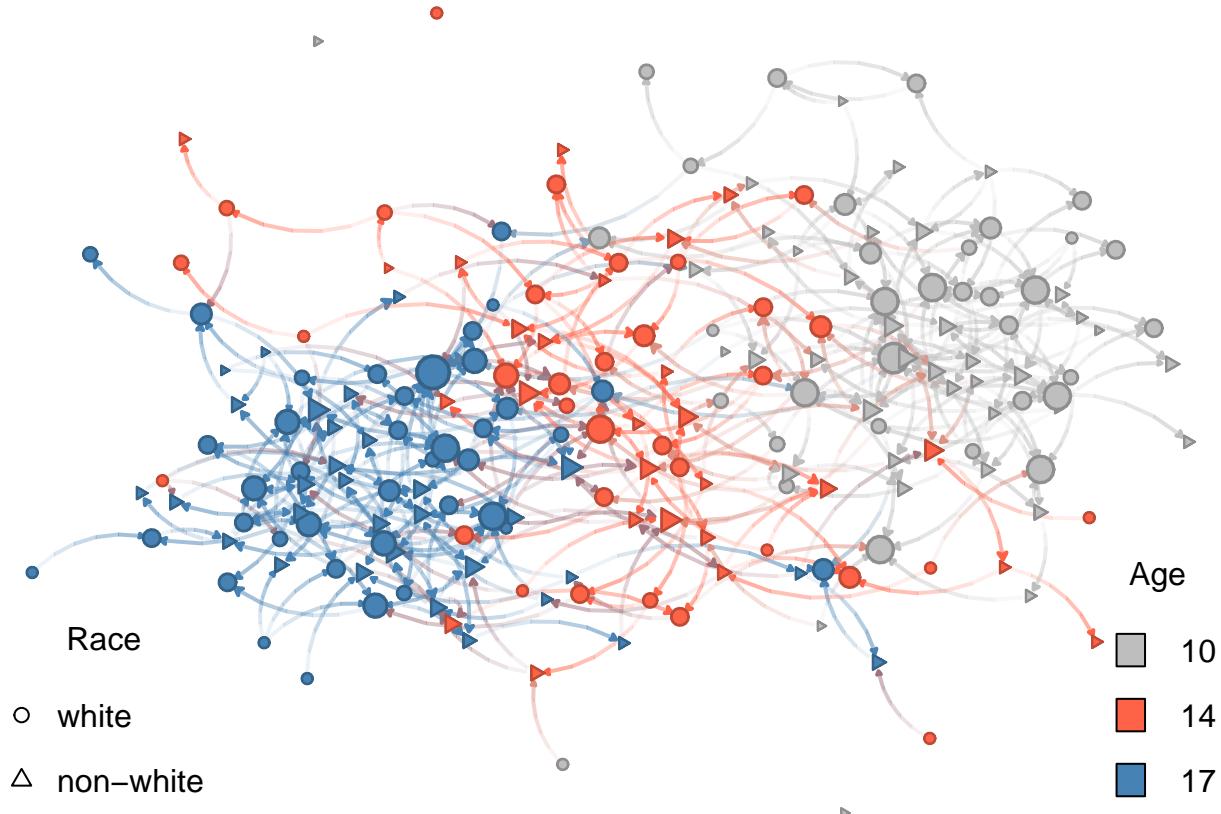
For the color, we will use vertex age. Although age is, by definition, continuous, we only have three values for age. Because of this, we can treat age as categorical.

```
# Specifying the shapes for each vertex
vshape_list <- c("white" = 15, "non-white" = 3)
vshape      <- vshape_list[as.character(net_sim %v% "race")]
net_sim %v% "shape" <- vshape

# Plotting
nplot_base(
  net_ig,
  layout = glayout,
  vertex.color = net_sim %v% "color",
  vertex.nsides = net_sim %v% "shape"
)

# Color legend
legend(
  "bottomright",
  legend = names(vcolors_palette),
  fill   = vcolors_palette,
  bty    = "n",
  title  = "Age"
)

# Shape legend
legend(
  "bottomleft",
  legend = names(vshape_list),
  pch   = c(1, 2),
  bty    = "n",
  title  = "Race"
)
```



Let's now compare the figure to our original ERGM:

1. **Low density (edges)** Without low density, the figure would be a hairball.
2. **Race homophily (nodematch("race"))** Although not surprisingly evident, nodes tend to form small clusters by shape, which, in our model, represents race.
3. **Structural balance (ttriad)** A force, in this case, opposite to low density, higher prevalence of transitive triads makes individuals cluster.
4. **Age homophily (absdiff("age"))** This is the most prominent feature of the graph. In it, nodes are clustered by age.

Of the four features, **age homophily** is the one that stands out. Why is this the case? If we look again at the parameters used in the ERGM and how these interact with vertices' attributes, we will find the answer:

- The log-odds of a new race-homophilic tie are $1 \times \theta_{\text{race-homophily}} = 0.5$.
- But, the log-odd of an age heterophilic tie between, say, 14 and 17 year olds is $|17 - 14| \theta_{\text{age-homophily}} = 3 \times -0.5 = -1.5$.
- Therefore, the effect of heterophily (which is just the opposite of homophily) is significantly larger, actually three times in this case, than the race-homophily effect.

This observation becomes clear if we run another simulation with the same seed, but adjusting for the maximum size the effect of age-homophily can take. A quick-n-dirty way to achieve this

is to re-run the simulation with the nodematch term instead of the absdiff term. This way, we (a) explicitly operationalize the term as homophily (before it was heterophily,) and (b) have both homophily effects have the same influence in the model:

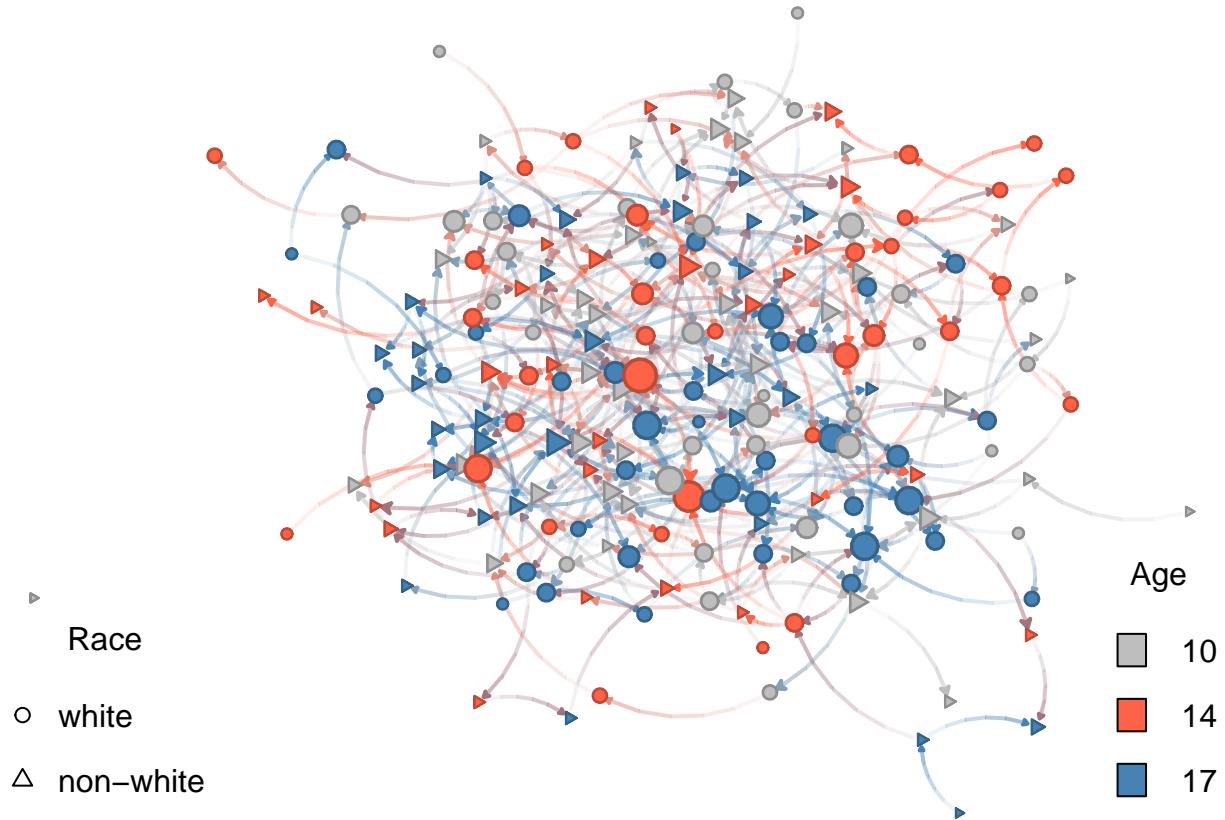
```
net_sim2 <- simulate(
  net ~ edges +
  nodematch("race") +
  ttriad +
  nodematch("age"),
  coef = c(-5, .5, .25, .5) # This line changed
)
```

Re-doing the plot. From the previous graph-drawing, only the graph structure changed. The vertex attributes are the same so we can go ahead and re-use them. Like I mentioned earlier, the nplot_base function currently supports igraph objects, so we will use intergraph::asIgraph to make it work:

```
# Plotting
nplot_base(
  asIgraph(net_sim2),
  # We comment this out to allow for a new layout
  # layout = glayout,
  vertex.color = net_sim %v% "color",
  vertex.nsides = net_sim %v% "shape"
)

# Color legend
legend(
  "bottomright",
  legend = names(vcolors_palette),
  fill = vcolors_palette,
  bty = "n",
  title = "Age"
)

# Shape legend
legend(
  "bottomleft",
  legend = names(vshape_list),
  pch = c(1, 2),
  bty = "n",
  title = "Race"
)
```



As expected, there is no longer a dominant effect in homophily. One important thing we can learn from this final example is that phenomena will not always show themselves in graph visualization. Careful analysis in complex networks is a must.

Chapter 10

Egocentric networks

In egocentric social network analysis (ESNA, for our book,) instead of dealing with a single network, we have as many networks as participants in the study. Egos—the main study subjects—are analyzed from the perspective of their local social network. For a more extended view of ESNA, look at Raffaele Vacca's "*Egocentric network analysis with R*".

In this chapter, I show how to work with one particular type of ESNA data: information generated by the tool Network Canvas. You can download an “artificial” ZIP file containing the outputs from a Network Canvas project here¹. During the entire chapter, we assume that the ZIP file was extracted to the path `data-raw/egonets`. You can go ahead and extract the ZIP by point-and-click, or use the following R code to automate the process:

```
## [1] FALSE  
  
unzip(  
  zipfile = "data-raw/networkCanvasExport-fake.zip",  
  exdir   = "data-raw/egonets"  
)
```

This will extract all the files in `networkCanvasExport-fake.zip` to the subfolder `egonets`. Let's take a look at the first few files:

```
head(list.files(path = "data-raw/egonets"))  
## [1] "I_-59190_BRB9111_attributeList_Person.csv"  
## [2] "I_-59190_BRB9111_edgeList_Knows.csv"  
## [3] "I_-59190_BRB9111_ego.csv"  
## [4] "I_-59190_BRB9111.graphml"  
## [5] "I-100BB_00B95-90_attributeList_Person.csv"  
## [6] "I-100BB_00B95-90_edgeList_Knows.csv"
```

As you can see, for each ego in the dataset there are four files:

¹I thank Jacqueline Kent-Marwick who provided me with what I used as a baseline to generate the artificial Network Canvas export.

- ...attributeList_Person.csv: Attributes of the alters.
- ...edgeList_Knows.csv: Edgelist indicating the ties between the alters.
- ...ego.csv: Information about the egos.
- ...graphml: And a graphml file that contains the egonets.

The next sections will illustrate, file by file, how to read the information into R, apply any required processing, and store the information for later use. We start with the graphml files.

10.1 Network files (graphml)

The graphml files can be read directly with igraph's `read_graph` function. The key is to take advantage of R's lists to avoid writing over and over the same block of code, and instead, manage the data through lists.

Just like any data-reading function, `read_graph` function requires a file path to the network file.

The function we will use to list the required files is `list.files()`:

```
# We start by loading igraph
library(igraph)

# Listing all the graphml files
graph_files <- list.files(
  path      = "data/raw/egonets", # Where are these files
  pattern   = "*.graphml",       # Specify a pattern for only listing graphml
  full.names = TRUE             # And we make sure we use the full name
                                # (path.) Otherwise, we would only get names.
)

# Taking a look at the first three files we got
graph_files[1:3]
## [1] "data/raw/egonets/I--59190_BRB9111.graphml"
## [2] "data/raw/egonets/I-100BB_00B95-90.graphml"
## [3] "data/raw/egonets/I-1BB79950-0-7.graphml"

# Applying igraph's read_graph
graphs <- lapply(
  X      = graph_files,        # List of files to read
  FUN    = read_graph,         # The function to apply
  format = "graphml"          # Argument passed to read_graph
)
```

If the operation succeeded, then the previous code block should generate a list of igraph objects named `graphs`. Let's take a peek at the first two:

```

graphs[[1]]
## IGRAPH 41d3cfa U--- 12 25 --
## + attr: age (v/n), healthy_diet (v/n), gender_1 (v/l), eat_with_2
## | (v/l), id (v/c)
## + edges from 41d3cfa:
## [1] 1-- 3 1-- 2 1-- 6 1-- 5 1-- 4 1-- 8 1--11 1--10 2-- 3 3-- 7 3-- 4 3-- 5
## [13] 3-- 6 2-- 7 2-- 4 2-- 5 2-- 6 5-- 6 6--10 7-- 9 4-- 5 5-- 7 4--11 6-- 7
## [25] 4-- 7
graphs[[2]]
## IGRAPH f4204e3 U--- 16 47 --
## + attr: age (v/n), healthy_diet (v/n), gender_1 (v/l), eat_with_2
## | (v/l), id (v/c)
## + edges from f4204e3:
## [1] 7--13 1-- 5 1-- 6 1-- 4 1-- 2 7--15 1-- 3 11--13 1--10 1--16
## [11] 4-- 6 2-- 6 6-- 7 1--11 11--15 6-- 9 6-- 8 3-- 9 5--15 4-- 5
## [21] 2-- 5 5-- 8 5-- 7 5--10 3-- 5 6--14 12--13 6--13 3--13 2-- 3
## [31] 3-- 4 3--16 3--11 10--14 7--14 2-- 4 2--10 2--15 10--12 4-- 7
## [41] 6--10 5--11 9--10 1-- 9 1--12 3--12 4--14

```

As always, one of the first things we do with networks is visualize them. We will use the `netplot` R package (by yours truly) to draw the figures:


```

library(netplot)
library(gridExtra)

# Graph layout is random
set.seed(1231)

# The grid.arrange allows putting multiple netplot graphs into the same page
grid.arrange(
  nplot(graphs[[1]]),
  nplot(graphs[[2]]),
  nplot(graphs[[3]]),
  nplot(graphs[[4]]),
  ncol = 2, nrow = 2
)

```



Great! Since nodes in our network have features, we can add a little bit of color. We will use the `eat_with_2` variable, which is coded as TRUE or FALSE. Vertex colors can be specified using the `vertex.color` argument of the `nplot` function. In our case, we will specify colors passing a vector of length equal to the number of nodes in the graph. Furthermore, since we will be doing this multiple times, it is worthwhile writing a function:

```
# A function to color by the eat with variable
color_it <- function(net) {

  # Coding eat_with_2 to be 1 (FALSE) or 2 (TRUE)
  eatswith <- V(net)$eat_with_2

  # Subsetting the color
  ifelse(eatswith, "purple", "darkgreen")

}
```


This function takes two arguments: a network and a vector of two colors. Vertex attributes in igraph can be accessed through the `V(...)$...` function. For this example, to access the attribute `eat_with_2` in the network `net`, we type `V(net)$eat_with_2`. Finally, individuals who have `eat_with_2` equal to true will be colored purple, otherwise, if equal to FALSE, will be colored darkgreen. Before plotting the networks, let's see what we get when we access the

eat_with_2 attribute in the first graph:

```
V(graphs[[1]])$eat_with_2
## [1] TRUE TRUE FALSE FALSE TRUE TRUE FALSE FALSE TRUE TRUE FALSE FALSE
```

A logical vector. Now let's redraw the figures:

```
grid.arrange(
  nplot(graphs[[1]], vertex.color = color_it(graphs[[1]])),
  nplot(graphs[[2]], vertex.color = color_it(graphs[[2]])),
  nplot(graphs[[3]], vertex.color = color_it(graphs[[3]])),
  nplot(graphs[[4]], vertex.color = color_it(graphs[[4]])),
  ncol = 2, nrow = 2
)
```



Since most of the time, we will be dealing with a large number of egonets, you may want to draw each network independently; the following code-block does exactly that. First, if needed, will create a folder to store the networks. Then, using the lapply function, it will use netplot::nplot() to draw the networks, add a legend, and save the graph as .../graphml_[number].png, where [number] will go from 01 to the total number of networks in graphs.

```
if (!dir.exists("egonets/figs/egonets"))
  dir.create("egonets/figs/egonets", recursive = TRUE)
```

```

lapply(seq_along(graphs), function(i) {

  # Creating the device
  png(sprintf("egonets/figs/egonets/graphml_%02i.png", i))

  # Drawing the plot
  p <- nplot(
    graphs[[i]],
    vertex.color = color_it(graphs[[i]])
  )

  # Adding a legend
  p <- nplot_legend(
    p,
    labels = c("eats with: FALSE", "eats with: TRUE"),
    pch     = 21,
    packgrob.args = list(side = "bottom"),
    gp       = gpar(
      fill = c("darkgreen", "purple")
    ),
    ncol = 2
  )

  print(p)

  # Closing the device
  dev.off()
})

```

10.2 Person files

Like before, we list the files ending in Person.csv (with the full path,) and read them into R. While R has the function `read.csv`, here I use the function `fread` from the `data.table` R package. Alongside `dplyr`, `data.table` is one of the most popular data-wrangling tools in R. Besides syntax, the biggest difference between the two is performance; `data.table` is significantly faster than any other data management package in R, and is a great alternative for handling large datasets. The following code block loads the package, lists the files, and reads them into R.

```

# Loading data.table
library(data.table)

```

```

# Listing the files
person_files <- list.files(
  path      = "data-raw/egonets",
  pattern   = "*Person.csv",
  full.names = TRUE
)

# Loading all into a single list
persons <- lapply(person_files, fread)

# Looking into the first element
persons[[1]]
##      nodeID age
## 1:      1 45
## 2:      2 32
## 3:      3 31
## 4:      4 45
## 5:      5 43
## 6:      6 47
## 7:      7 45
## 8:      8 62
## 9:      9 28
## 10:     10 41
## 11:     11 41
## 12:     12 46
## 13:     13 46
## 14:     14 46
## 15:     15 62
## 16:     16 41

```

A common task is adding an identifier to each dataset in `persons` so we know from to which ego they belong. Again, the `lapply` function is our friend:

```

persons <- lapply(seq_along(persons), function(i) {
  persons[[i]][, dataset_num := i]
})

```

In `data.table`, variables are created using the `:=` symbol. The previous code chunk is equivalent to this:

```

for (i in 1:length(persons)) {
  persons[[i]]$dataset_num <- i
}

```

If needed, we can transform the list `persons` into a `data.table` object (i.e., a single

`data.frame`) using the `rbindlist` function². The next code block uses that function to combine the `data.tables` into a single dataset.


```
# Combining the datasets
persons <- rbindlist(persons)
persons
##      nodeID age dataset_num
## 1:      1   45         1
## 2:      2   32         1
## 3:      3   31         1
## 4:      4   45         1
## 5:      5   43         1
## ...
## 271:     7   43        19
## 272:     8   48        19
## 273:     9   70        19
## 274:    10   46        19
## 275:    11   50        19
```

Now that we have a single dataset, we can do some data exploration. For example, we can use the package `ggplot2` to draw a histogram of alters' ages.

```
# Loading the ggplot2 package
library(ggplot2)

# Histogram of age
ggplot(persons, aes(x = age)) +          # Starting off the plot
  geom_histogram(fill = "purple") +        # Adding a histogram
  labs(x = "Age", y = "Frequency") +      # Changing the x/y axis labels
  labs(title = "Alter's Age Distribution") # Adding a title
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

²Although not the same, `rbindlist` (almost always) yields the same result as calling the function `do.call`. In particular, instead of executing the call `rbindlist(persons)`, we could have used `do.call(rbind, persons)`.



10.3 Ego files

The ego files contain information about egos (duh!). Again, we will read them all at once using `list.files` + `lapply`:

```
# Listing files ending with *ego.csv
ego_files <- list.files(
  path      = "data-raw/egonets",
  pattern   = "*ego.csv",
  full.names = TRUE
)

# Reading the files with fread
egos <- lapply(ego_files, fread)

# Combining them
egos <- rbindlist(egos)
head(egos)
##                                     networkCanvasEgoUUID networkCanvasCaseID
## 1: I-11ca3a78c-62f131f37169-c139217a1f6    I--59190_BRB9111
## 2: I-fef-ab-4-5a--7-35c4f23-96eb32-34ea    I-100BB_00B95-90
```

```

## 3: I2f1bd0b6d-f71f4664cf-d-26-97408f22d      I-1BB79950-0-7
## 4: Id36bb-3b2bcfd2a6239b1103134c6b3d1d6      I000091I_RB010B5
## 5: I436d32fc67fb5c6-23-244f353849b120cd      I019051R0_RRR0-0
## 6: Ibff1f-2-34162bb5f2c36b8241--316a-fff      I01B11-I1101_44R
##                           networkCanvasSessionID
## 1: I612b7a1af---0880b-70698204-b-8dbf09
## 2: If5e0-f-26cbec070760f-e6b6d26ebfb06f
## 3: I825c293a1304-e5-cbea8a80aae05b305fa
## 4: I1b8a7d0f6b4-8298c9-848-9186d68a7f3c
## 5: Ie620be37b75983c49ac63-38-425227c959
## 6: Ie3-134323ed40-0e-d954b3d-febbcb9363
##                           networkCanvasProtocolName      sessionStart
## 1: Postpartum social networks with sociogram_V5 2023-02-22 23:41:59
## 2: Postpartum social networks with sociogram_V5 2023-02-10 21:46:02
## 3: Postpartum social networks with sociogram_V5 2023-03-01 16:52:09
## 4: Postpartum social networks with sociogram_V5 2023-01-26 20:38:07
## 5: Postpartum social networks with sociogram_V5 2023-02-06 14:55:57
## 6: Postpartum social networks with sociogram_V5 2023-03-16 18:20:02
##                           sessionFinish      sessionExported
## 1: 2023-02-23 01:47:00 2023-02-23 01:47:08
## 2: 2023-02-11 01:29:32 2023-02-11 01:34:12
## 3: 2023-03-02 16:51:20 2023-03-02 17:04:42
## 4: 2023-01-26 22:03:20 2023-01-26 22:03:34
## 5: 2023-02-06 15:49:38 2023-02-06 15:56:42
## 6: 2023-03-17 21:11:09 2023-03-17 21:16:15

```

A cool thing about `data.table` is that, within square brackets, we can manipulate the data referring to the variables directly. For example, if we wanted to calculate the difference between `sessionFinish` and `sessionStart`, using base R we would do the following:

```
egos$total_time <- egos$sessionFinish - egos$sessionStart
```


Whereas with `data.table`, variable creation is much more straightforward (notice that instead of using `<-` or `=` to assign a variable, we use the `:=` operator):

```
# How much time?
egos[, total_time := sessionFinish - sessionStart]
```

We can also visualize this using `ggplot2`:

```
ggplot(egos, aes(x = total_time)) +
  geom_histogram() +
  labs(x = "Time in minutes", y = "Count") +
  labs(title = "Total time spent by egos")
## Don't know how to automatically pick scale for object of type <difftime>. Defaulting to
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Total time spent by egos



10.4 Edgelist files

As I mentioned earlier, since we are reading the graphml files, using the edgelist may not be needed. Nevertheless, the process to import the edgelist file to R is the same we have been applying: list the files and read them all at once using `lapply`:

```
# Listing all files ending in Knows.csv
edgelist_files <- list.files(
  path = "data-raw/egonets",
  pattern = "*Knows.csv",
  full.names = TRUE
)

# Reading all files at once
edgelists <- lapply(edgelist_files, fread)
```

To avoid confusion, we can also add ids corresponding to the file number. Once we do that, we can combine all files into a single `data.table` object using `rbindlist`:

```
edgelists <- lapply(seq_along(edgelists), function(i) {
  edgelists[[i]][, dataset_num := i]
})
```

```

edgelists <- rbindlist(edgelists)

head(edgelists)
##   edgeID from to           networkCanvasEgoUUID
## 1:      1   1  5 I839f-8fa8f8aeb8-eaf---ba8-cf3908f3a
## 2:      2   1 10 If81a9c0f-9f4f28ccf-c4c923a-8-0f5fce
## 3:      3   1  9 I899ffe-27-3-a3-ca2fb7f7-ca8e7715ce9
## 4:      4   1 10 I814efaba88cbb02caa8c89790-83beeaf9-
## 5:      5   7  6 Ifd-0eec2e08974eaf2b79f-9efb7e3-8998
## 6:      6   2  6 I-28fe89cc-fc5db3825b92-ae87c-c18e3d
##
##           networkCanvasSourceUUID           networkCanvasTargetUUID dataset_num
## 1: I720400eb19bccce-77cee773289b02-fe7e I4d5--16a08f8ba463c6458f8979e-65fa9d
## 2: I-b469c0-60f8bbb543-32-628-216f9-038 I-6cf8-f3da-4-96-87efaf5daaa48ba5e5c
## 3: Ifa4933-9baaf5fc-f-e4f5c5e5-ff34-f-f I5-f69a6eaa-5956e8897ca999-ffb6ed-e1
## 4: I4cb-904496b1-6194bcb51b58444b40-ef8 I3e5-6c8d5e0f086--e-5ab45-4-5aaa5-0e
## 5: I0ab7--b7a0ee71e54c1e93cdb-4ca5ab1-b I5-b-9-7eca5ab5-91915ba9b6565a6e42cc
## 6: Ic80142fc4c431009e84b3-ab3f-9b0eab03 Ie0a24eea4e01a4340343a0-66723-a-9970
##
##           networkCanvasTargetUUID dataset_num
## 1: Id1c8befd46bdd195c-ce91a8-bc0---4f0e          1
## 2: I757b4a-3ea4d95--b9ebb9db3d55dcba-f-c          1
## 3: I92a62925ff9-e2f27-6ef97d-29fb729624          1
## 4: I7f--da48-46a64-b972c-ef6bbec--64cb4          1
## 5: I-eaa7e95659-9cf01a4f5fd69af54e6-d60          1
## 6: I69060e8a-454609-faa04cd3eeb-5-9550-          1

```

10.5 Putting all together

In this last part of the chapter, we will use the `igraph` and `ergm` packages to generate features (covariates, controls, independent variables, or whatever you call them) at the ego-network level. Once again, the `lapply` function is our friend

10.5.1 Generating statistics using `igraph`

The `igraph` R package has multiple high-performing routines to compute graph-level statistics. For now, we will focus on the following statistics: vertex count, edge count, number of isolates, transitivity, and modularity based on betweenness centrality:

```

net_stats <- lapply(graphs, function(g) {

  # Calculating modularity
  groups <- cluster_edge_betweenness(g)

```

```
# Computing the stats
data.table(
  size      = vcount(g),
  edges     = ecount(g),
  nisolates = sum(degree(g) == 0),
  transit   = transitivity(g, type = "global"),
  modular   = modularity(groups)
)
})
```

Observe we count isolates using the degree() function. We can combine the statistics into a single data.table using the rbindlist function:

```
net_stats <- rbindlist(net_stats)

head(net_stats)
##   size edges nisolates    transit    modular
## 1: 12    25       1 0.6750000 0.012000000
## 2: 16    47       0 0.4332130 0.003395201
## 3: 16    58       0 0.5612009 0.002675386
## 4: 15    75       0 0.8515112 0.000000000
## 5: 15    52       0 0.5780488 0.000000000
## 6: 17    68       0 0.6291161 0.025735294
```

10.5.2 Generating statistics based on ergm

The ergm R package has a much larger set of graph-level statistics we can add to our models.³ The key to generating statistics based on the ergm package is the summary_formula function. Before we start using that function, we first need to convert the igraph networks to network objects, which are the native object class for the ergm package. We use the intergraph R package for that, and in particular, the asNetwork function:

```
# Loading the required packages
library(intergraph)
library(ergm)

## Loading required package: network
##
## 'network' 1.18.2 (2023-12-04), part of the Statnet Project
## * 'news(package="network")' for changes since last version
## * 'citation("network")' for citation information
## * 'https://statnet.org' for help, support, and other information
##
## Attaching package: 'network'
```

³There's an obvious reason, ERGMs are all about graph-level statistics!

```

## The following objects are masked from 'package:igraph':
##
##      %C%, %S%, add.edges, add.vertices, delete.edges, delete.vertices,
##      get.edge.attribute, get.edges, get.vertex.attribute, is.bipartite,
##      is.directed, list.edge.attributes, list.vertex.attributes,
##      set.edge.attribute, set.vertex.attribute
##
## 'ergm' 4.6-7262 (2023-09-04), part of the Statnet Project
## * 'news(package="ergm")' for changes since last version
## * 'citation("ergm")' for citation information
## * 'https://statnet.org' for help, support, and other information
## 'ergm' 4 is a major update that introduces some backwards-incompatible
## changes. Please type 'news(package="ergm")' for a list of major
## changes.

# Converting all "igraph" objects in graphs to network "objects"
graphs_network <- lapply(graphs, asNetwork)

```

With the network objects ready, we can proceed to compute graph-level statistics using the summary_formula function. Here we will only look into: the number of triangles, gender homophily, and healthy-diet homophily:

```

net_stats_ergm <- lapply(graphs_network, function(n) {

  # Computing the statistics
  s <- summary_formula(
    n ~ triangles +
      nodematch("gender_1") +
      nodematch("healthy_diet")
  )

  # Saving them as a data.table object
  data.table(
    triangles      = s[1],
    gender_homoph = s[2],
    healthyd_homoph = s[3]
  )
})

```

Once again, we use rbindlist to combine all the network statistics into a single data.table object:

```

net_stats_ergm <- rbindlist(net_stats_ergm)
head(net_stats_ergm)
##   triangles gender_homoph healthyd_homoph

```

## 1:	27	11	3
## 2:	40	30	20
## 3:	81	40	29
## 4:	216	33	38
## 5:	79	44	19
## 6:	121	38	16

10.6 Saving the data

We end the chapter saving all our work into four datasets:

- Network statistics (as a csv file)
- Igraph objects (as a rda file, which we can read back using `read.rds`)
- Network objects (idem)
- Person files (alter's information, as a csv file.)

CSV files can be saved either using `write.csv` or, as we do here, `fwrite` from the `data.table` package:

```
# Network attributes
master <- cbind(egos, net_stats, net_stats_ergm)
fwrite(master, file = "data/network_stats.csv")

# Networks
saveRDS(graphs, file = "data/networks_igraph.rds")
saveRDS(graphs_network, file = "data/networks_network.rds")

# Attributes
fwrite(persons, file = "data/persons.csv")
```


Chapter 11

Hypothesis testing in networks

Overall, there are many ways in which we can see hypothesis testing within the networks context:

1. **Comparing two or more networks**, e.g., we want to see if the density of two networks are *equal*.
2. **Prevalence of a motif/pattern**, e.g., check whether the observed number of transitive triads is different from that expected as of by chance.
3. **Multivariate using ERGMs**, e.g., jointly test whether homophily and two stars are the motifs that drive network structure.

The latter we already review in the ERGM chapter. In this part, we will look at types one and two; both using non-parametric methods.

11.1 Comparing networks

Imagine that we have two graphs, $(G_1, G_2) \in \mathcal{G}$, and we would like to assess whether a given statistic $s(\cdot)$, e.g., density, is equal in both of them. Formally, we would like to asses whether $H_0 : s(G_1) - s(G_2) = k$ vs $H_a : s(G_1) - s(G_2) \neq k$.

As usual, the true distribution of $s(\cdot)$ is unknown, thus, one approach that we could use is a non-parametric bootstrap test.

11.1.1 Network bootstrap

The non parametric bootstrap and jackknife methods for social networks were introduced by (T. A. B. Snijders and Borgatti 1999). The method itself is used to generate standard errors for network level statistics. Both methods are implemented in the R package `netdiffuseR`.

11.1.2 When the statistic is normal

When we deal with things that are normally distributed, e.g., sample means like density¹, we can make use of the Student's distribution for making inference. In particular, we can use Bootstrap/Jackknife to approximate the standard errors of the statistic for each network:

1. Since $s(G_i) \sim N(\mu_i, \sigma_i^2/m_i)$ for $i \in \{1, 2\}$, in the case of the density, $m_i = n_i * (n_i - 1)$. The statistic is then:

$$s(G_1) - s(G_0) \sim N(\mu_1 - \mu_0, \sigma_1^2/m_1 + \sigma_0^2/m_0)$$

Thus

$$\frac{s(G_1) - s(G_0) - \mu_1 + \mu_0}{\sqrt{\sigma_1^2/m_1 + \sigma_0^2/m_0}} \sim t_{m_1+m_0-2}$$

But, if we are testing $H_0 : \mu_1 - \mu_0 = k$, then, under the null

$$\frac{s(G_1) - s(G_0) - k}{\sqrt{\sigma_1^2/m_1 + \sigma_0^2/m_0}} \sim t_{m_1+m_0-2}$$

Where We now proceede to approximate the variances.

2. Using the *plugin principle* (Efron and Tibshirani 1994), we can approximate the variances using Bootstrap/Jackknife, i.e., compute $\hat{\sigma}_1^2 \approx \sigma_1^2/m_1$ and $\hat{\sigma}_0^2 \approx \sigma_0^2/m_0$. Using netdiffuseR

```
library(netdiffuseR)

# Obtain a 100 replicates
sg1 <- bootnet(g1, function(i, ...) sum(i)/(nnodes(i) * (nnodes(i) - 1)), R = 100)
sg2 <- bootnet(g2, function(i, ...) sum(i)/(nnodes(i) * (nnodes(i) - 1)), R = 100)

# Retrieving the variances
hat_sigma1 <- sg1$var_t
hat_sigma2 <- sg2$var_t

# And the actual values
sg1 <- sg1$t0
sg2 <- sg2$t0
```

3. With the approximates in hand, we can then use the the “t-test table” to retrieve the corresponding value, in R:

¹Density is indeed a sample mean as we are, in principle computing the average of a sequence of Bernoulli variables. Formally: $\text{density}(G) = \frac{1}{n(n-1)} \sum_{ij} A_{ij}$.

```
# Building the statistic
k <- 0 # For equal variances
tstat <- (sg1 - sg2 - k)/(sqrt(hat_sigma1 + hat_sigma2))

# Computing the pvalue
m1 <- nnodes(g1)*(nnodes(g1) - 1)
m2 <- nnodes(g2)*(nnodes(g2) - 1)
pt(tstat, df = m1 + m2 - 2)
```

11.1.3 When the statistic is NOT normal

In the case that the statistic is not normally distributed, we cannot use the t-statistic any longer. Nevertheless, the Bootstrap can come to help. While in general it is better to use distributions of pivot statistics (see (Efron and Tibshirani 1994)), we can still leverage the power of this method to make inferences. For this example, $s(\cdot)$ will be the range of the threshold in a diffusion graph.

As before, imagine that we are dealing with an statistic $s(\cdot)$ for two different networks, and we would like to asses whether we can reject H_0 or fail to reject it. The procedure is very similar:

1. One approach that we can test is whether $k \in \text{ConfInt}(s(G_1) - s(G_2))$. Building confidence intervals with bootstrap could be more intuitive.
2. Like before, we use bootstrap to generate a distribution of $s(G_1)$ and $s(G_2)$, in R:

```
# Obtain a 1000 replicates
sg1 <- bootnet(g1, function(i, ...) range(threshold(i)), R = 1000)
sg2 <- bootnet(g2, function(i, ...) range(threshold(i)), R = 1000)

# Retrieving the distributions
sg1 <- sg1$boot$t
sg2 <- sg2$boot$t

# Define the statistic
sdiff <- sg1 - sg2
```

3. Once we have $sdiff$, we can proceed and compute the, for example, 95% confidence interval, and evaluate whether k falls within. In R:

```
diff_ci <- quantile(sdiff, probs = c(0.025, .975))
```

This corresponds to what Efron and Tibshirani call “percentile interval.” This is easy to compute, but a better approach is using the “BCa” method, “Bias Corrected and Accelerated.” (TBD)

11.2 Examples

11.2.1 Average of node-level stats

Supposed that we would like to compare something like average indegree. In particular, for both networks, G_1 and G_2 , we compute the average indegree per node:

$$s(G_1) = \text{AvgIndeg}(G_1) = \frac{1}{n} \sum_i \sum_{j \neq i} A_{ji}^1$$

where A_{ji}^1 equals one if vertex j sends a tie to i . In this case, since we are looking at an average, we have that $\text{AvgIndeg}(G_1) \sim N(\mu_1, \sigma_1^2/n)$. Thus, taking advantage of the normality of the statistic, we can build a test statistic as follows:

$$\frac{s(G_1) - s(G_2) - k}{\sqrt{\hat{\sigma}_1^2 + \hat{\sigma}_2^2}} \sim t_{n_1+n_2-2}$$

Where $\hat{\sigma}_i$ is the bootstrap standard error, and $k = 0$ when we are testing equality. This distributes t with $n_1 + n_2 - 2$ degrees of freedom. As a difference from the previous example using density, the degrees of freedom for this test are less as, instead of having an average across all entries of the adjacency matrix, we have an average across all vertices.

Chapter 12

Network diffussion with netdiffuseR

This chapter mainly was based on the 2018 and 2019 tutorials of netdiffuseR at the Sunbelt conference. The source code of the tutorials, taught by Thomas W. Valente and George G. Vega Yon (author of this book), is available [here](#).

12.1 Network diffusion of innovation

12.1.1 Diffusion networks

- Explains how new ideas and practices (innovations) spread within and between communities.
- While a lot of factors have been shown to influence diffusion (Spatial, Economic, Cultural, Biological, etc.), Social Networks is a prominent one.
- There are many components in the diffusion network model including network exposures, thresholds, infectiousness, susceptibility, hazard rates, diffusion rates (bass model), clustering (Moran's I), and so on.

12.1.2 Thresholds

- One of the canonical concepts is the network threshold. Network thresholds (Valente, 1995; 1996), τ , are defined as the required proportion or number of neighbors that lead you to adopt a particular behavior (innovation), $a = 1$. In (very) general terms

$$a_i = \begin{cases} 1 & \text{if } \tau_i \leq E_i \\ 0 & \text{Otherwise} \end{cases} \quad E_i \equiv \frac{\sum_{j \neq i} \mathbf{X}_{ij} a_j}{\sum_{j \neq i} \mathbf{X}_{ij}}$$

Where E_i is i's exposure to the innovation and \mathbf{X} is the adjacency matrix (the network).

- This can be generalized and extended to include covariates and other network weighting schemes (that's what **netdiffuseR** is all about).

12.2 The **netdiffuseR** R package

12.2.1 Overview

netdiffuseR is an R package that:

- Is designed for Visualizing, Analyzing, and Simulating network diffusion data (in general).
- Depends on some pretty popular packages:
 - *RcppArmadillo*: So it's fast,
 - *Matrix*: So it's big,
 - *statnet* and *igraph*: So it's not from scratch
- Can handle big graphs, e.g., an adjacency matrix with more than 4 billion elements (PR for RcppArmadillo)
- Already on CRAN with ~6,000 downloads since its first version, Feb 2016,
- A lot of features to make it easy to read network (dynamic) data, making it a companion of other net packages.

12.2.2 Datasets

- **netdiffuseR** has the three classic Diffusion Network Datasets:

- `medInnovationsDiffNet` Doctors and the innovation of Tetracycline (1955).
- `brfarmersDiffNet` Brazilian farmers and the innovation of Hybrid Corn Seed (1966).
- `kfamilyDiffNet` Korean women and Family Planning methods (1973).

```
brfarmersDiffNet
```

```
## Dynamic network of class -diffnet-
## Name : Brazilian Farmers
## Behavior : Adoption of Hybrid Corn Seeds
## # of nodes : 692 (1001, 1002, 1004, 1005, 1007, 1009, 1010, 1020, ...)
## # of time periods : 21 (1946 - 1966)
## Type : directed
## Final prevalence : 1.00
## Static attributes : village, idold, age, liveout, visits, contact, coo... (146)
## Dynamic attributes : -
```

```
medInnovationsDiffNet
```

```
## Dynamic network of class -diffnet-
## Name : Medical Innovation
```

```

## Behavior : Adoption of Tetracycline
## # of nodes : 125 (1001, 1002, 1003, 1004, 1005, 1006, 1007, 1008, ...)
## # of time periods : 18 (1 - 18)
## Type : directed
## Final prevalence : 1.00
## Static attributes : city, detail, meet, coll, attend, proage, length, ... (58)
## Dynamic attributes : -
kfamilyDiffNet

```

```

## Dynamic network of class -diffnet-
## Name : Korean Family Planning
## Behavior : Family Planning Methods
## # of nodes : 1047 (10002, 10003, 10005, 10007, 10010, 10011, 10012, 10014,
## # of time periods : 11 (1 - 11)
## Type : directed
## Final prevalence : 1.00
## Static attributes : village, recnol, studnol, areal, id1, nmagel, nmag... (430)
## Dynamic attributes : -


```

```

set.seed(12315)
x <- rdifffnet(
  400, t = 6, rgraph.args = list(k=6, p=.3),
  seed.graph = "small-world",
  seed.nodes = "central", rewire = FALSE, threshold.dist = 1/4
)
plot(x)


```

Diffusion network in time 1




```
plot_diffnet(x)
```

Diffusion Network




```
plot_diffnet2(x)
```


Diffusion dynamics



Adopters and Cumulative Adopters





```
plot_threshold(x)
```



```
plot_infectsuscep(x, K=2)
```

```
## Warning in plot_infectsuscep.list(graph$graph, graph$toa, t0, normalize, : When
## applying logscale some observations are missing.
```



12.2.4 Problems

1. Using the diffnet object in intro.rda, use the function plot_threshold specifying shapes and colors according to the variables ItrustMyFriends and Age. Do you see any pattern?

12.3 Simulation of diffusion processes

Before we start, a review of the concepts we will be using here

1. Exposure: Proportion/number of neighbors that have adopted an innovation at each point in time.
2. Threshold: The proportion/number of your neighbors who had adopted at or one time period before ego (the focal individual) adopted.
3. Infectiousness: How much i 's adoption affects her alters.
4. Susceptibility: How much i 's alters' adoption affects her.
5. Structural equivalence: How similar is i to j in terms of position in the network.

12.3.1 Simulating diffusion networks

We will simulate a diffusion network with the following parameters:

1. Will have 1,000 vertices,
2. Will span 20 time periods,
3. The initial adopters (seeds) will be selected at random,
4. Seeds will be a 10% of the network,
5. The graph (network) will be small-world,
6. Will use the WS algorithm with $p = .2$ (probability of rewiring).
7. Threshold levels will be uniformly distributed between [0.3, 0.7]

To generate this diffusion network, we can use the rdiffnet function included in the package:

```
# Setting the seed for the RNG
set.seed(1213)

# Generating a random diffusion network
net <- rdiffnet(
  n           = 1e3,                      # 1.
  t           = 20,                       # 2.
  seed.nodes  = "random",                 # 3.
  seed.p.adopt = .1,                      # 4.
  seed.graph   = "small-world",            # 5.
  rgraph.args  = list(p=.2),                # 6.
  threshold.dist = function(x) runif(1, .3, .7) # 7.
)

## Warning in (function (graph, p, algorithm = "endpoints", both.ends = FALSE, :
```

```
## The option -copy.first- is set to TRUE. In this case, the first graph will be
## treated as a baseline, and thus, networks after T=1 will be replaced with T-1.
```

- The function `rdiffnet` generates random diffusion networks. Main features:
 1. Simulating random graph or using your own,
 2. Setting threshold levels per node,
 3. Network rewiring throughout the simulation, and
 4. Setting the seed nodes.
- The simulation algorithm is as follows:
 1. If required, a baseline graph is created,
 2. Set of initial adopters and threshold distribution are established,
 3. The set of t networks is created (if required), and
 4. Simulation starts at $t=2$, assigning adopters based on exposures and thresholds:
 - a. For each $i \in N$, if its exposure at $t-1$ is greater than its threshold, then adopts, otherwise, continue without change.
 - b. next i

12.3.2 Rumor spreading

```
library(netdiffuseR)

set.seed(09)
diffnet_rumor <- rdiffnet(
  n = 5e2,
  t = 5,
  seed.graph = "small-world",
  rgraph.args = list(k = 4, p = .3),
  seed.nodes = "random",
  seed.p.adopt = .05,
  rewire = TRUE,
  threshold.dist = function(i) 1L,
  exposure.args = list(normalized = FALSE)
)


## Warning in (function (graph, p, algorithm = "endpoints", both.ends = FALSE, :
## The option -copy.first- is set to TRUE. In this case, the first graph will be
## treated as a baseline, and thus, networks after T=1 will be replaced with T-1.

summary(diffnet_rumor)
```

```

## Diffusion network summary statistics
## Name      : A diffusion network
## Behavior  : Random contagion
## -----
##   Period Adopters Cum Adopt. (%) Hazard Rate Density Moran's I (sd)
## -----
##   1       25      25 (0.05)      -     0.01 -0.00 (0.00)
##   2       78      103 (0.21)     0.16    0.01  0.01 (0.00) ***
##   3      187      290 (0.58)     0.47    0.01  0.01 (0.00) ***
##   4      183      473 (0.95)     0.87    0.01  0.01 (0.00) ***
##   5       27      500 (1.00)     1.00    0.01      -
## -----
##   Left censoring : 0.05 (25)
##   Right centoring : 0.00 (0)
##   # of nodes      : 500
## 
## Moran's I was computed on contemporaneous autocorrelation using 1/geodesic
## values. Significane levels *** <= .01, ** <= .05, * <= .1.
plot_diffnet(diffnet_rumor, slices = c(1, 3, 5))


```



```
# We want to use igraph to compute layout
igdf <- diffnet_to_igraph(diffnet_rumor, slices=c(1,2))[[1]]
pos <- igraph::layout_with_drl(igdf)

plot_diffnet2(diffnet_rumor, vertex.size = dgr(diffnet_rumor)[,1], layout=pos)
```

Diffusion dynamics




12.3.3 Diffusion

```
set.seed(09)
diffnet_complex <- rdiffnet(
  seed.graph = diffnet_rumor$graph,
  seed.nodes = which(diffnet_rumor$toa == 1),
  rewire = FALSE,
  threshold.dist = function(i) rbeta(1, 3, 10),
  name = "Diffusion",
  behavior = "Some social behavior"
)

plot_adopters(diffnet_rumor, what = "cumadopt", include.legend = FALSE)
plot_adopters(diffnet_complex, bg="tomato", add=TRUE, what = "cumadopt")
legend("topleft", legend = c("Disease", "Complex"), col = c("lightblue", "tomato"),
       bty = "n", pch=19)
```

Adopters and Cumulative Adopters



12.3.4 Mentor Matching

```
# Finding mentors
mentors <- mentor_matching(diffnet_rumor, 25, lead.ties.method = "random")

# Simulating diffusion with these mentors
set.seed(09)
diffnet_mentored <- rdiffnet(
  seed.graph = diffnet_complex,
  seed.nodes = which(mentors$`1`$isleader),
  rewire = FALSE,
  threshold.dist = diffnet_complex[["real_threshold"]],
  name = "Diffusion using Mentors"
)

summary(diffnet_mentored)

## Diffusion network summary statistics
## Name      : Diffusion using Mentors
## Behavior : Random contagion
## -----
##  Period   Adopters   Cum Adopt. (%)   Hazard Rate   Density   Moran's I (sd)
##  -----   -----   -----   -----   -----   -----   -----
```

```

##      1      25      25 (0.05)      -      0.01 -0.00 (0.00)
##      2      92      117 (0.23)     0.19      0.01 0.01 (0.00) ***
##      3     152      269 (0.54)     0.40      0.01 0.01 (0.00) ***
##      4     150      419 (0.84)     0.65      0.01 0.01 (0.00) ***
##      5      73      492 (0.98)     0.90      0.01 -0.00 (0.00) **
## -----
##  Left censoring : 0.05 (25)
##  Right centoring : 0.02 (8)
##  # of nodes      : 500
##
##  Moran's I was computed on contemporaneous autocorrelation using 1/geodesic
##  values. Significane levels *** <= .01, ** <= .05, * <= .1.

cumulative_adopt_count(diffnet_complex)

```

```

##      1      2      3      4      5
## num  25.00 80.00 183.0000 338.0000000 470.0000000
## prop  0.05  0.16   0.3660   0.6760000   0.9400000
## rate  0.00  2.20   1.2875   0.8469945   0.3905325

```

```
cumulative_adopt_count(diffnet_mentored)
```

```

##      1      2      3      4      5
## num  25.00 117.000 269.000000 419.0000000 492.0000000
## prop  0.05  0.234   0.538000   0.8380000   0.9840000
## rate  0.00  3.680   1.299145   0.5576208   0.1742243

```


12.3.5 Example by changing threshold

```

# Simulating a scale-free homophilic network
set.seed(1231)
X <- rep(c(1,1,1,1,1,0,0,0,0,0), 50)
net <- rgraph_ba(t = 499, m=4, eta = X)

# Taking a look in igraph
ig <- igraph::graph_from_adjacency_matrix(net)
plot(ig, vertex.color = c("azure", "tomato")[X+1], vertex.label = NA,
     vertex.size = sqrt(dgr(net)))

```



```

# Now, simulating a bunch of diffusion processes
nsim <- 500L
ans_1and2 <- vector("list", nsim)
set.seed(223)
for (i in 1:nsim) {
  # We just want the cum adopt count
  ans_1and2[[i]] <-
    cumulative_adopt_count(
      rdiffnet(
        seed.graph = net,
        t = 10,
        threshold.dist = sample(1:2, 500L, TRUE),
        seed.nodes = "random",
        seed.p.adopt = .10,
        exposure.args = list(outgoing = FALSE, normalized = FALSE),
        rewire = FALSE
      )
    )

  # Are we there yet?
  if (!(i %% 50))
    message("Simulation ", i, " of ", nsim, " done.")
}
## Simulation 50 of 500 done.
## Simulation 100 of 500 done.
## Simulation 150 of 500 done.
## Simulation 200 of 500 done.
## Simulation 250 of 500 done.
## Simulation 300 of 500 done.
## Simulation 350 of 500 done.

```

```

## Simulation 400 of 500 done.
## Simulation 450 of 500 done.
## Simulation 500 of 500 done.

# Extracting prop
ans_1and2 <- do.call(rbind, lapply(ans_1and2, "[", i="prop", j=))

ans_2and3 <- vector("list", nsim)
set.seed(223)
for (i in 1:nsim) {
  # We just want the cum adopt count
  ans_2and3[[i]] <-
    cumulative_adopt_count(
      rdiffnet(
        seed.graph = net,
        t = 10,
        threshold.dist = sample(2:3, 500L, TRUE),
        seed.nodes = "random",
        seed.p.adopt = .10,
        exposure.args = list(outgoing = FALSE, normalized = FALSE),
        rewire = FALSE
      )
    )
}

# Are we there yet?
if (!(i %% 50))
  message("Simulation ", i, " of ", nsim, " done.")
}

## Simulation 50 of 500 done.
## Simulation 100 of 500 done.
## Simulation 150 of 500 done.
## Simulation 200 of 500 done.
## Simulation 250 of 500 done.
## Simulation 300 of 500 done.
## Simulation 350 of 500 done.
## Simulation 400 of 500 done.
## Simulation 450 of 500 done.
## Simulation 500 of 500 done.

ans_2and3 <- do.call(rbind, lapply(ans_2and3, "[", i="prop", j=))


```

We can simplify by using the function `rdiffnet_multiple`. The following lines of code accomplish the same as the previous code avoiding the for-loop (from the user's perspective).

Besides of the usual parameters passed to `rdiffnet`, the `rdiffnet_multiple` function requires `R` (number of repetitions/simulations), and `statistic` (a function that returns the statistic of interest). Optionally, the user may choose to specify the number of clusters to run it in parallel (multiple CPUs):

```
ans_1and3 <- rdiffnet_multiple(
  # Num of sim
  R           = nsim,
  # Statistic
  statistic    = function(d) cumulative_adopt_count(d)[ "prop", ],
  seed.graph   = net,
  t            = 10,
  threshold.dist = sample(1:3, 500, TRUE),
  seed.nodes   = "random",
  seed.p.adopt = .1,
  rewire       = FALSE,
  exposure.args = list(outgoing=FALSE, normalized=FALSE),
  # Running on 4 cores
  ncpus        = 4L
)

boxplot(ans_1and2, col="ivory", xlab = "Time", ylab = "Threshold")
boxplot(ans_2and3, col="tomato", add=TRUE)
boxplot(t(ans_1and3), col = "steelblue", add=TRUE)
legend(
  "topleft",
  fill = c("ivory", "tomato", "steelblue"),
  legend = c("1/2", "2/3", "1/3"),
  title = "Threshold range",
  bty = "n"
)
```



12.3.6 Problems

- Given the following types of networks: Small-world, Scale-free, Bernoulli, what set of n initiators maximizes diffusion?

12.4 Statistical inference

12.4.1 Moran's I


- Moran's I tests for spatial autocorrelation.
- netdiffuseR** implements the test in `moran`, which is suited for sparse matrices.
- We can use Moran's I as a first look to whether there is something happening: let that be influence or homophily.

12.4.2 Using geodesics

- One approach is to use the geodesic (shortest path length) matrix to account for indirect influence.
- In the case of sparse matrices, and furthermore, in the presence of structural holes it is more convenient to calculate the distance matrix taking this into account.
- netdiffuseR** has a function to do so, the `approx_geodesic` function, which, using graph powers, computes the shortest path up to n steps. This could be faster (if you only care up to n steps) than `igraph` or `sns`:

```
# Extracting the large adjacency matrix (stacked)
dgc <- diag_expand(medInnovationsDiffNet$graph)
ig <- igraph::graph_from_adjacency_matrix(dgc)
mat <- network::as.network(as.matrix(dgc))

# Measuring times
times <- microbenchmark::microbenchmark(
  netdiffuseR = netdiffuseR::approx_geodesic(dgc),
  igraph = igraph::distances(ig),
  sna = sna::geodist(mat),
  times = 50, unit="ms"
)
```



- The summary.diffnet method already runs Moran's for you. What happens under the hood is:

```
# For each time point we compute the geodesic distances matrix
W <- approx_geodesic(medInnovationsDiffNet$graph[[1]])

# We get the element-wise inverse
W@x <- 1/W@x

# And then compute moran
moran(medInnovationsDiffNet$cumadopt[,1], W)

## $observed
## [1] 0.06624028
```

```

## 
## $expected
## [1] -0.008064516
##
## $sd
## [1] 0.03265066
##
## $p.value
## [1] 0.02286087
##
## attr(,"class")
## [1] "diffnet_moran"

```

12.4.3 Structural dependence and permutation tests

- A novel statistical method (work-in-progress) that allows conducting inference.
- Included in the package, tests whether a particular network statistic depends on network structure
- Suitable to be applied to network thresholds (you can't use thresholds in regression-like models!)

12.4.4 Idea

- Let $\mathcal{G} = (V, E)$ be a graph, γ a vertex attribute, and $\beta = f(\gamma, \mathcal{G})$, then

$$\gamma \perp \mathcal{G} \implies \mathbb{E}[\beta(\gamma, \mathcal{G}) | \mathcal{G}] = \mathbb{E}[\beta(\gamma, \mathcal{G})]$$

- This is, if for example time of adoption is independent on the structure of the network, then the average threshold level will be independent from the network structure as well.
- Another way of looking at this is that the test will allow us to see how probable is to have this combination of network structure and network threshold (if it is uncommon then we say that the diffusion model is highly likely)

12.4.4.1 Example Not random TOA

- To use this test, `__netdiffuseR__` has the `struct_test` function.
- It simulates networks with the same density, and computes a particular statistic every

```

# Simulating network
set.seed(1123)
net <- rdiffnet(n=500, t=10, seed.graph = "small-world")

## Warning in (function (graph, p, algorithm = "endpoints", both.ends = FALSE, :
## The option -copy.first- is set to TRUE. In this case, the first graph will be

```

```

## treated as a baseline, and thus, networks after T=1 will be replaced with T-1.

# Running the test
test <- struct_test(
  graph      = net,
  statistic = function(x) mean(threshold(x), na.rm = TRUE),
  R          = 1e3,
  ncpus=4, parallel="multicore"
)


## Warning in (function (graph, p, algorithm = "endpoints", both.ends = FALSE, :
## The option -copy.first- is set to TRUE. In this case, the first graph will be
## treated as a baseline, and thus, networks after T=1 will be replaced with T-1.

# See the output
test

##
## Structure dependence test
## # Simulations      : 1,000
## # nodes           : 500
## # of time periods : 10
## -----
## H0: E[beta(Y,G)|G] - E[beta(Y,G)] = 0 (no structure dependency)
##      observed     expected    p.val
##      0.5513      0.2504    0.0000

```

Empirical Distribution of Statistic



- Now we shuffle times of adoption, so that is random

```
# Resetting TOAs (now will be completely random)
diffnet.toa(net) <- sample(diffnet.toa(net), nnodes(net), TRUE)
```

Running the test

```
test <- struct_test(
  graph      = net,
  statistic = function(x) mean(threshold(x), na.rm = TRUE),
  R          = 1e3,
  ncpus=4, parallel="multicore"
)
```

```
## Warning in (function (graph, p, algorithm = "endpoints", both.ends = FALSE, :
## The option -copy.first- is set to TRUE. In this case, the first graph will be
## treated as a baseline, and thus, networks after T=1 will be replaced with T-1.
```

See the output

```
test
```

```
##
```

```
## Structure dependence test
```

```
## # Simulations      : 1,000
```

```
## # nodes           : 500
```

```
## # of time periods : 10
```


```
## -----
```

```
## H0: E[beta(Y,G)|G] - E[beta(Y,G)] = 0 (no structure dependency)
```

```
##      observed    expected     p.val
```

```
##      0.2714      0.2590      0.4420
```

Empirical Distribution of Statistic



12.4.5 Regression analysis

- In regression analysis, we want to see if exposure, once we control for other covariates had any effect on the adoption of a behavior.
- In general, the big problem here is when we have a latent variable that co-determines both network and behavior.
- Unless we can control for such variable, regression analysis will be generically biased.
- On the other hand, if you can claim that either such variable doesn't exist or you actually can control for it, then we have two options: lagged exposure models or contemporaneous exposure models. We will focus on the former.

12.4.5.1 Lagged exposure models

- In this type of model, we usually have the following

$$y_t = f(W_{t-1}, y_{t-1}, X_i) + \varepsilon$$

Furthermore, in the case of adoption, we have

$$y_{it} = \begin{cases} 1 & \text{if } \rho \sum_{j \neq i} \frac{W_{ijt-1} y_{jt-1}}{\sum_{j \neq i} W_{ijt-1}} + X_{it}\beta > 0 \\ 0 & \text{otherwise} \end{cases}$$

- In netdiffuseR is as easy as doing the following:

```
# fakedata
set.seed(121)

W   <- rgraph_ws(1e3, 8, .2)
X   <- cbind(var1 = rnorm(1e3))
toa <- sample(c(NA,1:5), 1e3, TRUE)

dn  <- new_diffnet(W, toa=toa, vertex.static.attrs = X)

## Warning in new_diffnet(W, toa = toa, vertex.static.attrs = X): -graph- is
## static and will be recycled (see ?new_diffnet).

# Computing exposure and adoption for regression
dn[["cohesive_expo"]] <- cbind(NA, exposure(dn)[,-nslices(dn)])
dn[["adopt"]]        <- dn$cumadopt

# Generating the data and running the model
dat <- as.data.frame(dn)
ans <- glm(adopt ~ cohesive_expo + var1 + factor(per),
           data = dat,
           family = binomial(link="probit"),
           subset = is.na(toa) | (per <= toa))
summary(ans)

##
## Call:
## glm(formula = adopt ~ cohesive_expo + var1 + factor(per), family = binomial(link =
##           data = dat, subset = is.na(toa) | (per <= toa))
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.92777   0.05840 -15.888 < 2e-16 ***
## cohesive_expo  0.23839   0.17514   1.361 0.173452
## var1         -0.04623   0.02704  -1.710 0.087334 .
## factor(per)3  0.29313   0.07715   3.799 0.000145 ***
## factor(per)4  0.33902   0.09897   3.425 0.000614 ***
## factor(per)5  0.59851   0.12193   4.909 9.18e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2745.1  on 2317  degrees of freedom
```

```

## Residual deviance: 2663.5 on 2312 degrees of freedom
## (1000 observations deleted due to missingness)
## AIC: 2675.5
##
## Number of Fisher Scoring iterations: 4

```

Alternatively, we could have used the new function `diffreg`

```

ans <- diffreg(dn ~ exposure + var1 + factor(per), type = "probit")
summary(ans)

```

```

##
## Call:
## glm(formula = Adopt ~ exposure + var1 + factor(per), family = binomial(link = "prob
##       data = dat, subset = ifelse(is.na(toa), TRUE, toa >= per))
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.92777   0.05840 -15.888 < 2e-16 ***
## exposure     0.23839   0.17514   1.361 0.173452
## var1        -0.04623   0.02704  -1.710 0.087334 .
## factor(per)3 0.29313   0.07715   3.799 0.000145 ***
## factor(per)4 0.33902   0.09897   3.425 0.000614 ***
## factor(per)5 0.59851   0.12193   4.909 9.18e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2745.1 on 2317 degrees of freedom
## Residual deviance: 2663.5 on 2312 degrees of freedom
## (1000 observations deleted due to missingness)
## AIC: 2675.5
##
## Number of Fisher Scoring iterations: 4

```

12.4.5.2 Contemporaneous exposure models

- Similar to the lagged exposure models, we usually have the following

$$y_t = f(W_t, y_t, X_t) + \varepsilon$$

Furthermore, in the case of adoption, we have

$$y_{it} = \begin{cases} 1 & \text{if } \rho \sum_{j \neq i} \frac{W_{ijt} y_{jt}}{\sum_{j \neq i} W_{ijt}} + X_{it}\beta > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Unfortunately, since y_t is in both sides of the equation, this models cannot be fitted using a standard probit or logit regression.
- Two alternatives to solve this:
 - a. Using Instrumental Variables Probit (ivprobit in both R and Stata)
 - b. Use a Spatial Autoregressive (SAR) Probit (SpatialProbit and ProbitSpatial in R).
- We won't cover these here.

12.4.6 Problems

Using the dataset stats.rda:

1. Compute Moran's I as the function `summary.diffnet` does. For this you'll need to use the function `toa_mat` (which calculates the cumulative adoption matrix), and `approx_geodesic` (which computes the geodesic matrix). (see `?summary.diffnet` for more details).
2. Read the data as `diffnet` object, and fit the following logit model $\text{adopt} = \text{Exposure} * \gamma + \text{Measure} * \beta + \varepsilon$. What happens if you exclude the time-fixed effects?

Chapter 13

Stochastic Actor Oriented Models

Stochastic Actor Oriented Models (SOAM), also known as Siena models were introduced by CITATION NEEDED.

As a difference from ERGMs, Siena models look at the data generating process from the individuals' point of view. Based on McFadden's ideas of probabilistic choice, the model is founded in the following equation

$$U_i(x) - U_i(x') \sim \text{Extreme Value Distribution}$$

In other words, individuals choose between states x and x' in a probabilistic way (with some noise),

$$\frac{\exp\{f_i^Z(\beta^z, x, z)\}}{\sum_{z' \in C} \exp\{f_i^Z(\beta, x, z')\}}$$

snijders_(sociological methodology 2001)

Ripley et al. (2011)

Chapter 14

Power calculation in network studies

In survey and study design, calculating the required sample size is critical. Nowadays, tools and methods for calculating the required sample size abound; nonetheless, sample size calculation for studies involving social networks is still underdeveloped. This chapter will illustrate how we can use computer simulations to estimate the required sample size. Chapter 17 provides a general overview of power analysis.

14.1 Example 1: Spillover effects in egocentric studies¹

Suppose we want to run an intervention over a particular population, and we are interested in the effects of such intervention on the egos' alters. In economics, this problem, which they call the "spillover effect," is actively studied.

We assume that alters only get exposed if egos acquire the behavior for the power calculation. Furthermore, for this first run, we will assume that there is no social reinforcement or influence between alters. We will later relax this assumption. To calculate power, we will do the following:

1. Simulate egos' behavior following a logit distribution.
2. Randomly drop some egos as a result of attrition.
3. Simulate alters' behavior using their egos as the treatment.
4. Fit a logistic regression based on the previous model.
5. Accept/reject the null and store the result.

The previous steps will be repeated 500 for each value of n we analyze. We will finalize by plotting power against sample sizes. Let's first start by writing down the simulation parameters:

¹The original problem was posed by Dr. Shinduk Lee from the School of Nursing at the University of Utah.

```

# Design
n_sims    <- 500 # Number of simulations
n_a       <- 4   # Number of alters
sizes     <-      # Sizes to try
seq(from = 100, to = 200, by = 25)

# Assumptions
odds_h_1 <- 2.0 # Odds of Increase/
attrition <- .3
baseline  <- .2 # Low prevalence in 1s

# Parameters
alpha     <- .05
beta_pow <- 0.2

```

As we discuss in 17, it is always a good idea to encapsulate the simulation into a function:

```

# The odds turned to a prob
theta_h_1 <- plogis(log(odds_h_1))

# Simulation function
sim_data <- function(n) {

  # Treatment assignment
  tr <- c(rep(1, n/2), rep(0, n/2))

  # Step 1: Sampling population of egos
  y_ego <- runif(n) < c(
    rep(theta_h_1, n/2),
    rep(0.5, n/2)
  )

  # Step 2: Simulating attrition
  todrop <- order(runif(n))[1:(n * attrition)]
  y_ego <- y_ego[-todrop]
  tr     <- tr[-todrop]
  n      <- n - length(todrop)

  # Step 3: Simulating alter's effect. We assume the same as in
  # ego
  tr.Alter <- rep(y_ego * tr, n_a)
  y.Alter  <- runif(n * n_a) < ifelse(tr.Alter, theta_h_1, 0.5)

  # Step 4: Computing test statistic

```

```

res_ego <- tryCatch(glm(y_ego ~ tr, family = binomial("logit")), error = function(e) e)
res.Alter <- tryCatch(glm(y.Alter ~ tr.Alter, family = binomial("logit")), error = function(e) e)

if (inherits(res_ego, "error") | inherits(res.Alter, "error"))
  return(c(ego = NA, alter = NA))

# Step 5: Reject?
c(
  ego = summary(res_ego)$coefficients["tr", "Pr(>|z|)"] < alpha,
  alter = summary(res.Alter)$coefficients["tr.Alter", "Pr(>|z|)"] < alpha
)

}

```

Now that we have the data generating function, we can run the simulations to approximate statistical power given the sample size. The results will be stored in the matrix `spower`. Since we are simulating data, it is crucial to set the seed so we can reproduce the results.

```

# We always set the seed
set.seed(88)

# Making space, and running!
spower <- NULL
for (s in sizes) {

  # Run the simulation for size s
  simres <- rowMeans(replicate(n_sims, sim_data(s)), na.rm = TRUE)

  # And store the results
  spower <- rbind(spower, simres)

}

```

The following figure shows the approximate power for finding effects at both levels, ego and alter:

```


library(ggplot2)

spower <- rbind(
  data.frame(size = sizes, power = spower[, "ego"], type = "ego"),
  data.frame(size = sizes, power = spower[, "alter"], type = "alter")
)

spower |>

```

```
ggplot(aes(x = size, y = power, colour = type)) +
  geom_point() +
  geom_smooth(method = "loess", formula = y ~ x) +
  labs(x = "Number of Egos", y = "Approx. Power", colour = "Node type") +
  geom_hline(yintercept = 1 - beta_pow)
```



As shown in Chapter 17, we can use a linear regression model to predict sample size as a function of statistical power:

```
# Fitting the model
power_model <- glm(
  size ~ power + I(power^2),
  data = spower, family = gaussian(), subset = type == "alter"
)

summary(power_model)

## 
## Call:
## glm(formula = size ~ power + I(power^2), family = gaussian(),
##      data = spower, subset = type == "alter")
## 
```

```

## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1460     1342   1.088   0.390
## power        -3532     3124  -1.131   0.376
## I(power^2)    2293     1805   1.270   0.332
##
## (Dispersion parameter for gaussian family taken to be 317.0856)
##
## Null deviance: 6250.00 on 4 degrees of freedom
## Residual deviance: 634.17 on 2 degrees of freedom
## AIC: 46.404
##
## Number of Fisher Scoring iterations: 2

# Predict
predict(power_model, newdata = data.frame(power = .8), type = "response") |>
  ceiling()

##    1
## 102

```

From the figure, it becomes apparent that, although there is not enough power to identify effects at the ego level, because each ego brings in five alters, the alter sample size is high enough that we can reach above 0.8 statistical power with relatively small sample size.

14.2 Example 2: Spillover effects pre-post effect

Now the dynamics are different. Instead of having a treated and control group, we have a single group over which we will measure behavioral change. We will simulate individuals in their initial state, still 0/1, and then simulate that the intervention will make them more likely to have $y = 1$. We will also assume that subjects generally don't change their behavior and that the baseline prevalence of zeros is higher. The simulation steps are as follows:

1. For each individual in the population, draw the underlying probability that $y = 1$. With that probability, assign the value of y . This applies to both ego and alter.
2. Randomly drop some egos, and their corresponding alters due to attrition.
3. Simulate alters' behavior using their egos as the treatment. Both ego and alter's underlying probability are increased by the chosen odds.
4. To control for the underlying probability that an individual has $y = 1$, we use conditional logistic regression (also known as matched case-control logit,) to estimate the treatment effects.
5. Accept/reject the null and store the result.

```

beta_pars <- c(4, 6)
odds_h_1 <- 2.0

# Simulation function
library(survival)
sim_data_repost <- function(n) {

  # Step 1: Sampling population of egos
  y_ego_star <- rbeta(n, beta_pars[1], beta_pars[2])
  y_ego_0     <- runif(n) < y_ego_star

  # Step 2: Simulating attrition
  todrop      <- order(runif(n))[1:(n * attrition)]
  y_ego_0     <- y_ego_0[-todrop]
  n           <- n - length(todrop)
  y_ego_star <- y_ego_star[-todrop]

  # Step 3: Simulating alter's effect. We assume the same as in
  # ego
  y_alter_star <- rbeta(n * n_a, beta_pars[1], beta_pars[2])
  y_alter_0     <- runif(n * n_a) < y_alter_star

  # Simulating post
  y_ego_1     <- runif(n) < plogis(qlogis(y_ego_star) + log(odds_h_1))
  tr_alter    <- as.integer(rep(y_ego_1, n_a))
  y_alter_1 <- runif(n * n_a) < plogis(qlogis(y_alter_star) + log(odds_h_1) * tr_alter) #

  # Step 4: Computing test statistic
  y_ego_0 <- as.integer(y_ego_0)
  y_ego_1 <- as.integer(y_ego_1)
  y_alter_0 <- as.integer(y_alter_0)
  y_alter_1 <- as.integer(y_alter_1)

  d <- data.frame(
    y = c(y_ego_0, y_ego_1),
    tr = c(rep(0, n), rep(1, n)),
    g = c(1:n, 1:n)
  )

  res_ego   <- tryCatch(
    clogit(y ~ tr + strata(g), data = d, method = "exact"),
    error = function(e) e
  )
}

```

```

        )

d <- data.frame(
  y  = c(y_alter_0, y_alter_1),
  tr = c(rep(0, n * n_a), tr_alter),
  g  = c(1:(n * n_a), 1:(n * n_a))
)

res_alter <- tryCatch(
  clogit(y ~ tr + strata(g), data = d, method = "exact"),
  error = function(e) e
)

if (inherits(res_ego, "error") | inherits(res_alter, "error"))
  return(c(ego = NA, alter = NA))

# Step 5: Reject?
c(
  # ego      = res_ego$p.value < alpha,
  ego      = summary(res_ego)$coefficients["tr", "Pr(>|z|)"] < alpha,
  alter     = summary(res_alter)$coefficients["tr", "Pr(>|z|)"] < alpha,
  ego_test  = coef(res_ego),
  alter_glm  = coef(res_alter)
)
}

# We always set the seed
set.seed(88)

# Making space and running!
spower <- NULL
for (s in sizes) {

  # Run the simulation for size s
  simres <- rowMeans(
    replicate(n_sims, sim_data_prepot(s)),
    na.rm = TRUE
  )

  # And store the results
  spower <- rbind(spower, simres)
}


```

```
}
```

```
library(ggplot2)

spowerd <- rbind(
  data.frame(size = sizes, power = spower[, "ego"], type = "ego"),
  data.frame(size = sizes, power = spower[, "alter"], type = "alter")
)

spowerd |>
  ggplot(aes(x = size, y = power, colour = type)) +
  geom_point() +
  geom_smooth(method = "loess", formula = y ~ x) +
  labs(x = "Number of Egos", y = "Approx. Power", colour = "Node type") +
  geom_hline(yintercept = 1 - beta_pow)
```



As shown in Chapter 17, we can use a linear regression model to predict sample size as a function of statistical power:

```
# Fitting the model
power_model <- glm(
  size ~ power + I(power^2),
```

```

  data = spowerd, family = gaussian(), subset = type == "alter"
)

summary(power_model)

##
## Call:
## glm(formula = size ~ power + I(power^2), family = gaussian(),
##      data = spowerd, subset = type == "alter")
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  611.4     666.3   0.918   0.456
## power       -1553.8    1504.7  -1.033   0.410
## I(power^2)   1147.9     844.8   1.359   0.307
##
## (Dispersion parameter for gaussian family taken to be 52.1536)
##
## Null deviance: 6250.00 on 4 degrees of freedom
## Residual deviance: 104.31 on 2 degrees of freedom
## AIC: 37.379
##
## Number of Fisher Scoring iterations: 2

# Predict
predict(power_model, newdata = data.frame(power = .8), type = "response") |>
  ceiling()

##    1
## 104

```

14.3 Example 3: First difference

Now, instead of looking at a dichotomous outcome, let's evaluate what happens if the variable is continuous. The effects we are interested to identify are the ego and alter effect, γ_{ego} and γ_{alter} , respectively. Furthermore, the data generating process is

$$\begin{aligned}
 Y_{itg} &= \alpha_i + \kappa_g + X_i\beta + \varepsilon_{itg} \\
 Y_{itg} &= \alpha_i + \kappa_g + X_i\beta + D_i^{ego}\gamma_{ego} + D_i^{alter}\gamma_{alter} + \varepsilon_{itg}
 \end{aligned}$$

Where $D_i^{ego/alter}$ is an indicator variable. Here, ego and alter behavior are correlated through a fixed effect. In other words, within each group, we are assuming that there's a shared baseline

prevalence of the outcome. The main difference is that ego and alter may have different results regarding the effect size of the treatment. Another way of approaching the group-level correlation could be through an autocorrelation process, like in a spatial Autocorrelated model; nonetheless, estimating such models is computationally expensive, so we opted to use the former.

For simplicity, we assume that there is no time effect. Two essential components here, α_i and κ_g are individual and group-level unobserved fixed effects. The most straightforward approach here is to use a first difference estimator:

$$(y_{it+1g} - y_{itg}) = D_i^{ego} \gamma_{ego} + D_i^{alter} \gamma_{alter} + \varepsilon'_i, \quad \varepsilon'_i = \varepsilon_{it+1g} - \varepsilon_{itg}$$

By taking the first difference, the fixed effects are removed from the equation, and we can proceed with a regular linear model.

```
effect_size_ego <- 0.5
effect_size_alter <- 0.25
sizes <- seq(10, 100, by = 10)

# Simulation function
sim_data_prep <- function(n) {

  # Applying attrition
  n <- floor(n * (1 - attrition))

  # Step 1: Sampling fixed effects
  alpha_i <- rnorm(n * (n_a + 1))
  kappa_g <- rep(rnorm(n_a + 1), n)

  # Step 2: Generating the outcome at t = 1
  is_ego <- rep(c(1, rep(0, n_a)), n)
  is_alter <- 1 - is_ego
  y_0 <- alpha_i + kappa_g + rnorm(n * (n_a + 1))
  y_1 <- alpha_i + kappa_g +
    is_ego * effect_size_ego +
    is_alter * effect_size_alter +
    rnorm(n * (n_a + 1))

  # Step 4: Computing test statistic
  res <- tryCatch(
    glm(I(y_1 - y_0) ~ -1 + is_ego + is_alter, family = gaussian("identity")),
    error = function(e) e
  )
}
```

```

if (inherits(res, "error"))
  return(c(ego = NA, alter = NA))

# Step 5: Reject?
c(
  # ego      = res_ego$p.value < alpha,
  ego      = summary(res)$coefficients["is_ego", "Pr(>|t|)"] < alpha,
  alter     = summary(res)$coefficients["is_alter", "Pr(>|t|)"] < alpha,
  coef(res)[1],
  coef(res)[2]
)

}

# We always set the seed
set.seed(88)

# Making space and running!
spower <- NULL
for (s in sizes) {

  # Run the simulation for size s
  simres <- rowMeans(
    replicate(n_sims, sim_data_repost(s)),
    na.rm = TRUE
  )

  # And store the results
  spower <- rbind(spower, simres)
}

library(ggplot2)

spowerd <- rbind(
  data.frame(size = sizes, power = spower[, "ego"], type = "ego"),
  data.frame(size = sizes, power = spower[, "alter"], type = "alter")
)


spowerd |>
  ggplot(aes(x = size, y = power, colour = type)) +
  geom_point() +

```

```

geom_smooth(method = "loess", formula = y ~ x) +
  labs(x = "Number of Egos", y = "Approx. Power", colour = "Node type") +
  geom_hline(yintercept = 1 - beta_pow) +
  labs(
    caption = sprintf(
      "Ego effect: %.2f; Alter effect: %.2f", effect_size_ego, effect_size.Alter
    )
  )

```



From the inferential point of view, we could still use a demean operator to estimate individual-level effects. In particular, we would require to use the demean operator at the group level and then fit a fixed effect model to estimate individual-level parameters.

Part II

Statistical Foundations

Chapter 15

Bayes' Rule

Bayes' Rule is a fundamental equation in Bayesian statistics. With it, we can reformulate inferential problems by writing probabilities in terms of known quantities. Bayes' rule can be stated as follows:

$$\mathbb{P}(X = x|Y = y) = \frac{\mathbb{P}(Y = y|X = y)\mathbb{P}(X = x)}{\mathbb{P}(Y = y)} \quad (15.1)$$

Here we say that the conditional probability of X given Y can be expressed in terms of the conditional probability of Y given X . For example, let X be an unknown vector of parameters $\theta \in \Theta$ and Y a dataset $D \sim f(\theta)$ whose data generating process depends on the unobserved θ . As the posterior distribution of model parameters is in general elusive, instead, we use Bayes' rule to rephrase the problem:

$$\mathbb{P}(\theta|D) = \frac{\mathbb{P}(\theta|D)\mathbb{P}(\theta)}{\mathbb{P}(D)}$$

Since the denominator of the equation does not depend on θ , we can, instead, write

$$\mathbb{P}(\theta|D) \propto \mathbb{P}(\theta|D)\mathbb{P}(\theta)$$

In the Bayesian world, the unconditional distribution of the model parameters is assumed to come from a particular distribution, whereas in the frequentist world, no distributional assumptions are made about the model parameters. The latter is then equivalent to saying that $\theta \sim \text{Uniform}(-\infty, +\infty)$; therefore, even frequentists assume something about the model parameters!¹

Bayes' rule can be derived using conditional probabilities. In particular, $\mathbb{P}(x = x|Y = y)$ is defined as $\mathbb{P}(x = x, Y = y)/\Pr(Y = y)$. Likewise, $\mathbb{P}(y = y|X = x)$ is defined as $\mathbb{P}(y = y, X = x)/\Pr(X = x)$,

¹The discussion about differences and similarities between frequentists and Bayesians has a long tradition. Bottom line, no one can say 100% that they are either-or. In rigor, frequentists say model parameters are not random but deterministic.

which can be re-written as $\mathbb{P}(x = x, Y = y) = \mathbb{P}(y = y|X = x)Pr(X = x)$. Replacing the last equality in the first equation we get

$$\mathbb{P}(x = x|Y = y) = \frac{\mathbb{P}(x = x, Y = y)}{\mathbb{P}(Y = y)}$$
$$= \frac{\mathbb{P}(y = y|X = x)Pr(X = x)}{\mathbb{P}(Y = y)}$$

Chapter 16

Markov Chain

A Markov Chain is a sequence of random variables in which the conditional distribution of the n -th element only depends on $n - 1$.

16.1 Metropolis Algorithm

In the Metropolis Algorithm, or Metropolis MCMC, builds a Markov Chain that under certain conditions converges to the target distribution. The key of the Algorithm is in accepting a proposed move from θ to θ' with probability equal to:

$$r = \min\left(1, \frac{\mathbb{P}(\theta'|D)}{\mathbb{P}(\theta|D)}\right) \quad (16.1)$$

The resulting sequence converges to the target distribution. We can prove convergence by showing that (a) the sequence is ergodic, and (b) the posterior distribution matches the target distribution. Ergodicity describes three properties of a chain:

- Irreducibility: There is no zero probability of transitioning between any pair of states.
- Aperiodicity: As the term suggests, the chain has no repetitive periods/sequences.
- Non-transient: Transient refers to a chain having non-zero probability of never returning to a starting state.

The three properties are reached by any random walk based on a well-defined probability distribution, so we will focus on showing that the posterior matches the target distribution.

16.2 Metropolis-Hastings

$$\min\left(1, \frac{\mathbb{P}(d|\theta')\mathbb{P}(\theta')\mathbb{P}(\theta'|\theta)}{\mathbb{P}(d|\theta)\mathbb{P}(\theta)\mathbb{P}(\theta|\theta')}\right)$$

If the transition probability is symmetric, then the previous equation reduces to the Metropolis probability.

16.3 Likelihood-free MCMC

1. Initialize the algorithm with θ_0 , $\theta^* = \theta_0$ —the current accepted state,—and observed summary statistic $s_0 = S(D_{observed})$:
2. For $t = 1$ to T do:
 - a. Draw θ_t from the proposal distribution $J(\theta_t|\theta^*)$
 - b. Draw a simulated data D_t from model $M(\theta_t)$
 - c. Calculate the summary statistics $s_t = S(D_t)$
 - d. Accept the proposed state with probability
If accepted, set $\theta^* = \theta_t$.
 - e. Next t

Chapter 17

Power and sample size

Computing power and sample size are common tasks in study design. This chapter will walk you through power analysis for network studies. First, we will start with some preliminaries regarding error types and statistical power.

17.1 Error types

One of the most important tables we'll see around is the contingency table of accept/reject the null hypothesis conditional on the true state:

	Accept H0	Reject H0
H0 is true	True positive	False negative
H1 is true	False positive	True negative

A better way, more statistically accurate version of this table would be

	Accept H0	Reject H0
H0 is true	Correct inference	Type I error
H1 is true	Type II error	Correct Inference

With $\mathbb{P}(\text{Type I error}) = \alpha$ and $\mathbb{P}(\text{Type II error}) = \beta$. This way, power can be defined as the probability of rejecting the null given the alternative is true, $\mathbb{P}(\text{Reject H0}|H1 \text{ is true}) = 1 - \beta$.

17.2 Example 1: Sample size for a proportion

Let's imagine we are preparing a study in which we would like to estimate the proportion of individuals with a given status. Formally, we then say that the variable $Y \sim \text{Bernoulli}(p)$. To do so, we will need to survey n individuals and estimate such a number by taking the sample average. Furthermore, we hypothesize that under the null the proportion is $H_0 : p = p_0$.

The key here is to think about a simple rejection rule. Again, power is the probability of **rejecting the null** given that **the alternative is true**. So, to write down the equation, we

need to think about acceptance and rejection regions. Let \hat{p} be our estimate for the population parameter, furthermore, $\hat{p} = n^{-1} \sum_i y_i$. Our test statistic can be—and will be, most of the cases—standardized to leverage the law of large numbers; under the null, we write the following:

$$\begin{aligned}\mathbb{E}(\hat{p}) &= p_0 \\ \text{Var}(\hat{p}) &= \sqrt{p_0(1-p_0)/n}\end{aligned}$$

Therefore, the statistic:

$$\frac{\hat{p} - p_0}{\sqrt{p_0(1-p_0)/n}} = \frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_0(1-p_0)}} \sim N(0, 1)$$

Since the statistic is normally distributed, we can then say when we will reject the null. For this case, that depends on the critical value, which most of the time is defined in terms of the type I error rate. Formally, we reject the null if

$$\frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_0(1-p_0)}} > Z_{1-\alpha/2}$$

This is equivalent to saying that the **test statistic fell into the rejection region**. With this in hand, we can now write out the equation that we will be using for calculating the sample size. Going back to the definition of power:

$$\begin{aligned}\mathbb{P}(\text{Reject } H_0 | H_1 \text{ is true}) &= 1 - \beta \\ \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_0(1-p_0)}} > Z_{1-\alpha/2} \mid p = p_1\right) &= 1 - \beta\end{aligned}$$

Observe that we cannot compute the power for all $p \neq p_0$; instead, we look at a given parameter value. A good idea is to start from one previously known or identified in other studies. The key idea here is to be able to manipulate the argument of the probability to turn it into a known distribution, for example, the normal distribution:

For a given Type I of 0.05 and power of 0.8, the required sample size can be computed as follows:

$$\begin{aligned}1 - \beta &= \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_0(1-p_0)}} > Z_{1-\alpha/2} \mid p = p_1\right) \\ &= \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_0(1-p_0)}} < Z_{\alpha/2} \mid p = p_1\right)\end{aligned}$$

$$\begin{aligned}
&= \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_0)}{\sqrt{p_1(1-p_1)}} < \frac{Z_{\alpha/2} \sqrt{p_0(1-p_0)}}{\sqrt{p_1(1-p_1)}} \mid p = p_1\right) \\
&= \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_0 + p_0 - p_1)}{\sqrt{p_1(1-p_1)}} < \frac{Z_{\alpha/2} \sqrt{p_0(1-p_0)} + \sqrt{n}(p_0 - p_1)}{\sqrt{p_1(1-p_1)}} \mid p = p_1\right) \\
&= \mathbb{P}\left(\frac{\sqrt{n}(\hat{p} - p_1)}{\sqrt{p_1(1-p_1)}} < \frac{Z_{\alpha/2} \sqrt{p_0(1-p_0)} + \sqrt{n}(p_0 - p_1)}{\sqrt{p_1(1-p_1)}} \mid p = p_1\right) \\
&= \Phi\left(\frac{Z_{\alpha/2} \sqrt{p_0(1-p_0)} + \sqrt{n}(p_0 - p_1)}{\sqrt{p_1(1-p_1)}} \mid p = p_1\right)
\end{aligned}$$

The last equality follows from the quantity $\frac{\sqrt{n}(\hat{p} - p_1)}{\sqrt{p_1(1-p_1)}}$ distributing standard normal. We can now take the inverse of the cumulative distribution function (cdf) to isolate the sample size n :

$$\begin{aligned}
\Phi^{-1}(1-\beta) &= \frac{Z_{\alpha/2} \sqrt{p_0(1-p_0)} + \sqrt{n}(p_0 - p_1)}{\sqrt{p_1(1-p_1)}} \\
Z_{1-\beta} \sqrt{p_1(1-p_1)} &= Z_{\alpha/2} \sqrt{p_0(1-p_0)} + \sqrt{n}(p_0 - p_1) \\
\frac{(Z_{1-\beta} \sqrt{p_1(1-p_1)} - Z_{\alpha/2} \sqrt{p_0(1-p_0)})^2}{(p_0 - p_1)^2} &= n
\end{aligned}$$

Therefore, for the parameters $(1-\beta, \alpha, p_0, p_1) = (0.8, 0.05, 0.5, 0.6)$, the required sample size is $193.8473 \sim 194$.

17.3 Example 2: Sample size for a proportion (vis)

Now, what happens if the model we are planning to estimate does not have a close form? If analytical solutions are not available, simulations can be an excellent alternative to save the day. Let's re-do the sample size calculation using simulations.

The procedure to compute sample size based on simulations is computationally intensive. The concept is straightforward, pick a set of best guesses for sample size, and for each one of them, simulate the system to estimate power. Now, for a given value of n , we:

1. Simulate a sample of size n under the alternative.
2. Compute the test statistic corresponding to the null.
3. Accept or reject accordingly to the selected α , and store the result.
4. Repeat steps 1-3 many times. The obtained average is the corresponding power.

When running simulations, it is convenient to write a function for the data generating process.

In our case, the function will be called `sim_fun`. The following lines of code achieve our goal: approximate power by simulating 10,000 experiments for each sample size candidate:

```
# Model parameters
p0      <- .5
p1      <- .6
betapower <- 1 - 0.8
alpha    <- 0.05
nsims   <- 10000

# Step 1: Simulate the data under H1
z_one_minus_alpha_half <- qnorm(1 - alpha / 2)
sim_fun <- function(n) {

  # Generating the data
  y <- as.integer(runif(n) < p1)
  phat <- mean(y)

  # Accept or reject?
  sqrt(n) * (phat - p0) / sqrt(p0 * (1 - p0)) >
    z_one_minus_alpha_half

}

# Step 2: For an array of n, simulate multiple experiments
n_seq <- seq(from = 150, to = 250, by = 10)

simulations <- NULL
set.seed(12312)
for (n in n_seq) {

  # Run the nsims experiments
  res <- replicate(nsims, sim_fun(n))

  # Compute power and store the value
  simulations <- rbind(
    simulations,
    data.frame(size = n, power = mean(res))
  )
}

# Finding out what is the closets value
best <- which.min(
  abs((1 - betapower) - simulations$power)
```


```
)
simulations[best,,drop=FALSE]

##   size  power
## 5  190 0.7952
```

Let's visualize the power curve we generate from this simulation:

```
library(ggplot2)
ggplot(simulations, aes(x = size, y = power)) +
  geom_point() +
  geom_smooth() +
  geom_hline(yintercept = 1 - betapower)

## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



Alternatively, we can fit a linear regression model where we predict power as a function of sample size using linear and quadratic effects:

$$n = \theta_0 + \theta_1(1 - \beta) + \theta_2(1 - \beta)^2$$

```

# Fitting the model
power_model <- glm(
  size ~ power + I(power^2),
  data = simulations, family = gaussian()
)

# Printing the results
summary(power_model)

## 
## Call:
## glm(formula = size ~ power + I(power^2), family = gaussian(),
##      data = simulations)
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   632.5     232.9   2.715  0.02644 *
## power        -1590.3    598.1  -2.659  0.02885 *
## I(power^2)    1301.0    381.6   3.410  0.00923 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 34.83159)
##
## Null deviance: 11000.00 on 10 degrees of freedom
## Residual deviance: 278.65 on 8 degrees of freedom
## AIC: 74.769
##
## Number of Fisher Scoring iterations: 2

# Predict
predict(power_model, newdata = data.frame(power = .8), type = "response") |>
  ceiling()

##    1
## 193

```

According to our simulation study, the closest to our 80% power is using a sample size equal to 193, which is very close to the analytical solution of 194.

As a final comment for this example, remember that the more simulations the better.

Appendix A

Datasets

A.1 SNS data

A.1.1 About the data

- This data is part of the NIH Challenge grant # RC 1RC1AA019239 “Social Networks and Networking That Puts Adolescents at High Risk”.
- In general terms, the SNS’s goal was(is) “Understand the network effects on risk behaviors such as smoking initiation and substance use”.

A.1.2 Variables

The data has a *wide* structure, which means that there is one row per individual, and that dynamic attributes are represented as one column per time.

- photoid Photo id at the school level (can be repeated across schools).
- school School id.
- hispanic Indicator variable that equals 1 if the individual ever reported himself as hispanic.
- female1, ..., female4 Indicator variable that equals 1 if the individual reported to be female at the particular wave.
- grades1, ..., grades4 Academic grades by wave. Values from 1 to 5, with 5 been the best.
- eversmk1, ..., eversmk4 Indicator variable of ever smoking by wave. A one indicated that the individual had smoked at the time of the survey.
- everdrk1, ..., everdrk4 Indicator variable of ever drinking by wave. A one indicated that the individual had drink at the time of the survey.
- home1, ..., home4 Factor variable for home status by wave. A one indicates home ownership, a 2 rent, and a 3 a “I don’t know”.

During the survey, participants were asked to name up to 19 of their school friends:

- sch_friend11, ..., sch_friend119 School friends nominations (19 in total) for wave 1.
The codes are mapped to the variable photoid.
- sch_friend21, ..., sch_friend219 School friends nominations (19 in total) for wave 2.
The codes are mapped to the variable photoid.
- sch_friend31, ..., sch_friend319 School friends nominations (19 in total) for wave 3.
The codes are mapped to the variable photoid.
- sch_friend41, ..., sch_friend419 School friends nominations (19 in total) for wave 4.
The codes are mapped to the variable photoid.

References

- Admiraal, Ryan, and Mark S Handcock. 2006. "Sequential Importance Sampling for Bipartite Graphs with Applications to Likelihood-Based Inference." Department of Statistics, University of Washington.
- Bache, Stefan Milton, and Hadley Wickham. 2014. *Magrittr: A Forward-Pipe Operator for r*. <https://CRAN.R-project.org/package=magrittr>.
- Bojanowski, Michal. 2015. *Intergraph: Coercion Routines for Network Data Objects*. <http://mbojan.github.io/intergraph>.
- Brooks, Steve, Andrew Gelman, Galin Jones, and Xiao-Li Meng. 2011. *Handbook of Markov Chain Monte Carlo*. CRC press.
- Csardi, Gabor, and Tamas Nepusz. 2006. "The Igraph Software Package for Complex Network Research." *InterJournal Complex Systems*: 1695. <http://igraph.org>.
- Efron, Bradley, and Robert J Tibshirani. 1994. *An Introduction to the Bootstrap*. CRC press.
- Geyer, Charles J., and Elizabeth A. Thompson. 1992. "Constrained Monte Carlo Maximum Likelihood for Dependent Data." *Journal of the Royal Statistical Society. Series B (Methodological)* 54 (3): 657–99. <http://www.jstor.org/stable/2345852>.
- Handcock, Mark S., David R. Hunter, Carter T. Butts, Steven M. Goodreau, Pavel N. Krivitsky, Skye Bender-deMoll, and Martina Morris. 2016. *Statnet: Software Tools for the Statistical Analysis of Network Data*. The Statnet Project (<http://www.statnet.org>). CRAN.R-project.org/package=statnet.
- Handcock, Mark S., David R. Hunter, Carter T. Butts, Steven M. Goodreau, Pavel N. Krivitsky, and Martina Morris. 2017. *Ergm: Fit, Simulate and Diagnose Exponential-Family Models for Networks*. The Statnet Project (<http://www.statnet.org>). <https://CRAN.R-project.org/package=ergm>.
- Hunter, David R, Steven M Goodreau, and Mark S Handcock. 2008. "Goodness of Fit of Social Network Models." *Journal of the American Statistical Association* 103 (481): 248–58. <https://doi.org/10.1198/016214507000000446>.
- Hunter, David R., Mark S. Handcock, Carter T. Butts, Steven M. Goodreau, and Martina Morris. 2008. "ergm : A Package to Fit, Simulate and Diagnose Exponential-Family Models for Networks." *Journal of Statistical Software* 24 (3). <https://doi.org/10.18637/jss.v024.i03>.
- Lazega, Emmanuel, and Tom AB Snijders. 2015. *Multilevel Network Analysis for the Social Sciences: Theory, Methods and Applications*. Vol. 12. Springer.
- Leifeld, Philip. 2013. "texreg: Conversion of Statistical Model Output in R to LaTeX and HTML

- Tables." *Journal of Statistical Software* 55 (8): 1–24. <http://www.jstatsoft.org/v55/i08/>.
- Lusher, Dean, Johan Koskinen, and Garry Robins. 2012. *Exponential Random Graph Models for Social Networks: Theory, Methods, and Applications*. Cambridge University Press.
- Matloff, Norman. 2011. *The Art of r Programming: A Tour of Statistical Software Design*. No Starch Press.
- Morris, Martina, Mark Handcock, and David Hunter. 2008. "Specification of Exponential-Family Random Graph Models: Terms and Computational Aspects." *Journal of Statistical Software, Articles* 24 (4): 1–24. <https://doi.org/10.18637/jss.v024.i04>.
- Plummer, Martyn, Nicky Best, Kate Cowles, and Karen Vines. 2006. "CODA: Convergence Diagnosis and Output Analysis for MCMC." *R News* 6 (1): 7–11. <https://journal.r-project.org/archive/>.
- R Core Team. 2017a. *Foreign: Read Data Stored by 'Minitab', 's', 'SAS', 'SPSS', 'Stata', 'Systat', 'Weka', 'dBase', ...* <https://CRAN.R-project.org/package=foreign>.
- . 2017b. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Ripley, Ruth M., Tom AB Snijders, Paulina Preciado, and Others. 2011. "Manual for RSiena." *University of Oxford: Department of Statistics, Nuffield College*, no. 2007. https://www.uni-due.de/hummell/sna/R/RSiena%7B/_%7DManual.pdf.
- Sarkar, Deepayan, and Felix Andrews. 2016. *latticeExtra: Extra Graphical Utilities Based on Lattice*. <https://CRAN.R-project.org/package=latticeExtra>.
- Snijders, Tom A B, and Stephen P Borgatti. 1999. "Non-Parametric Standard Errors and Tests for Network Statistics." *Connections* 22 (2): 1–10. https://insna.org/PDF/Connections/v22/1999_I-2_61-70.pdf.
- Snijders, Tom A B, Gerhard G. van de Bunt, and Christian E G Steglich. 2010. "Introduction to stochastic actor-based models for network dynamics." *Social Networks* 32 (1): 44–60. <https://doi.org/10.1016/j.socnet.2009.02.004>.
- SNIJDERS, TOM A. B. 2010. "Conditional Marginalization for Exponential Random Graph Models." *The Journal of Mathematical Sociology* 34 (4): 239–52. <https://doi.org/10.1080/0022250X.2010.485707>.
- Snijders, Tom AB. 2002. "Markov Chain Monte Carlo Estimation of Exponential Random Graph Models." *Journal of Social Structure* 3.
- Ushey, Kevin, Jim Hester, and Robert Krzyzanowski. 2017. *Rex: Friendly Regular Expressions*. <https://CRAN.R-project.org/package=rex>.
- Wang, Peng, Ken Sharpe, Garry L. Robins, and Philippa E. Pattison. 2009. "Exponential Random Graph (p*) Models for Affiliation Networks." *Social Networks* 31 (1): 12–25. <https://doi.org/10.1016/j.socnet.2008.08.002>.
- Wickham, Hadley. 2017. *Stringr: Simple, Consistent Wrappers for Common String Operations*. <https://CRAN.R-project.org/package=stringr>.
- Wickham, Hadley, and Jennifer Bryan. 2017. *Readxl: Read Excel Files*. <https://CRAN.R-project.org/package=readxl>.
- Wickham, Hadley, Romain Francois, Lionel Henry, and Kirill Müller. 2017. *Dplyr: A Grammar of Data Manipulation*. <https://CRAN.R-project.org/package=dplyr>.
- Wickham, Hadley, and Lionel Henry. 2017. *Tidyr: Easily Tidy Data with 'Spread()' and 'Gather()'*

Functions. <https://CRAN.R-project.org/package=tidyr>.
Wickham, Hadley, Jim Hester, and Romain Francois. 2017. *Readr: Read Rectangular Text Data*.
<https://CRAN.R-project.org/package=readr>.